

15 – WEB INFORMATION RETRIEVAL

CS 1656

Introduction to Data Science

Alexandros Labrinidis – <http://labrinidis.cs.pitt.edu>

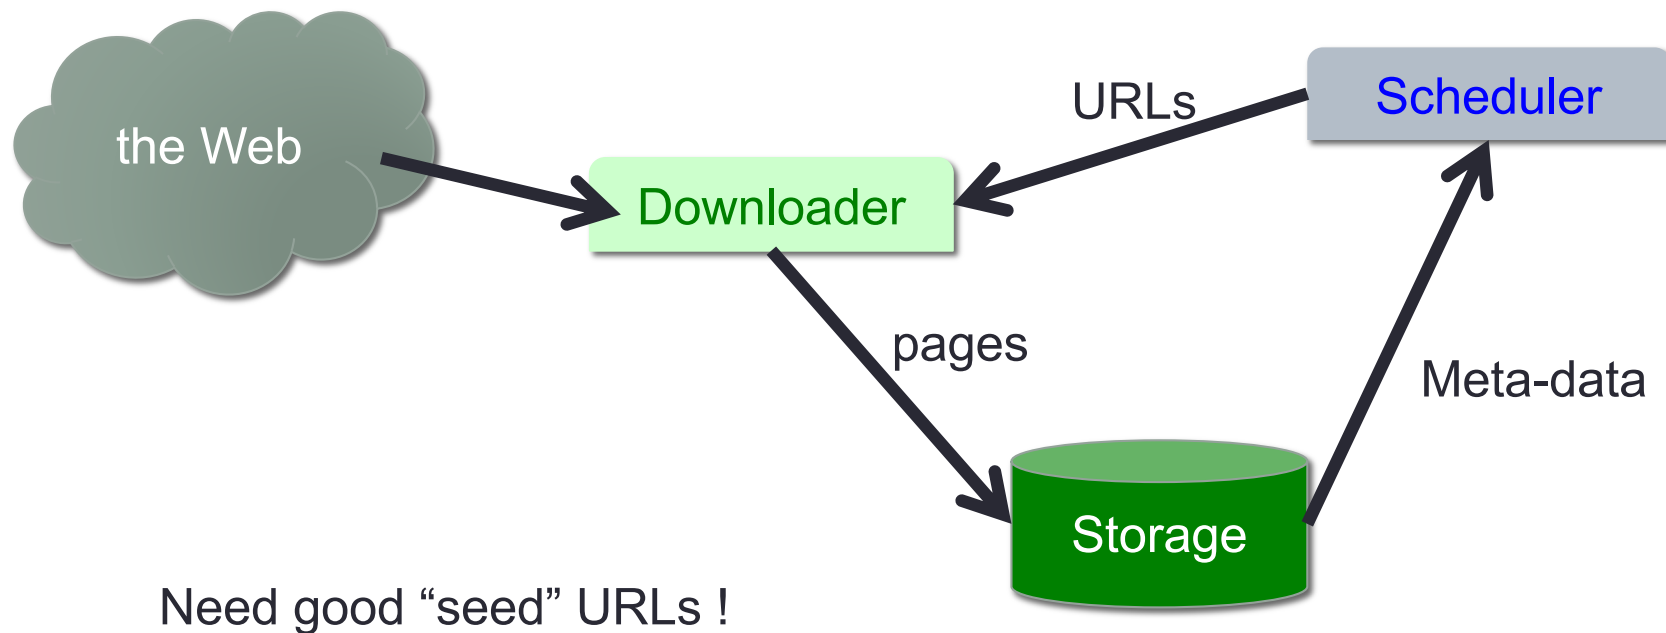
University of Pittsburgh

WEB CRAWLING

Web Information Retrieval

- Q) How to locate relevant web pages?
- A) Maintain a **directory**
- A1) Human-generated web directories
 - Rely on human users (experts or not experts) to submit web pages
 - Examples: Yahoo Web Directory <https://dir.yahoo.com/>
 - Open Directory Project <http://www.dmoz.org>
http://www.dmoz.org/Society/People/Personal_Homepages/L/
- A2) Machine-generated web directories or index
 - Use web crawler to collect web pages and store locally
 - Examples: Google, Bing, Yahoo Search

Anatomy of a Web Crawler



Can we store the entire Web?

- No, but why?
- Reason #1: Storage
- Estimate for size of Web is ASTRONOMICAL
 - 45 billion pages
 - [Source: <http://www.worldwidewebsize.com/>]
- Solution:
 - Store metadata, instead of storing entire web page

Can we store the entire Web?

- Reason #2: Cannot access **entire** Web!
- Because not sure where all pages are
 - E.g., if a web page has no other page linking to it
 - Could be on purpose (dark net)
- Because data is access-controlled
 - E.g., password protected slides
- Because data is made available only after filling out forms
 - Also known as the **Deep Web / Hidden Web**
 - E.g., Pitt user directory (<http://find.pitt.edu>)

Mechanics of Web Crawlers

- HTTP Protocol – Request

```
telnet db.cs.pitt.edu 80
```

- > Trying 136.142.50.166...
- > Connected to db9.cs.pitt.edu.
- > Escape character is '^]'.

```
GET /courses/cs1656/fall2016/test12.html HTTP/1.0
```

(need to hit return twice)

- HTTP Protocol – Response
 - (continued on next page)

Mechanics of Web Crawlers – II

- HTTP/1.1 200 OK ← STATUS CODE
- Date: Wed, 14 Sep 2016 03:16:12 GMT
- Server: Apache/2.2.3 (CentOS)
- Last-Modified: Wed, 14 Sep 2016 03:15:35 GMT
- ETag: "1f773c70-af-53c6f2480a3c0"
- Accept-Ranges: bytes
- Content-Length: 175
- Connection: close
- Content-Type: text/html; charset=UTF-8

- <html>
- <head> <title>This is a test file</title> </head>
- <body>
- <h1> This is a test file </h1>
- Created exclusively for the CS1656 class on September 14, 2016.
- </body> </html>
- Connection closed by foreign host.

DNS Lookup

- DNS = Domain Name Service
 - holds mapping of domain names (e.g., db.cs.pitt.edu) to IP addresses (e.g., 136.142.50.166) and vice versa

```
> nslookup 10.228.27.74
```

```
Server:          136.142.57.10  
Address:         136.142.57.10#53
```

```
74.27.228.10.in-addr.arpa      name = ipsec-10-228-27-  
74.vpn.pitt.edu.
```

```
> nslookup db.cs.pitt.edu
```

```
Server:          136.142.57.10  
Address:         136.142.57.10#53
```

```
db.cs.pitt.edu canonical name = db9.cs.pitt.edu.  
Name:   db9.cs.pitt.edu  
Address: 136.142.50.166
```

Web Server Access Logs

```
> ssh elements.cs.pitt.edu
```

```
> wget
```

```
http://db.cs.pitt.edu/courses/cs1656/fall2016/test12.html
```

```
> grep test12 db.access_log
```

```
136.142.227.11 - - [13/Sep/2016:23:23:45 -0400] "GET  
/courses/cs1656/fall2016/test12.html HTTP/1.0" 200 175
```


```
> nslookup 136.142.227.11
```

```
Server:          136.142.57.10
```

```
Address:         136.142.57.10#53
```

```
11.227.142.136.in-addr.arpa      name = hydrogen.cs.pitt.edu.
```

The Robots Exclusion Protocol



- Web crawlers are not always wanted by web site owners
 - Q: can you think of examples?
 - A: craigslist, eBay, other shopping web sites
- As a result, /robots.txt indicates what is allowed and what is not allowed to be fetched by web crawlers
 - Example:
User-agent: *
Disallow: /
- **Note**: web crawlers can misbehave and ignore robots.txt
- **Note**: robots.txt file is publicly available

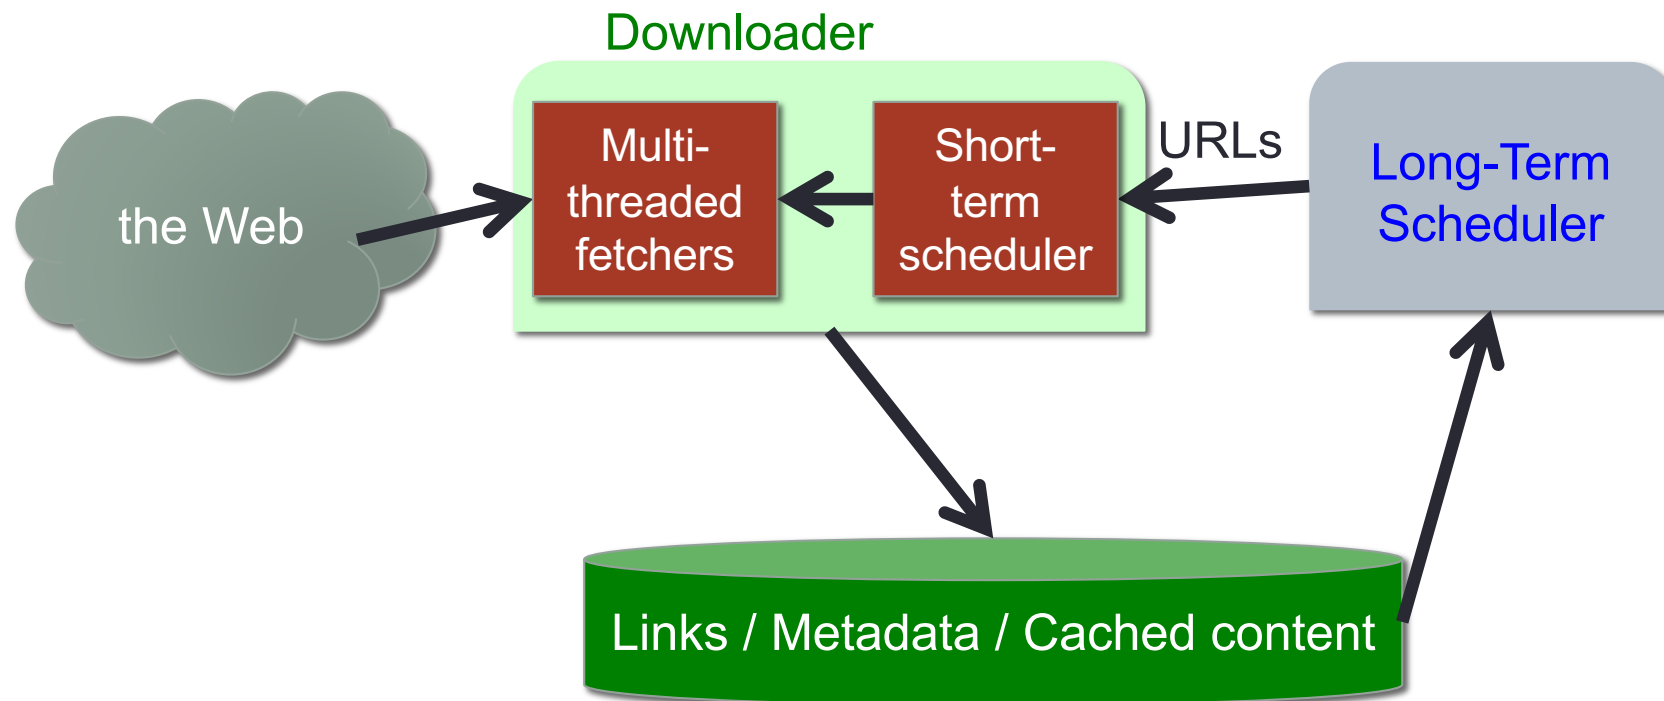
TYPES OF WEB CRAWLERS

Not all crawlers are created equal

- Create a full index of the Web
 - E.g., googlebot, bingbot
- Vertical crawling
 - E.g., news bots, spambots
 - E.g., based on file types
- Focused crawling
 - E.g., based on driving query
- Mirror a specific page / check for updates
 - E.g., <https://www.changedetection.com/>

SCHEDULING

Anatomy of a Web Crawler – II



How to schedule web page crawls?

- **Q1: How often to crawl a page?**
 - Too often → wasted bandwidth and crawler resources
 - Too sparse → missed potentially important updates
- Long-term scheduling:
 - Determine based on page quality/freshness estimations
- Short-term scheduling:
 - Determine based on politeness policy
 - In order not to overwhelm web site
 - Spread requests over time
 - Spread requests across multiple web sites

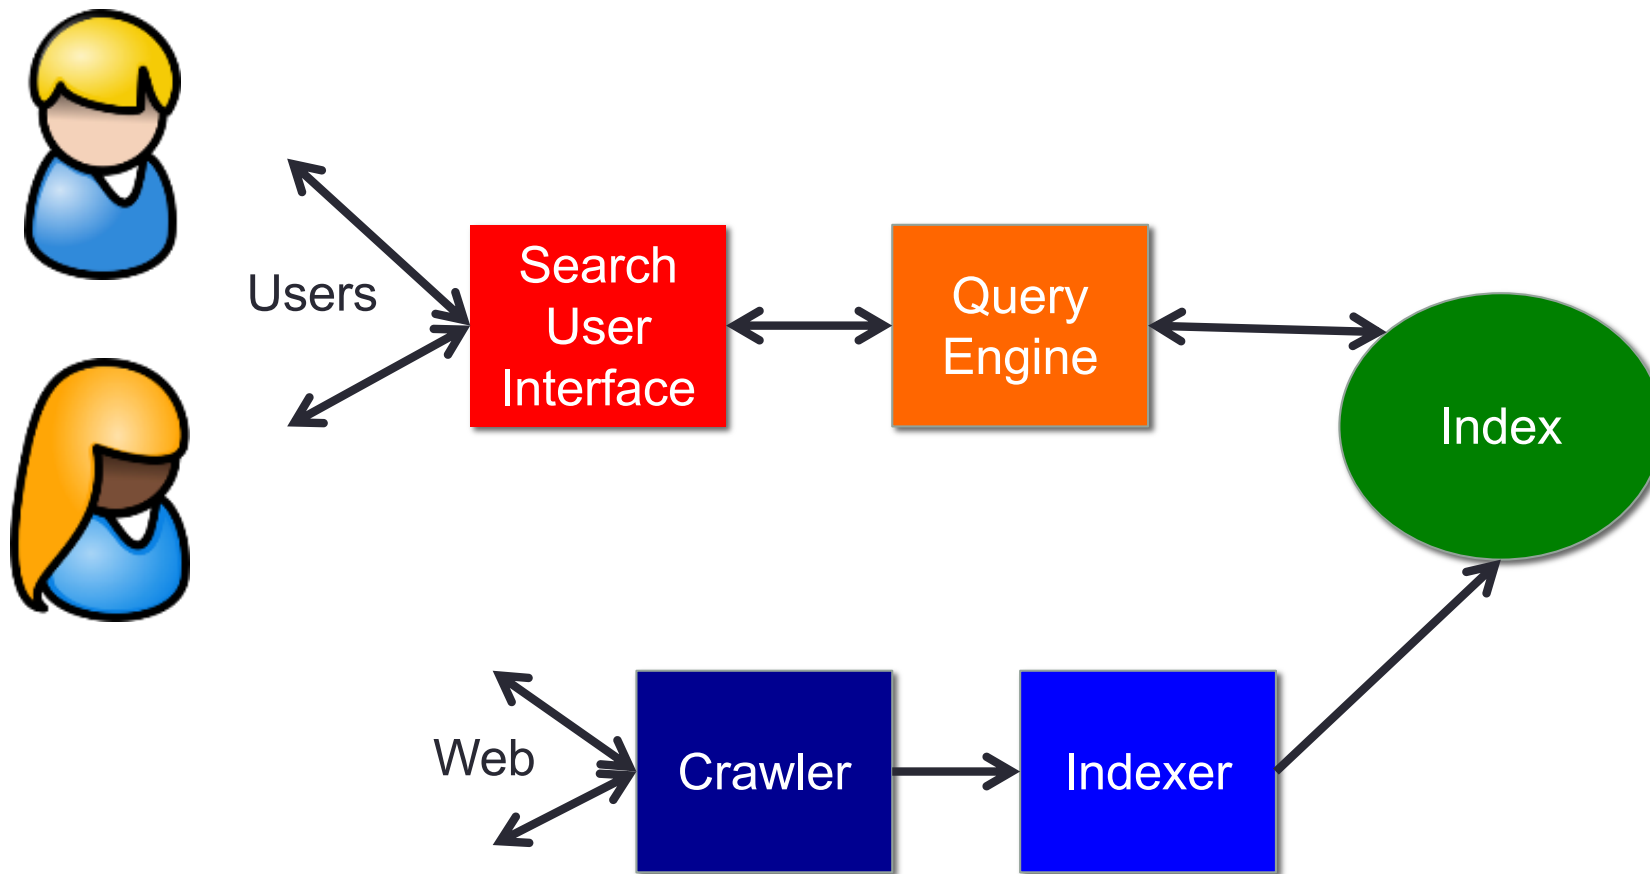
Crawling the Deep Web

- Also called the *Hidden Web*
- Need to provide form input for page to be generated
 - E.g., weather for 15213 zip code
- General web crawling of Deep Web is very difficult
- Special cases for known forms/input:
 - E.g., can request same form and see if results changed
 - Called *screen scraping*
 - Challenging for pages with Javascript, especially AJAX

[More info: http://en.wikipedia.org/wiki/Web_crawler]

WEB INFORMATION RETRIEVAL

Web Information Retrieval Architecture



Search Engine Ranking

- Simple approach: use relevance ranking (TF-IDF)
- Question: What are the challenges / shortcomings?
- Answers:
 - Web site owners can “pad” pages with more keywords
 - Web site owners can present different pages to crawler vs to humans
 - Are not considering web page importance
 - Are not taking advantage of hypertext link structure

How about using link structure?

- Main idea:
 - View incoming links as “**votes**” of confidence for web page
 - Similar to citations in academic publishing
- Important question remains:
 - How to count incoming links – are all equally important?
- Must have different levels of importance:
 - Link from authoritative web site (e.g., nytimes.com)
 - vs
 - Link from one’s own web site (e.g., iloveapples.com)


In-classroom activity

- Students organize into teams
 - The most popular team wins
- Each team gets **unlimited “votes”**
- Each team decides **how many votes** to “cast”
 - One or multiple votes per team
- Each team decides **where to give** votes to
 - Vote can go to any team (including own / duplicates)
- What voting strategy do you think is better?

In-classroom activity

- Vote only for your team once
- Give multiple votes to your team
- Give multiple votes to another team
- Give one vote to your team and another to another team
- Give multiple votes to multiple teams (not including your team)
- Give multiple votes to multiple teams (also voting for your team)


Google's PageRank – Receivers



$$PR(a) = (1 - d) + d(V_1 + V_2 + \dots + V_n)$$

PageRank of page a is the sum of values of incoming links

Google's PageRank – Senders



How much “value” should each page give?

$$V_i = \frac{PR(b)}{C(b)}$$

PageRank of page b


Number of outgoing links from page b

Putting it all together

- PageRank of page **a** that has **n** incoming links from pages **T₁** to **T_n** is given by the following formula:

$$PR(a) = (1 - d) + d \left(\frac{PR(T_1)}{C(T_1)} + \frac{PR(T_2)}{C(T_2)} + \dots + \frac{PR(T_n)}{C(T_n)} \right)$$


A simple example



Notice the directional edges!

- Guess initial values of 1.0, i.e., $PR(A) = PR(B) = 1.0$
- $d=0.85$
- $PR(A) = (1-d) + d(PR(B) / 1) = 0.15 + 0.85*1 = 1$
- $PR(B) = (1-d) + d(PR(A) / 1) = 0.15 + 0.85*1 = 1$


A simple example – Take 2



Notice the directional edges!

- Guess initial values of 0.0, i.e., $PR(A) = PR(B) = 0.0$
- **$d=0.85$**
- $PR(A) = (1-d) + d(PR(B) / 1) = 0.15 + 0 = 0.15$
- $PR(B) = (1-d) + d(PR(A) / 1) = 0.15 + 0.85 \cdot 0.15 = 0.2775$
- $PR(A) = 0.15 + 0.85 \cdot 0.2775 = 0.385875$
- $PR(B) = 0.15 + 0.85 \cdot 0.385875 = 0.47799375$
- $PR(A) = 0.15 + 0.85 \cdot 0.47799375 = 0.5562946875$
- $PR(B) = 0.15 + 0.85 \cdot 0.5562946875 = 0.622850484375$

A simple example – Take 3



Notice the directional edges!

- Guess initial values of 40, i.e., $PR(A) = PR(B) = 40.0$
- **$d=0.85$**
- $PR(A) = (1-d) + d(PR(B) / 1) = 0.15 + 0.85*40 = 34.15$
- $PR(B) = (1-d) + d(PR(A) / 1) = 0.15 + 0.85*34.15 = 29.1775$
- $PR(A) = 0.15 + 0.85 * 29.1775 = 24.950875$
- $PR(B) = 0.15 + 0.85 * 24.950875 = 21.35824375$
- Seems numbers will get to 1.0 and stop

More examples

- https://webworkshop.net/seo-tools/pagerank_calculator


Modern search engine ranking

- Modern search engines utilize multiple signals:
 - PageRank
 - Name of query term in domain name
 - Location
 - https vs http
- This is a result of many years of improvements
- And reacting to Search Engine Optimization techniques!
 - E.g., link farms

[More info: http://en.wikipedia.org/wiki/Search_engine_optimization]

[More info: http://en.wikipedia.org/wiki/Link_farm]

Understanding Question



- **Question:**
 - Which of these pages will have a higher PageRank value?
- **Possible Answers:**
 - A
 - B
 - C
 - D