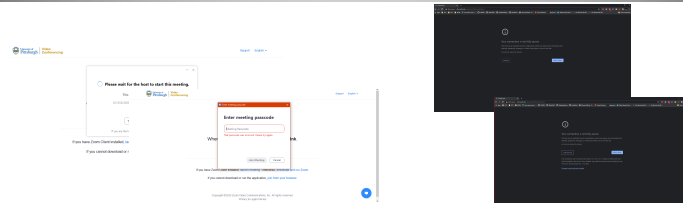


Announcements



- ❑ Zoom links issues on Friday
 - Zoom link for Recitations now available on Canvas and the website
 - No credit is lost 😊
- ❑ Website TLS issue
 - Please click the advance option and click [Proceed to db.cs.pitt.edu \(unsafe\)](https://db.cs.pitt.edu)

Overview of Database Management Systems

Part II

Database Vs. File Systems Approaches

- ❑ **Abstraction**
- ❑ **Reliability**
- ❑ **Efficiency/Performance**



Data Abstraction

- ❑ Data are structured in a way meaningful to applications
- ❑ **Data Model:**
 - A collection of high-level data description constructs that hide low-level storage details
- ❑ **The Relational / Object-Relational Model:**
 - Is the most widely used data model today
 - Main construct is a **relation**: table of records
 - Every relation has a **schema**:
 - Relation name
 - Names of fields
 - Types of fields

Example

Students				
	<i>SID</i>	<i>Name</i>	<i>Age</i>	<i>GPA</i>
Record or Tuple	546007	Peter	18	3.8
	546100	Bob	19	3.65
	546500	Bill	20	3.7

- Schema:
 - Students (*sid*: string, *name*: string, *age*: integer, *gpa*: real)
- State: Actual data at a given point in time

Database Languages

- **Data Definition Language (DDL):**
 - Define schemas
 - Define **Integrity Constraints**
 - Example: unique *SIDs*
 - More...
- **Data Manipulation Language (DML):**
 - To ask questions = **Query**
 - Example: Which students have GPA > 3.75?
 - To create and modify data
- **SQL:** Most widely used database language

Good DBMS ≠ Good Design!

Students				
	<i>SID</i>	<i>Name</i>	<i>Age</i>	<i>GPA</i>
Record or Tuple	546007	Peter	18	3.8
	546100	Bob	19	3.65
	546500	Bill	20	3.7

- Schema:
 - Students (*sid*: string, *name*: string, *age*: integer, *gpa*: real)
 - Alternative Schema:
 - Students (*sid*: integer, *lname*: string, *fname*: string, *dob*: date, *gpa*: real)

Now we know the color of the Database

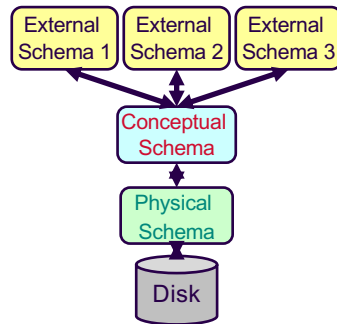


Levels of Data Abstraction in a DBMS

- The data in a DBMS is described at three levels of abstraction:

1. **Conceptual Schema**
DDL: data definition language
2. **Physical Schema**
SDL: storage definition language
3. **External Schema (Views)**
VDL: view definition language

- Many external schemas
- plus one conceptual
- plus one physical



External Schema - Views

- Allows data access to be customized and authorized at the user-level

- Defined in terms of data model
- Consists of a collection of **views**

- Example:

– CourseEnroll (*cid*: string, *enrollment*: int)

<i>CID</i>	<i>Enroll</i>
CS 1555	37
CS 2550	17
CS 3551	15

- Guided by end-user requirements
- Views are computed as needed
- Multiple Views of data allows each user/application to get different perspective of the database

3-level Architecture

- view level
 - CSMajors
 - MathMajors
- logical level: entire database schema
 - Courses (CourseNo, CourseName, Credits, Dept)
 - Students (StudentID, Lname, Name, Class, Major)
 - GradeReport (StudentID, CourseNo, Grade, Term)
- physical level:
 - how these tables are stored
 - how many bytes and attribute, etc.

Execution Abstraction


- A **transaction** is a **logical unit of work** in DBMSs
 - It is the execution of a **program segment** that performs some function or task by accessing shared data (e.g., a db)
 - logical grouping of query and update requests needed to perform a task
- Examples:
 - deposit, withdraw, transfer money (banking transaction)
 - reserve a seat on a flight (airline reservation)
 - print monthly payment checks (business transaction)
 - update inventory (inventory transaction)

ACID Properties

- ❑ **Atomicity** (alias failure atomicity)
 - Either all the operations associated with a transaction happen or none of them happens
- ❑ **Consistency Preservation**
 - A transaction is a correct program segment. It satisfies the integrity constraints on the database at the transaction's boundaries
- ❑ **Isolation** (alias concurrency atomicity / serializability)
 - Transactions are independent, the result of the execution of concurrent transactions is the same as if transactions were executed serially, one after the other
- ❑ **Durability** (alias persistence / permanence)
 - The effects of completed transactions become permanent surviving any subsequent failures

Transfer Money Example

- ❑ Two accounts:
 - A: \$100
 - B: \$200
- ❑ Client 1: transfer \$6 from A to B




Client 1
Read A: \$100
Write A: \$94
BOOM
Read B: \$200
Write B: \$206

ACID Properties

- ❑ **Atomicity** (alias failure atomicity)
 - Either all the operations associated with a transaction happen or none of them happens
- ❑ **Consistency Preservation**
 - A transaction is a correct program segment. It satisfies the integrity constraints on the database at the transaction's boundaries
- ❑ **Isolation** (alias concurrency atomicity / serializability)
 - Transactions are independent, the result of the execution of concurrent transactions is the same as if transactions were executed serially, one after the other
- ❑ **Durability** (alias persistence / permanence)
 - The effects of completed transactions become permanent surviving any subsequent failures

Transfer Money Example [Poll]

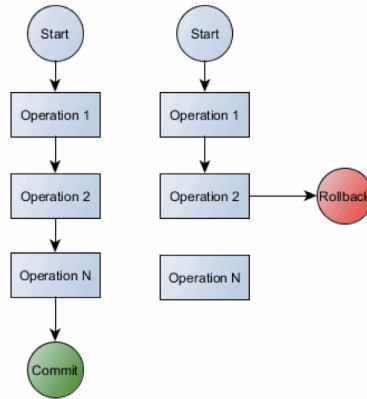
- ❑ Three accounts:
 - A: \$100
 - B: \$200
 - C: \$300
- ❑ Client 1: transfer \$6 from A to B
- ❑ Client 2: transfer \$4 from C to B



Client 1	Client 2
Read A: \$100	
Write A: \$94	
	Read C: \$300
	Write C: \$296
Read B: \$200	
	Read B: \$200
	Write B: \$204
Write B: \$206	

SQL TRANSACTIONS

- Begin:
 - Each SQL statement should *implicitly* start a transaction, unless one is active.
- COMMIT [WORK];
or END [WORK];
- ROLLBACK [WORK] ;
 - ROLLBACK default action



Sample Transfer Money SQL code

<pre>-- Client 1 BEGIN; UPDATE ACCOUNT SET BALANCE = BALANCE - 6 WHERE NAME = 'A'; UPDATE ACCOUNT SET BALANCE = BALANCE + 6 WHERE NAME = 'B'; COMMIT;</pre>	<pre>-- Client 2 BEGIN; UPDATE ACCOUNT SET BALANCE = BALANCE - 4 WHERE NAME = 'C'; UPDATE ACCOUNT SET BALANCE = BALANCE + 4 WHERE NAME = 'B'; COMMIT;</pre>
---	---

ACID in NoSQL Databases



- **ACID to BaSE**: No immediate consistency, data freshness & accuracy
 - **Basic Availability**: The database appears to work most of the time.
 - **Soft-state**: Stores don't have to be write-consistent, nor do different replicas have to be mutually consistent all the time.
 - **Eventual consistency**: Stores exhibit consistency at some later point (e.g., lazily at read time).