



University of Pittsburgh

ECE 1150: Computer Networks

The Transport Layer

Mai Abdelhakim, PhD

ECE Department



Swanson School of Engineering

University of Pittsburgh



This Unit

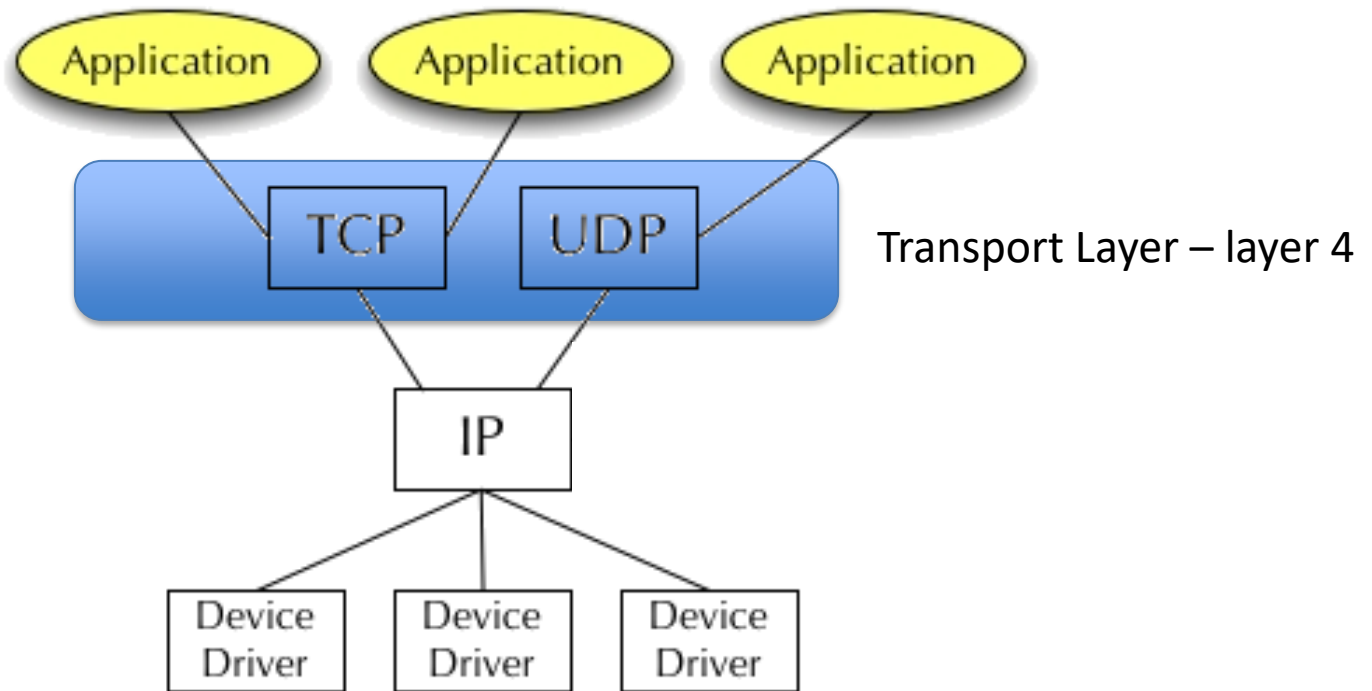
- Transport layer



We are here
(Layer 4)

Transport Layer

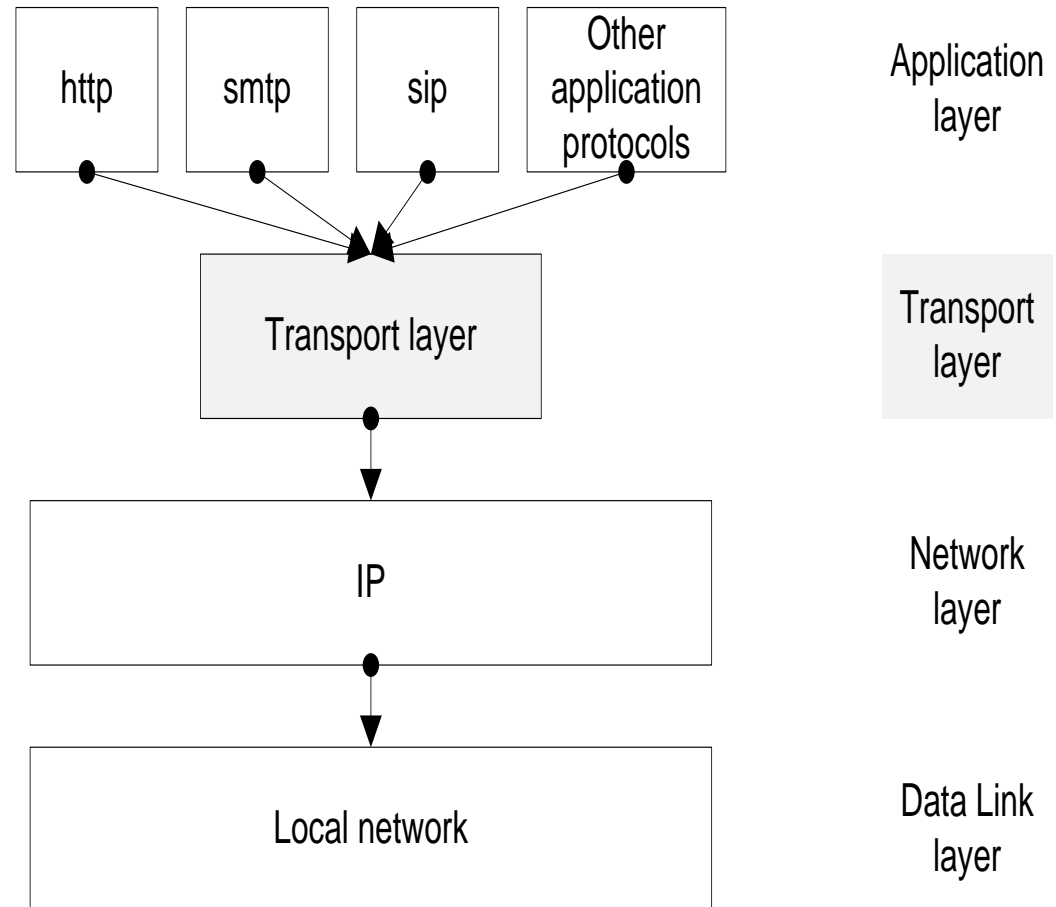
- Two types:
 - TCP: Transport control protocol (reliable)
 - UDP: User Datagram protocol (easy)



Schematic of TCP/IP Operation

TCP Functions

- Therefore TCP has to perform many tasks
 1. Segmentation
 2. Reliability
 3. Flow-control
 4. Multiplexing for applications
 5. Connection establishment



TCP Functions - Segmentation

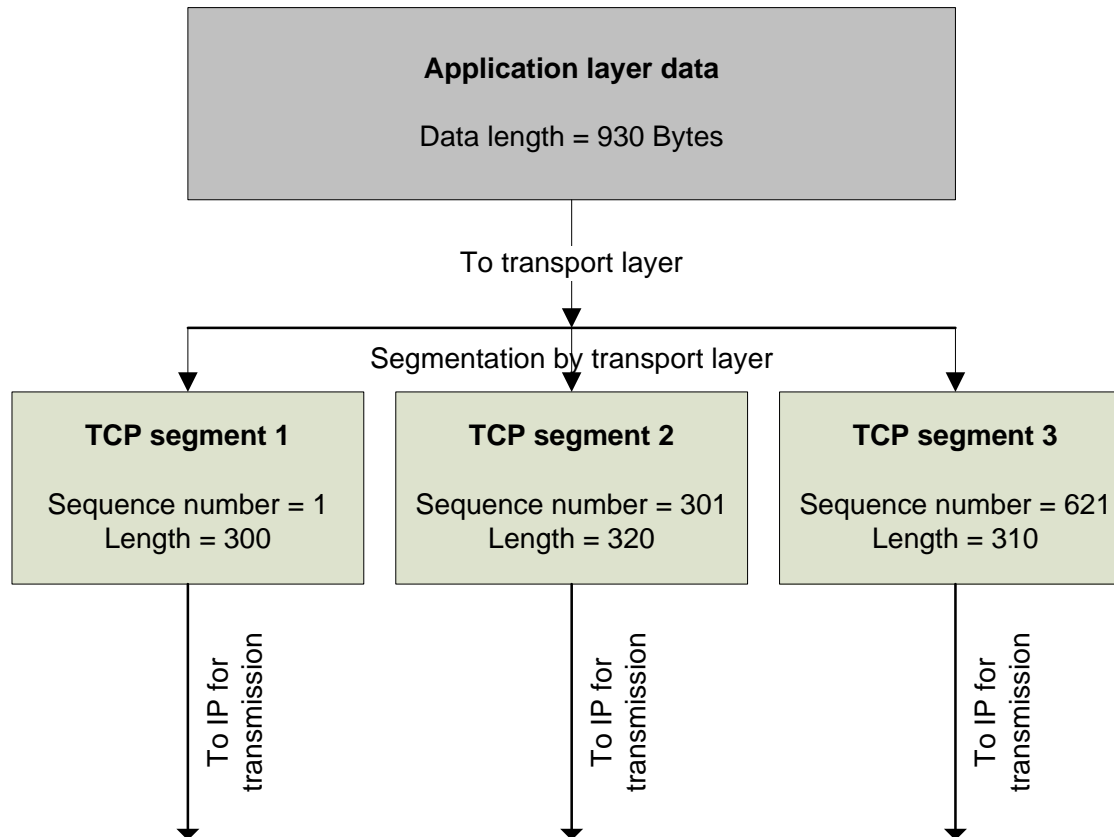
- TCP allows IP to transfer arbitrarily large data blocks
- However, the maximum packet size in network layer is 65,536 bytes
 - What happens if the application wants to send a file of size 5 Mega bytes?

TCP Functions - Segmentation

TCP does segmentation

- At the sender, **break files into smaller blocks** (called segments or datagrams)
- At the receiving end, **re-assemble** these blocks into the file
- A **sequence number** is assigned to each datagram
 - Sequence numbers help **receiver order datagrams** even when received out of order
 - Datagrams may also get **duplicated**
 - Sequence numbers help identify these duplicates

Sequence Numbers



Assume we start
with sequence
number 1

Sequence number of a TCP
segment

=

sequence number of the previous segment
+ length of previous segment

TCP functions - Segmentation

- Advantages of segmentation
 - **Errors are less** likely in smaller segments
 - **Less retransmission** if error is introduced in a segment
 - When error occurs in one segment, no need to transmit the entire file again. Just retransmit the segment.
 - **Easier for routers** to hold segments in memory
 - Reduce complexities in packet handling
- Disadvantages
 - **More computations** in sending and receiving

TCP Functions - Reliability

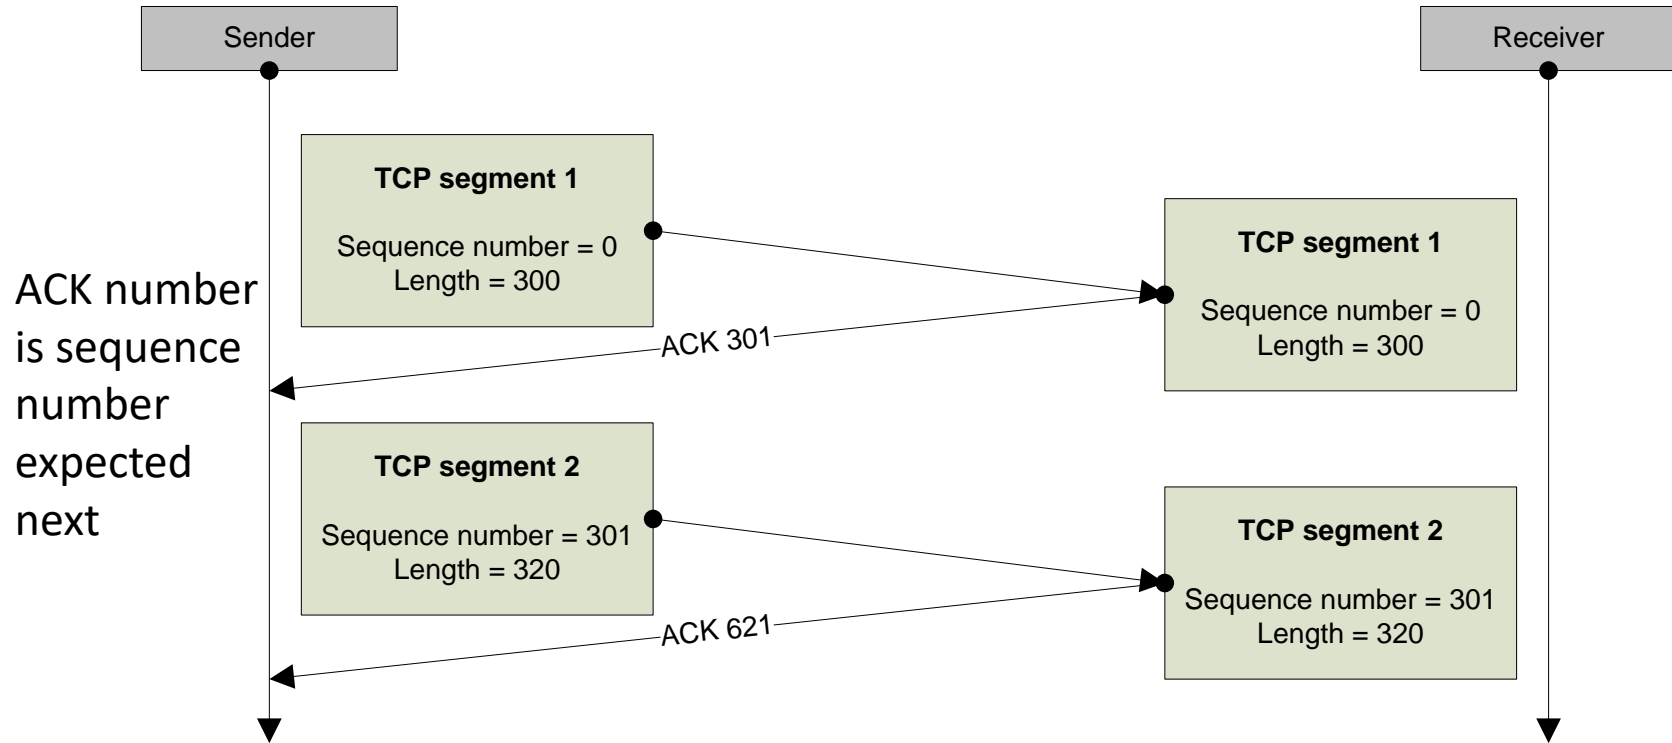
- Datagrams may get damaged during transmission
- **TCP adds a checksum** to detect errors. This checksum is not as robust as CRC.

TCP Functions - Reliability

- **Reliable end-to-end** communication service
- TCP **recovers from network damage** to data
- Basic mechanisms similar to error control at data link layer, but now **applied end-to-end**
 - Receiver sends a positive **acknowledgment (ACK)** if all goes well
 - If the ACK is not received within a **timeout** interval, the sender retransmits the data
- Two techniques: Stop-and-wait & sliding window
- Also used for **flow control**

Simple Flow Control Mechanism

Stop - and - Wait

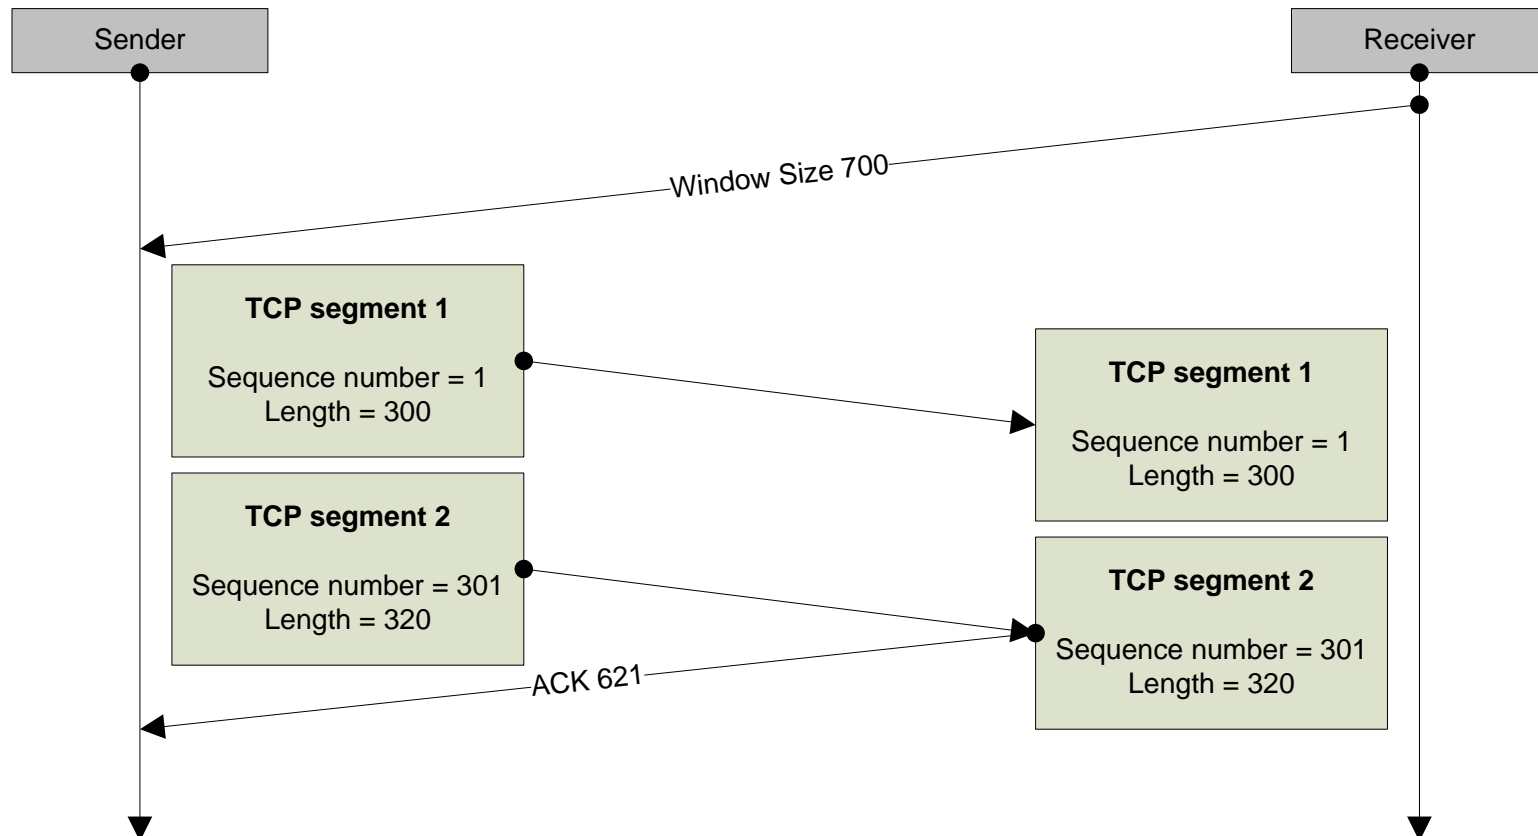


- Flow control mechanism shown here is called stop-and-wait
 - Sender waits for ACK before sending next datagram
- Very slow

Flow Control - Sliding window

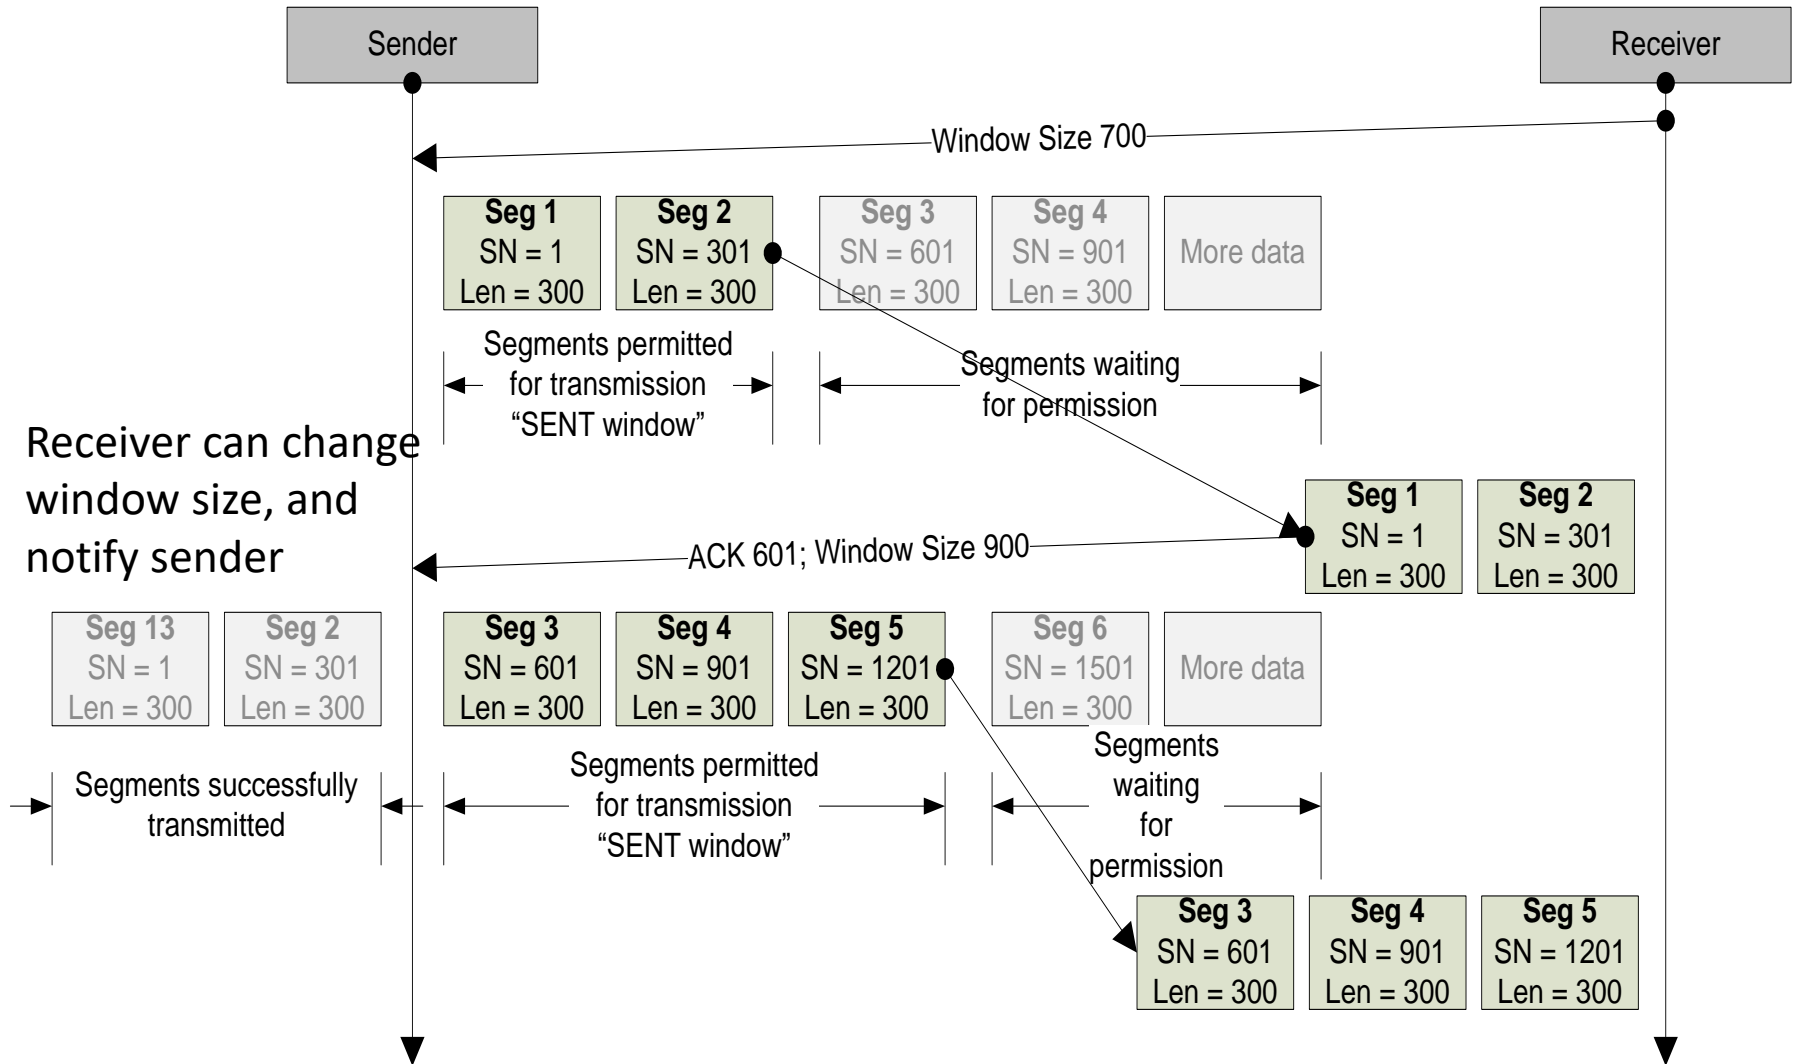
- Flow control: **receiver regulates** the amount of data the sender may send
- Accomplished by specifying a “**window**” with every **ACK**
 - “**Window**” indicates **how many bytes** of data the sender may transmit **before receiving any more ACK**
- Window slides as **receiver acknowledges** packets or modifies window size

TCP Flow Control with Window Size



- Window size is the amount of data the receiver is capable of processing
- Receiver begins by announcing the window size
- ACK 621: means that all data up to byte 620 is received correctly

Sliding Window

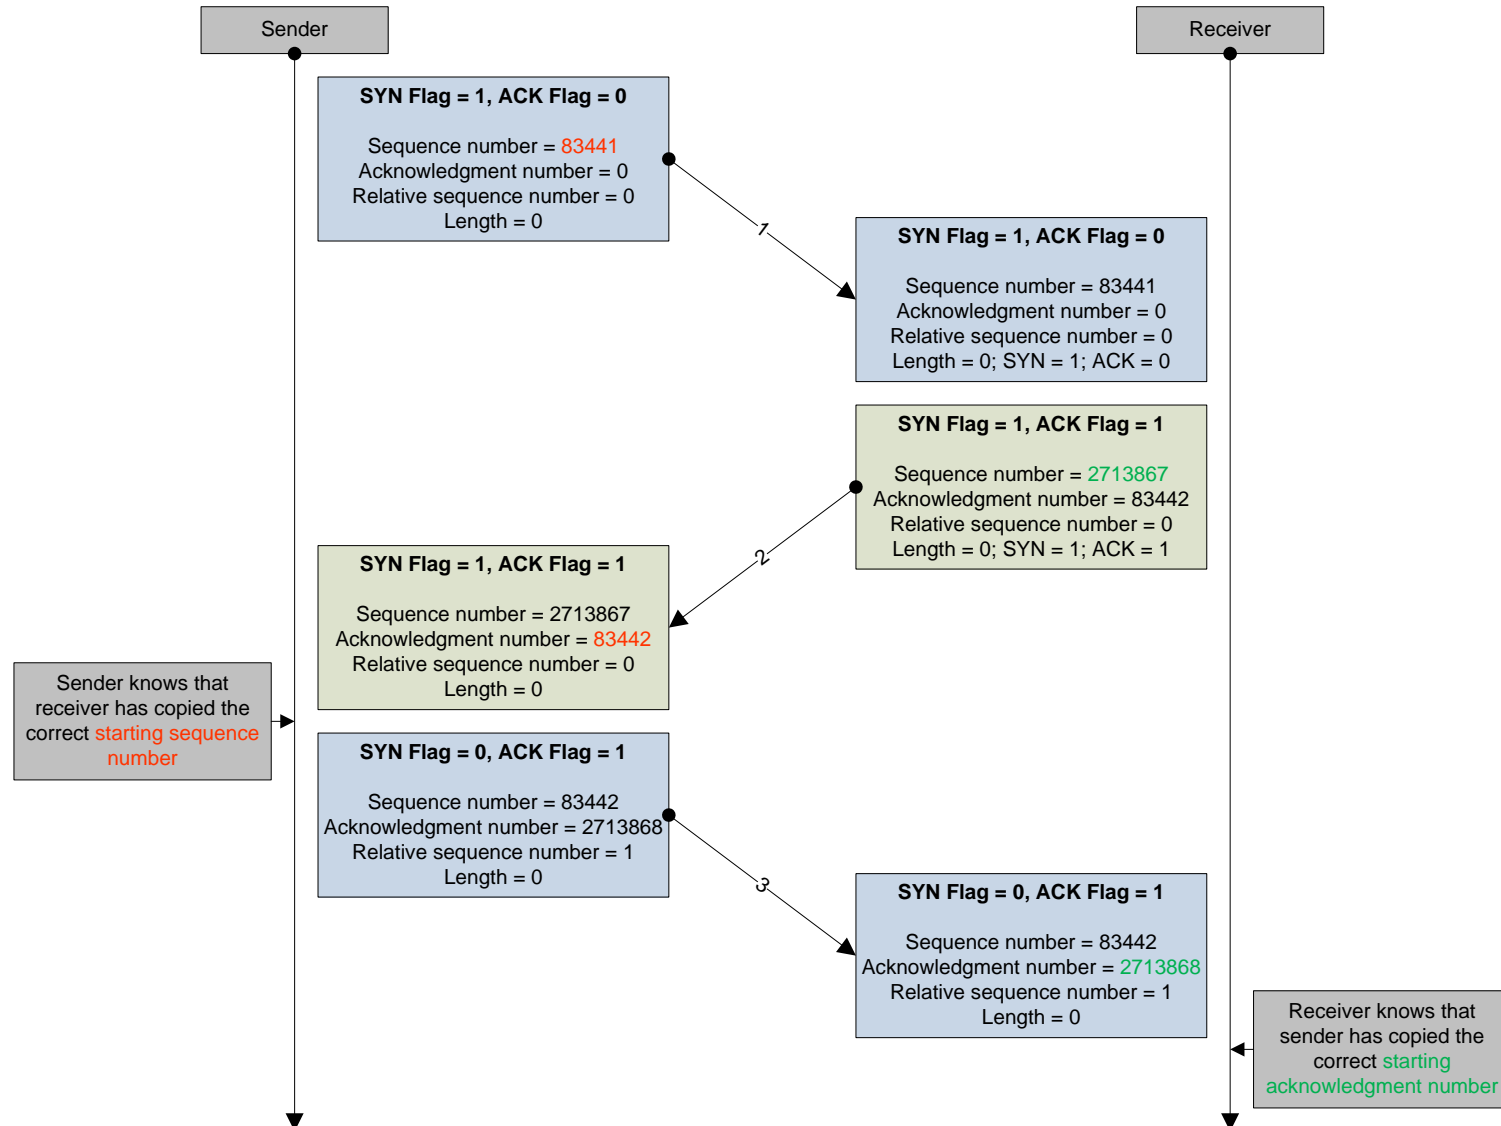


TCP Functions – Connection Establishment

- Sequence numbers are core part of TCP
- It is not a good idea to reuse the same sequence numbers in succession
 - Creates problems in detecting duplicates
- Hence, before communication starts, sender and receiver **negotiate the initial sequence numbers to use in TCP connection**
 - **Called 3-Way Handshake, which is made to establish a connection**

3-Way Handshake

1. Sender generates Initial Sequence Number (ISN).
E.g. ISN = 83441.
2. Receiver ACKs this reception by putting the (**ACK** = sender ISN + 1), and **inserts its own Rx ISN**.
3. When the sender receives the ACK, it also knows the ISN of the Rx, it then sends an ACK



Note: Acknowledgment number is increased by 1 if SYN or ACK flags are set

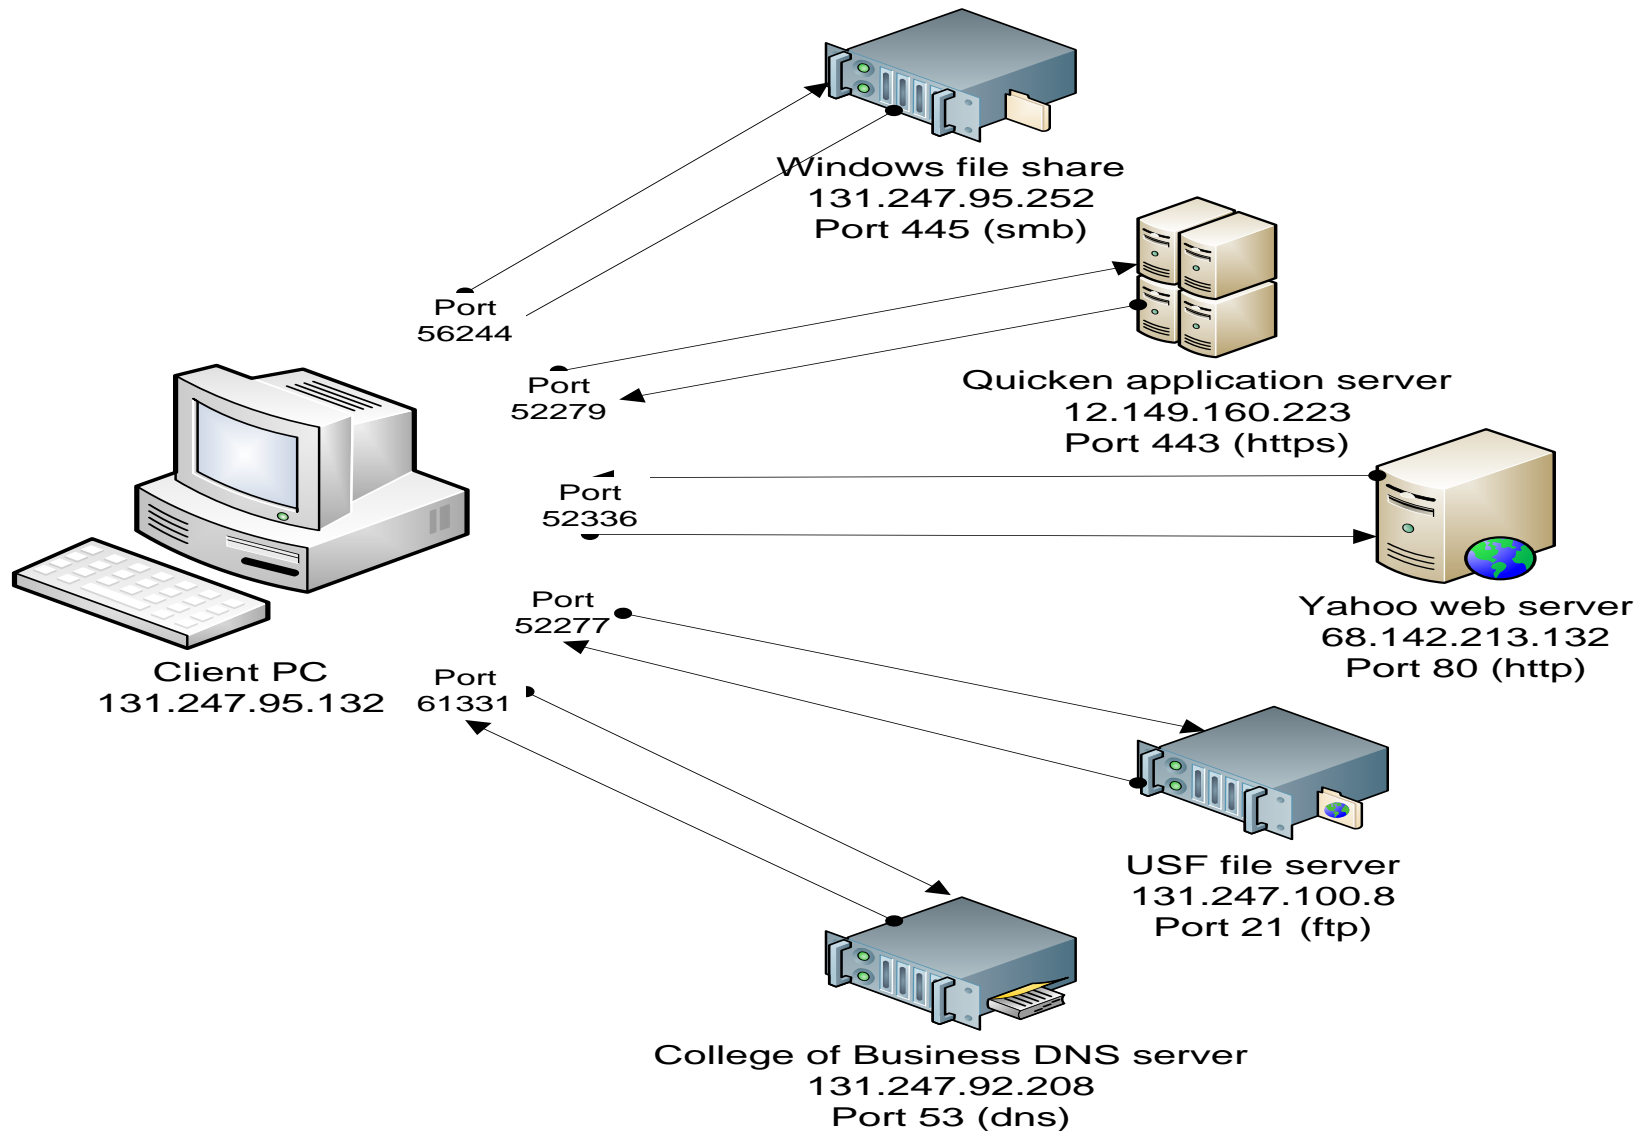
Connection Establishment with 3-Way Handshake

- Three way handshake
 1. Sender generate a random initial sequence number (ISN) and sends it to destination
 2. Destination generates a random initial sequence number and sends it along with the acknowledgment to sender
 3. Sender acknowledge

TCP Functions - Multiplexing

- Transport layer performs **application multiplexing**
 - Allowing multiple applications to use the same network interface card.
- This is accomplished by **assigning different port numbers** to different applications within a host
 - Each **application** that needs a network connection (each browser tab) is assigned a **port number**
 - Port number is 16 bits
- An IP address and port address together is called a **socket**

Port addresses and multiplexing



TCP Port Assignment

- On the **sender** side, the **operating system** assigns one of the free ports to an application that requires network connectivity
- How do you know what port to connect to on the receiver side?
 - On what port is the web server listening?
 - Standard ports
 - Assigned by IANA: Internet Assigned Numbers Authority
 - Common ports: <https://www.utilizewindows.com/list-of-common-network-port-numbers/>
 - 80 : http
 - 443: SSL (https)
 - 53: DNS
 - 23: Telnet

Viewing Ports Usage With Netstat Command

- Netstat: Utility provided by OS to view port usage

```
C:\>netstat -n
```

Active Connections

Foreign address: IP addresses the local device is connected to

Proto	Local Address	Foreign Address	State
TCP	127.0.0.1:5354	127.0.0.1:46133	ESTABLISHED
TCP	127.0.0.1:46133	127.0.0.1:5354	ESTABLISHED
TCP	192.168.1.5:46122	192.168.1.153:445	ESTABLISHED
TCP	192.168.1.5:46149	74.125.67.138:443	CLOSE_WAIT
TCP	192.168.1.5:46151	74.125.45.104:80	CLOSE_WAIT
TCP	192.168.1.5:51117	131.247.80.160:143	ESTABLISHED
TCP	192.168.1.5:51124	131.247.80.160:143	ESTABLISHED
TCP	192.168.1.5:51292	131.247.16.141:22	ESTABLISHED
TCP	192.168.1.5:51369	131.247.80.29:445	SYN_SENT

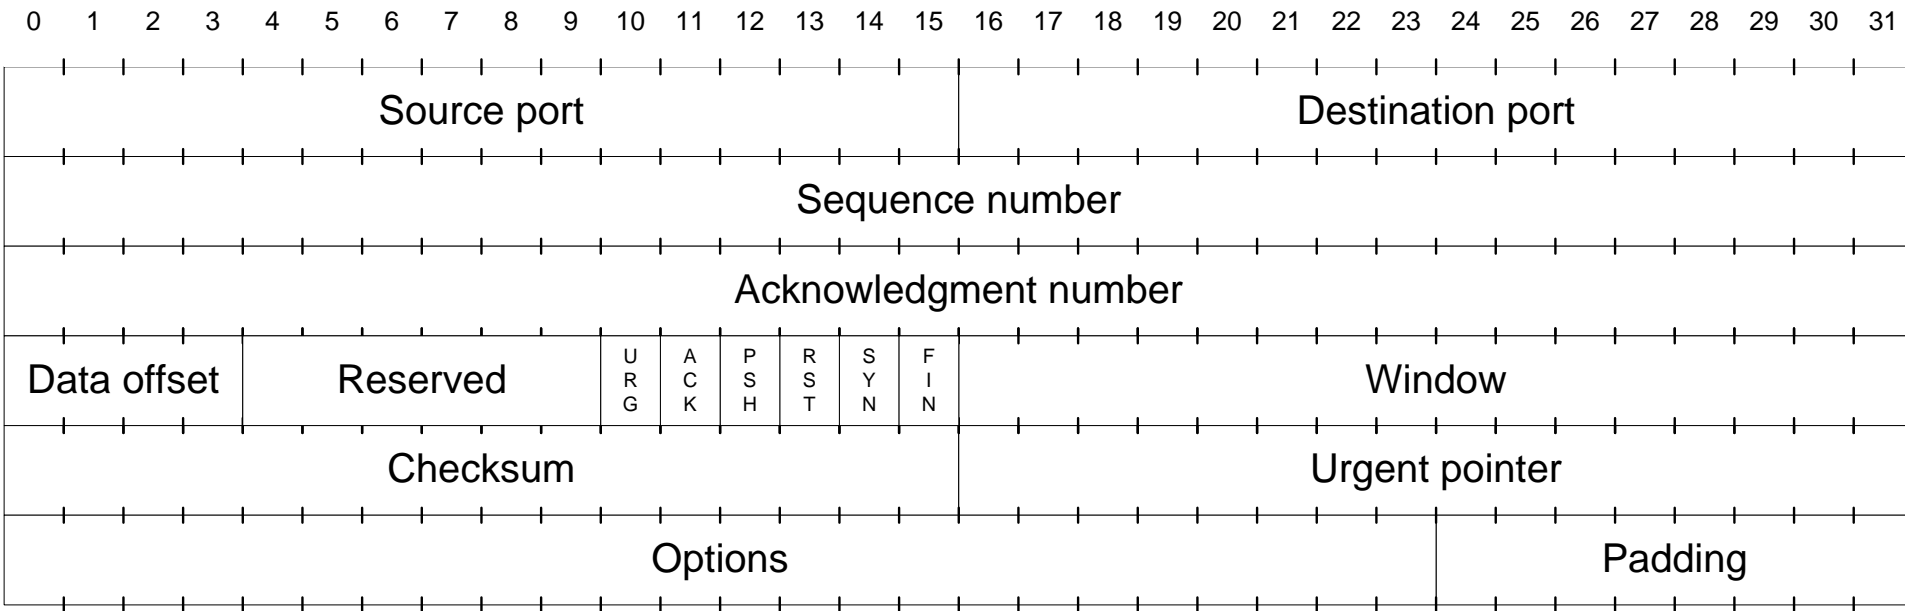
CLOSE_WAIT: Wait to request of terminating connection

Ports on local device

Ports on remote servers

https://en.wikipedia.org/wiki/Transmission_Control_Protocol

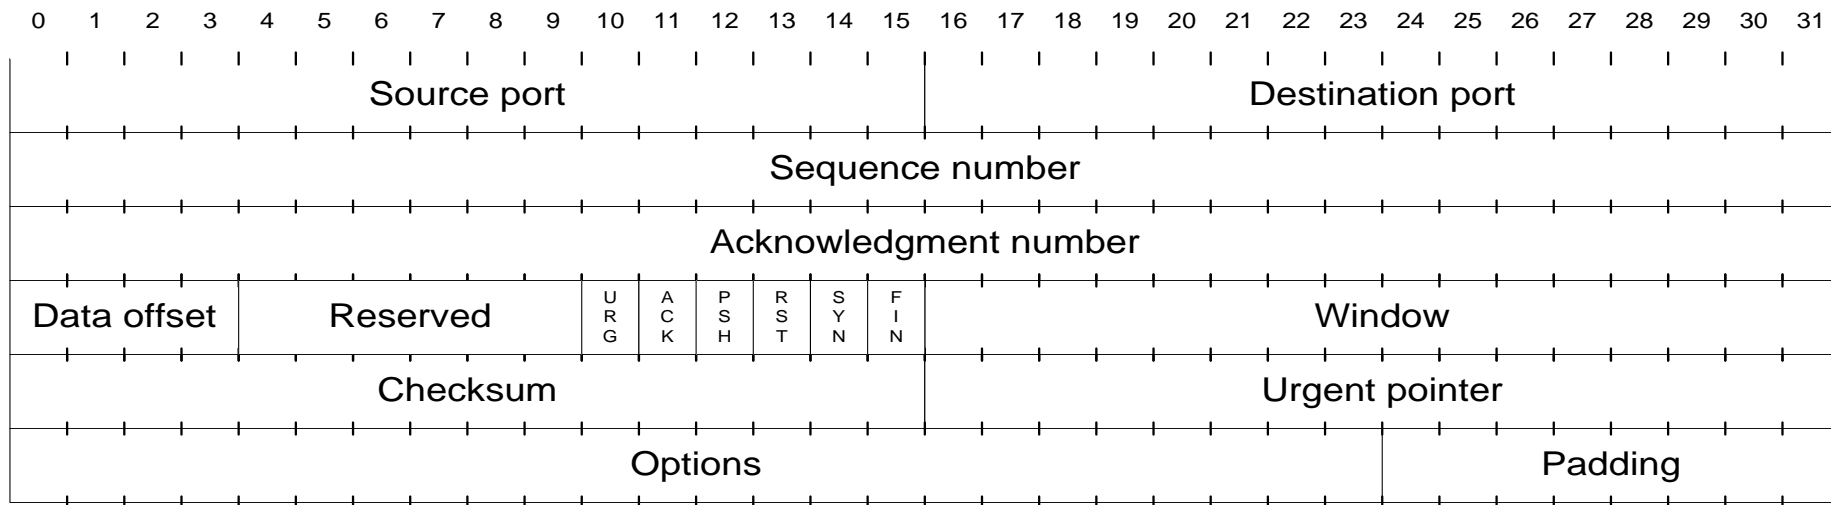
TCP Header



Note: Each tick mark represents a bit position

TCP Header Fields

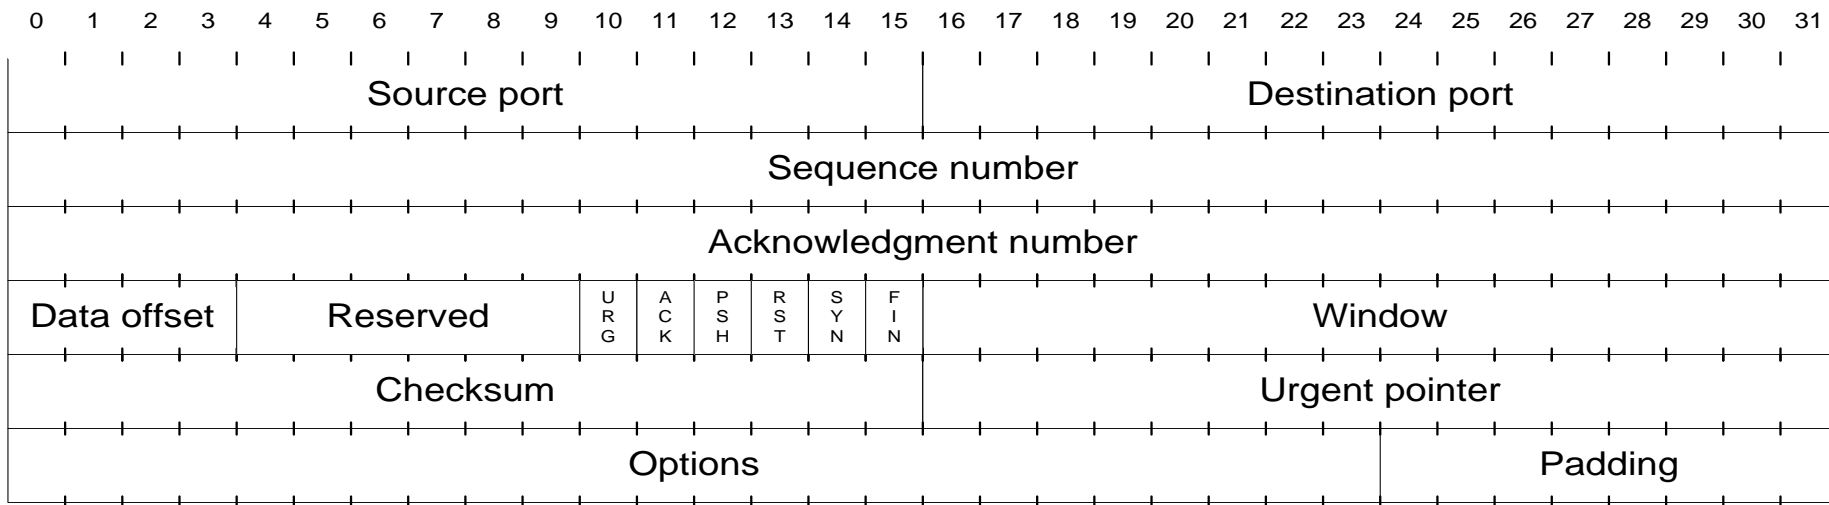
- Port addresses have 16 bits
 - $2^{16} = 65,536$ ports possible per host
 - If computing resources are available, a single computer running TCP can support 65,536 simultaneous network connections



Note: Each tick mark represents a bit position

TCP Header Fields

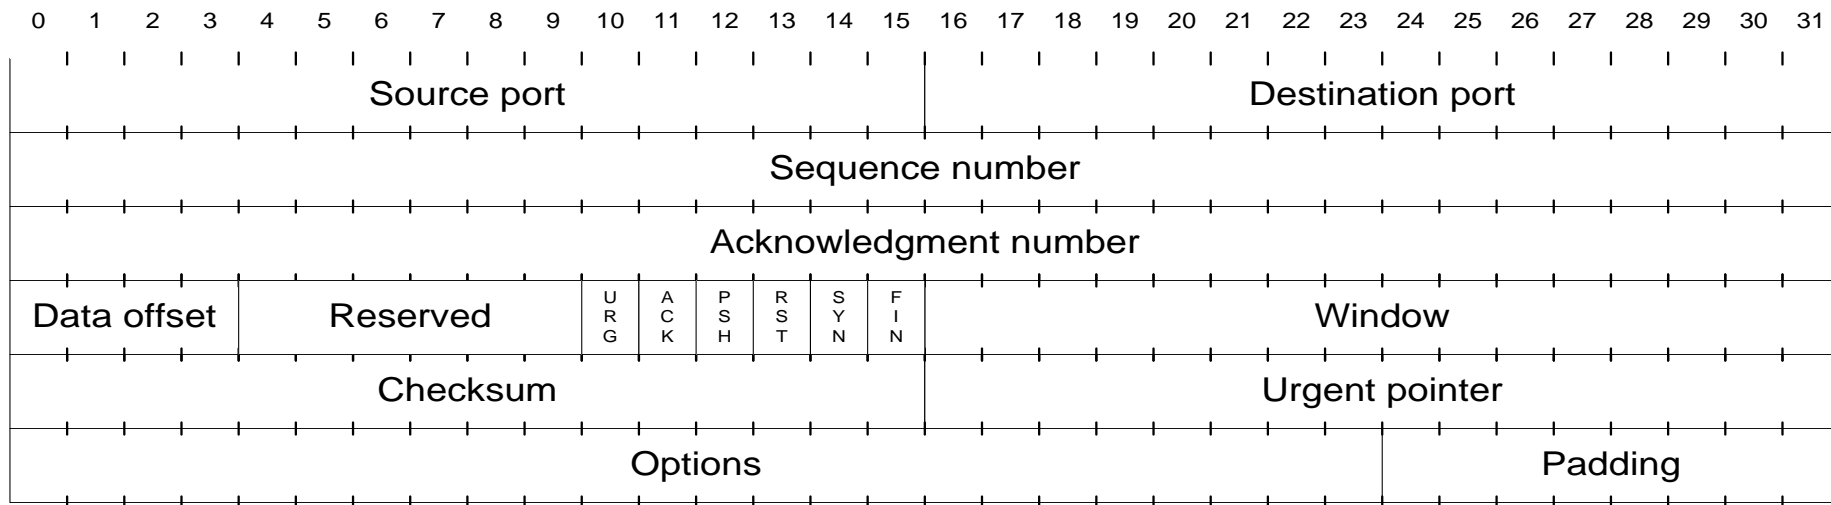
- Sequence and acknowledgment numbers have 32 bits
 - Sequence number** is the sequence number of the **first data byte** in the datagram
 - Acknowledgments are cumulative
 - ACK 1079** implies that all data till byte number **1078** have been received correctly



Note: Each tick mark represents a bit position

TCP Header Fields

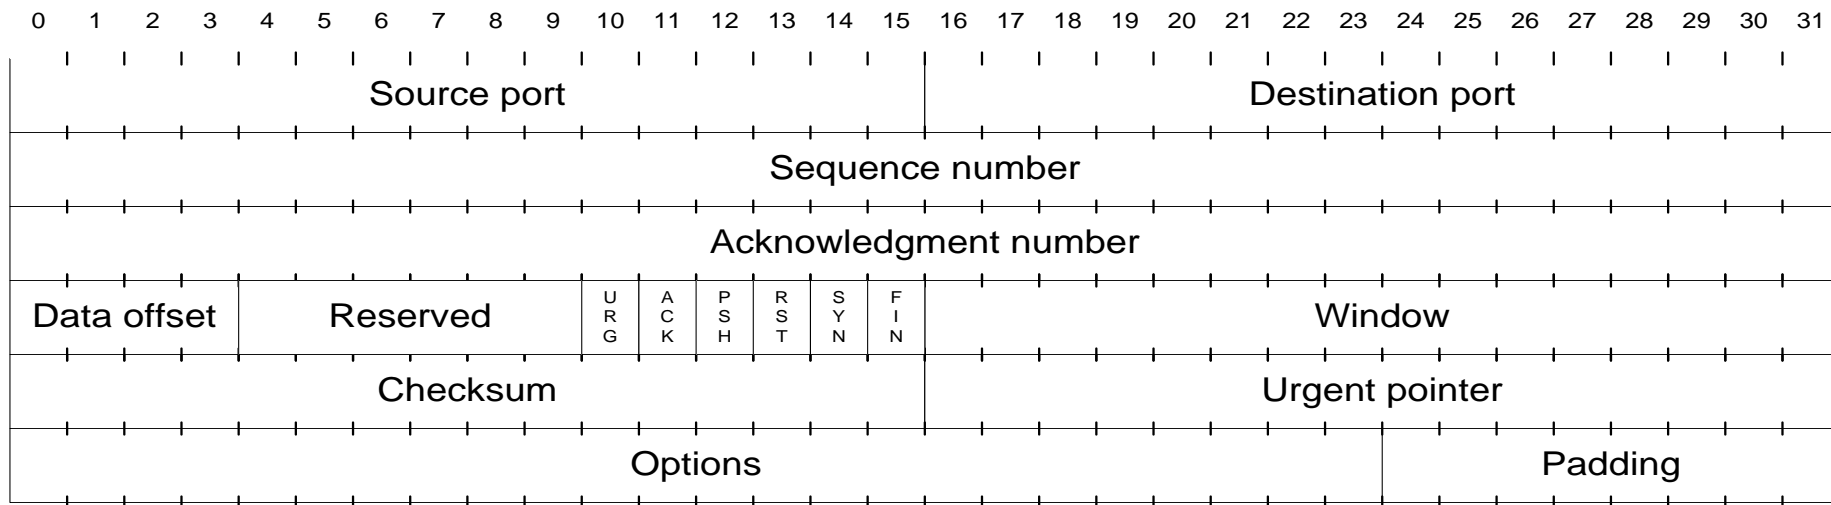
- Control fields
 - **ACK: 1** implies that the value of the acknowledgment field is meaningful
 - **SYN: 1** implies that the segment is trying to synchronize sequence numbers
 - **RST:** Reset the connection



Note: Each tick mark represents a bit position

TCP Header Fields

- Window size: Number of data octets the sender of this information is willing to accept beyond the ACK number field
- Checksum: For error detection
- Urgent pointer: Indicates that data must be processed immediately (earlier applications).

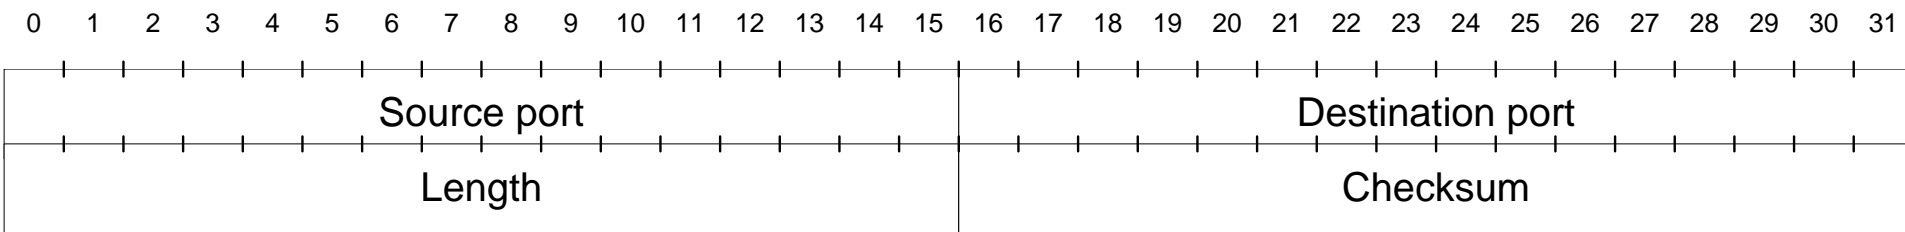


Note: Each tick mark represents a bit position

User Datagram Protocol (UDP)

- Defined in RFC 768 (1980)
- Many applications do not need TCP, such as
 - Applications sending very small amounts of data
 - Prefer speed to reliability (voice)
- In these cases, if we can avoid TCP, we **eliminate the overhead** of keeping track of sequence numbers, window sizes etc.
- With UDP: **No retransmissions, no sequence number validation, no window size defined.**

UDP Header



Note: Each tick mark represents a bit position

Checksum: similar to TCP, but often not used in computer-intensive applications such as video streaming

Length: of the datagram

Transport layer protocols

TCP service:

- Connection-oriented: setup required between client and server processes (3-way handshake)
- Reliable transport between sending and receiving process
- Flow control: sender won't overwhelm receiver

UDP service:

- **Unreliable data transfer** between sending and receiving process
- Faster transmission
- **Does not** provide: connection setup, reliability, flow control

Summary

- Segmentation
 - Need sequence numbers
- Flow control by sliding window
- Port numbers for application multiplexing
- TCP connection establishment through three-way handshake to agree on the initial sequence numbers
- UDP vs TCP