



University of Pittsburgh

# ECE 2195: Special Topics – Computers Machine Learning

## Logistic Regression

**Mai Abdelhakim, PhD**

ECE Department

Swanson School of Engineering

University of Pittsburgh

[maia@pitt.edu](mailto:maia@pitt.edu)



# Classification

- Response is qualitative – Very common problem
- Examples:
  - Email **spam detection** system: spam or not spam
  - **Medical system**: Set of symptoms attributed **to one of three possible** medical conditions. Which of these conditions does the individual have
  - **Handwritten digit recognition** system: Is the digit 0, 1, 2,..or 9?
- Response Y can be represented by unordered set C

# Can we use linear regression for binary classification?

- Assume two classes: We want to classify whether patient has **stroke** or **drug overdose** based on **symptoms**
- We can potentially use the dummy variable approach by defining

$$Y = \begin{cases} 0 & \text{if stroke;} \\ 1 & \text{if drug overdose.} \end{cases}$$

- If the **predicted**  $Y > 0.5$  (**not meaningful**), then classify as drug overdose, otherwise classify as stroke
  - Linear regression **may** work for two-level response

# Can we still use linear regression for more than 2 classes?

- Assume three-level response (3 classes) - Three possible diagnosis: **stroke**, **drug overdose** and **epileptic seizure**
- Can linear regression be applied if we used the following encoding?

$$Y = \begin{cases} 1 & \text{if stroke;} \\ 2 & \text{if drug overdose;} \\ 3 & \text{if epileptic seizure.} \end{cases} \quad \text{OR} \quad Y = \begin{cases} 1 & \text{if epileptic seizure;} \\ 2 & \text{if stroke;} \\ 3 & \text{if drug overdose.} \end{cases}$$


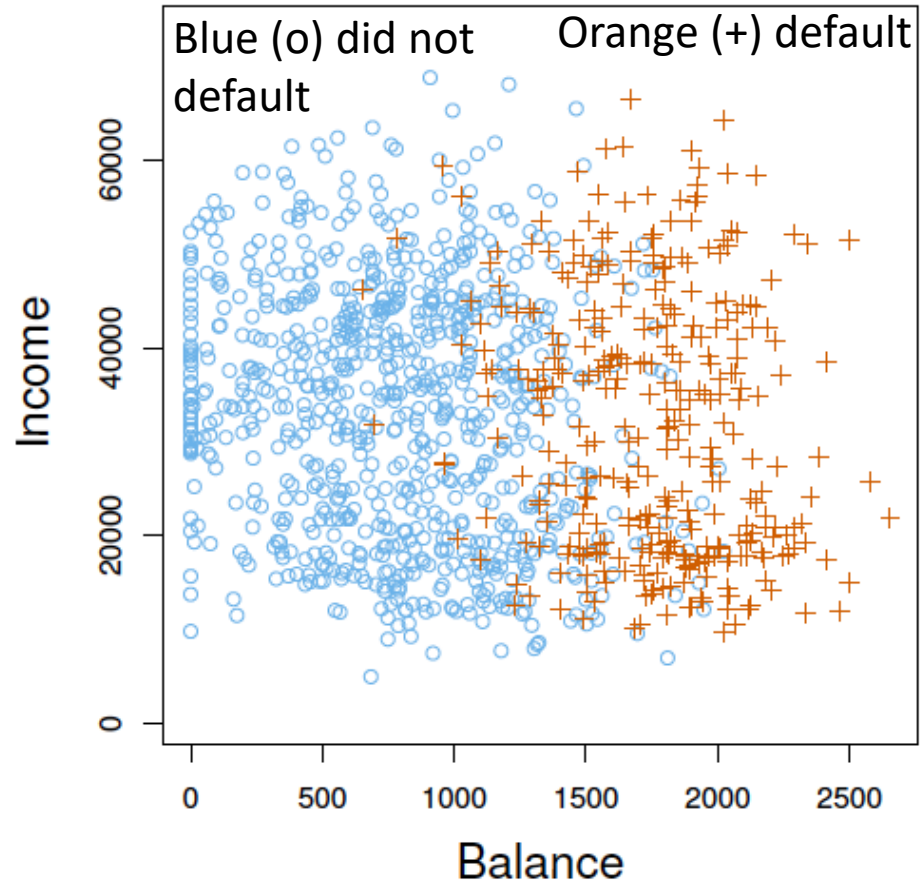
- Different coding result in different linear model and the encoding does not reflect a meaningful interpretation
  - Why error between stroke & seizure is higher than difference between stroke and overdose
- **Linear regression is not appropriate here**
- There is no general way to convert qualitative response into quantitative response

# Logistic Regression


- **A classifier estimates the conditional probability** that  $X$  belongs to each possible class label (elements in  $C$ )
  - Recall optimal Bayes classifier
  - What is the probability that the email is spam (or not) given features of email
- **The predicted class label is the one with the highest probability**
- **Logistic regression: models the probability** that the response  $Y$  belongs to a particular category

# Example: Default Data Set

- Predict whether an individual will default on his/her credit card payment
  - Two features: balance on the credit card and income
- Dataset contains information of  $n=10,000$  individuals



Box plots



# Logistic regression models the conditional probability

- Consider that we take a single feature: balance
- Logistic regression models the probability of default given the balance:

$$\Pr(\text{default} = \text{Yes} | \text{balance})$$

- Range?
  - This conditional probability is in the range [0,1]

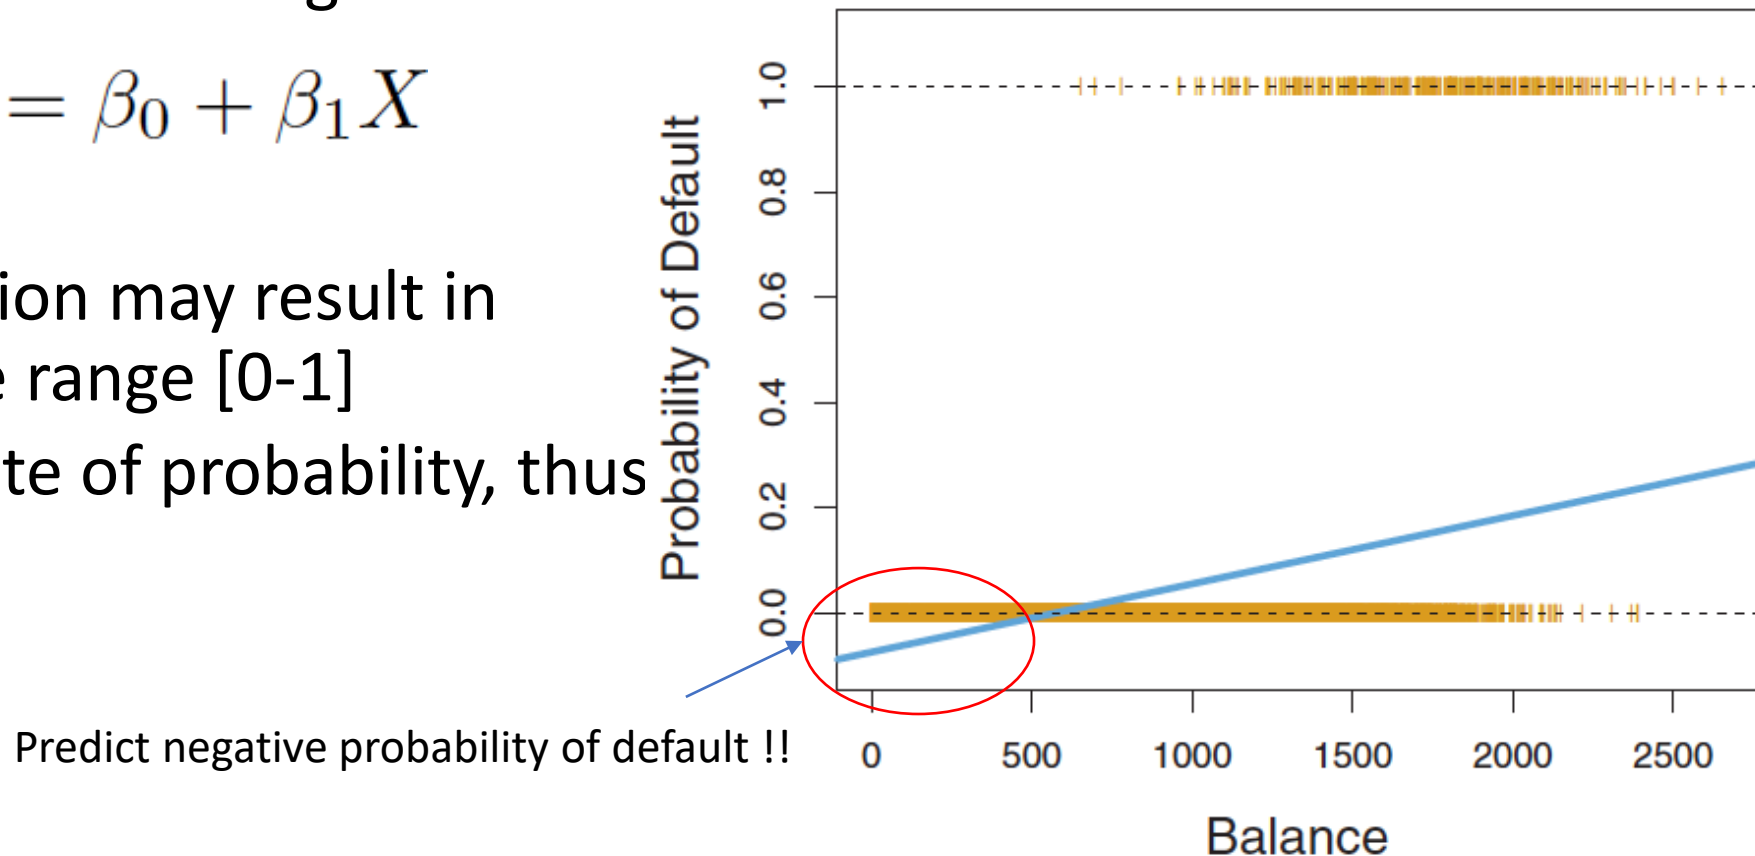


# Linear Model for Probability?

- Assume two possible classes, and denote  $f(X) = \Pr(Y = 1|X)$
- Now we can assume that this probability is the response we want to predict
- Can we use the linear regression model?

$$f(X) = \beta_0 + \beta_1 X$$

- Linear regression may result in values outside range [0-1]
  - Not estimate of probability, thus is not used






# Logistic Model

Logistic – Sigmoid function  $f(a)=1/(1+\exp(-a))$

- To solve this problem the logistic regression uses the form ( $e = 2.71828$  is a mathematical constant):

$$f(X) = \frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}}$$

- $f(X)$  is always in the range  $[0,1]$
- Large exp. power ( $\infty$ ): logistic regression output 1
- Small exp. power ( $-\infty$ ): logistic regression output 0



# Logit term is linear - hence the term logistic regression

- We can rearrange the terms, and get logit  $\log(f(x)/(1 - f(X)))$ :

$$\frac{f(X)}{1 - f(X)} = e^{\beta_0 + \beta_1 X}$$

$$\log \left( \frac{f(X)}{1 - f(X)} \right) = \beta_0 + \beta_1 X$$

This log is base e, i.e. ln

Logistic regression has a **logit** that is **linear** with X

What is the decision boundary?

# Estimating Coefficients – Maximum Likelihood

- It is more common to use maximum likelihood to estimate coefficients
- Likelihood gives probability of the observed zeros and ones in the training data.

$$\ell(\beta_0, \beta_1) = \prod_{i: y_i=1} f(x_i) \prod_{i': y_{i'}=0} (1 - f(x_{i'}))$$

Note that  $f(x_i)$  is the probability  $\Pr(Y_i = 1|x_i)$ ..

Therefore,  $1 - f(x_i) = \Pr(Y_i = 0|x_i)$ .

- Get coefficients that maximizes the likelihood, then use them for predictions

# Note

- Maximizing the likelihood function is equivalent to minimizing the cost function  $J(\beta)$  defined as

$$J(\beta) = - \sum_{i=1}^n [y_i \log(P(y_i = 1|x)) + (1 - y_i) \log(1 - P(y_i = 1|x))]$$

Called cross-entropy function

proof!

Non-linear optimization problem. Can be solved iteratively (for e.g. stochastic gradient descent)

No closed form solution due to the non-linear sigmoid function

# Coefficient statistics

- Similar aspects to linear regression: accuracy of coefficient estimate and p-value
- Here we have Z-statistics (instead of t-statistics as it have normal distribution):

$$\hat{\beta}_1 / SE(\hat{\beta}_1)$$

Credit card Default data using Balance

	Coefficient	Std. Error	Z-statistic	P-value
Intercept	-10.6513	0.3612	-29.5	< 0.0001
balance	0.0055	0.0002	24.9	< 0.0001

What is the probability of default for someone with balance \$1000 ?

# Predictions and estimating coefficients

- After estimating coefficients, we can make predictions.
- Example: estimated probability of default for someone with balance \$1000 is

$$\hat{f}(X) = \frac{e^{\hat{\beta}_0 + \hat{\beta}_1 X}}{1 + e^{\hat{\beta}_0 + \hat{\beta}_1 X}} = \frac{e^{-10.6513 + 0.0055 \times 1000}}{1 + e^{-10.6513 + 0.0055 \times 1000}} = 0.006$$

Coefficient values are in the previous table

# Multiple Features

$Y$  is 1 or zero, and is predicted using multiple features

Denote  $p(X)$  as the probability  $\Pr(Y = 1|X)$

$$\Pr(Y = 1|X) = f(X) = \frac{e^{\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p}}{1 + e^{\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p}}$$

$$\log \left( \frac{f(X)}{1 - f(X)} \right) = \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p$$



# Example

	Coefficient	Std. Error	Z-statistic	P-value
Intercept	-10.8690	0.4923	-22.08	< 0.0001
balance	0.0057	0.0002	24.74	< 0.0001
income	0.0030	0.0082	0.37	0.7115
student [Yes]	-0.6468	0.2362	-2.74	0.0062

- **Student** with card **balance** of \$1,500 and an **income** of \$40,000 (units of income in data is in \$1000) has an estimated probability of default:

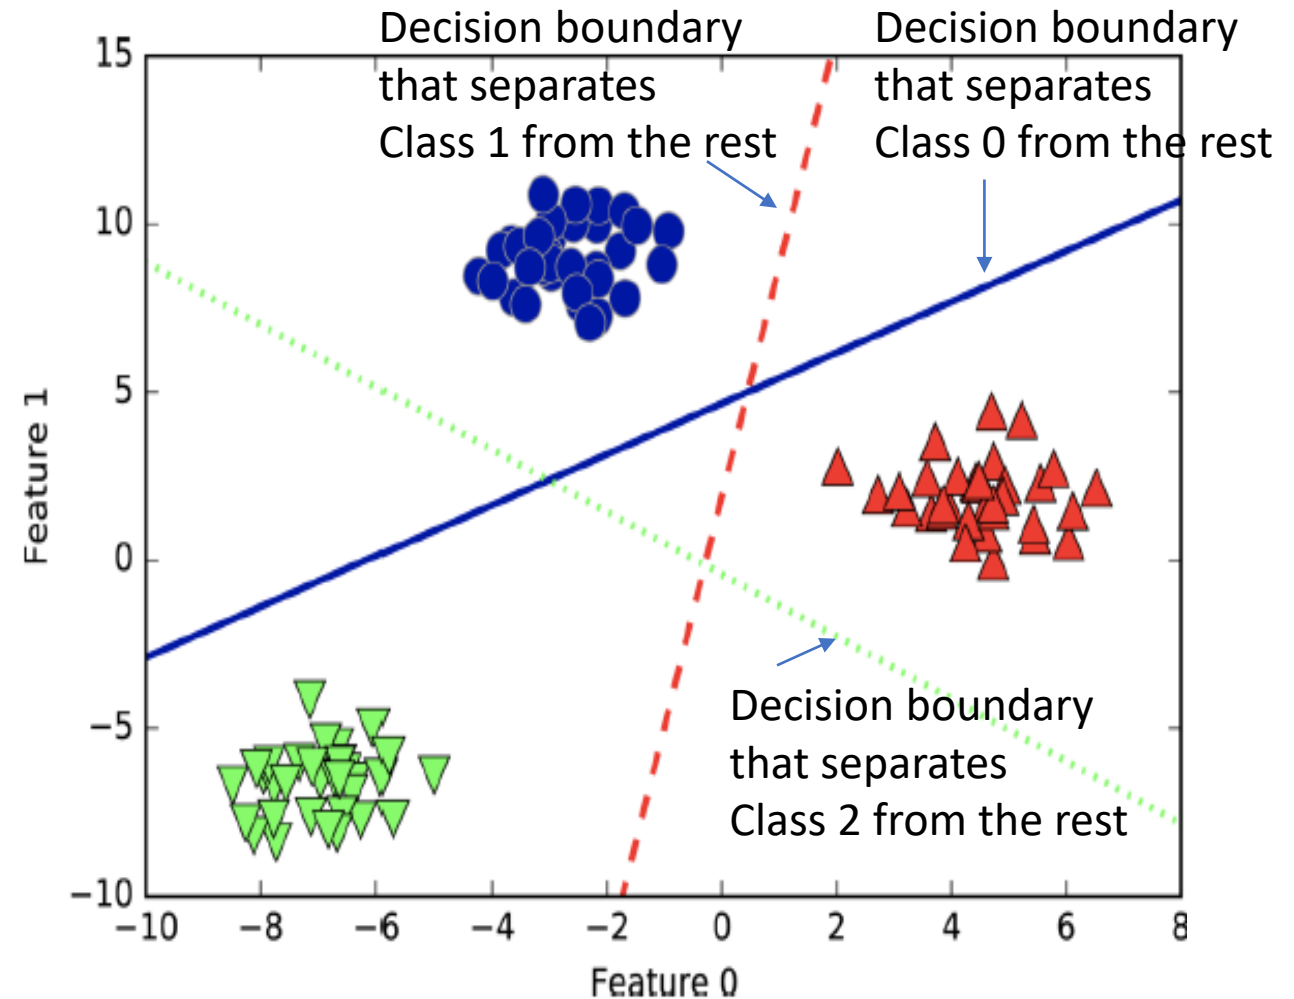

$$\hat{f}(X) = \frac{e^{-10.869+0.00574 \times 1,500+0.003 \times 40-0.6468 \times 1}}{1 + e^{-10.869+0.00574 \times 1,500+0.003 \times 40-0.6468 \times 1}} = 0.058.$$


# Multiple Classes

- One vs all
  - Predict whether response is from: class 1 or not, class 2 or not,...
- Train classifier for each class  $j$  to predict  $y=j$  or not
- Pick class that has  $\max P(y=j | X)$
- We'll talk about other methods that are more popular for multiclass classification

# One vs. All for Multiclass Classification

- Assume two features and three classes shown in the Figure
- One can build three binary classifiers





The region (triangle) where the classification is “rest” by all three binary classifiers: decision is made to assign to the class with the closest boundary

# Logistic Regression in Python

```
From sklearn.linear_model import LogisticRegression
```

```
LogRegModel= LogisticRegression( )
```

- Use .fit and .score as before.

# Regularization in Logistics Regression

- Regularization can be applied to logistic regression
- Same principles as before:
  - Ridge: All coefficients shrink towards zero
  - Lasso: Shrinks coefficients, and some coefficients will be forced to be zero (feature selection)

# Logistic Regression in Python - Ridge

From `sklearn.linear_model` import **LogisticRegression**

LogRegModel= LogisticRegression(**C=100** )



Regularization strength, Ridge  
Large C => less regularization

- By **default**, logistic regression in scikit-learn implements **Ridge regularization** to the classification problem, with strength defined by parameters C.
  - By **default**, **C=1** in LogisticRegression( )
  - **Large C** means **less regularization** strength
    - very large C means is close to the no regularization case
  - **Small C** means **more regularization** and coefficients will be close to zero
  - Note that **C is opposite to alpha or  $\lambda$**  in the regression functions



# Logistic Regression in Python - Lasso

- We can implement Lasso regularization (also called L1 regularization) which limits the model to few features

```
LogRegModel= LogisticRegression(C=0.1, penalty="l1")
```

# Finding Class probabilities in Python

- `predict_proba`: gives the probability of each of the classes given the observation
  - First row is: [probability that first observation is in class 1, probability that the first observation is in class 2,..]
- In python:

```
FittedLogRegModel1= LogisticRegression().fit(X_train,Y_train)  
Probabilities=FittedLogRegModel1.predict_proba(X_test)
```