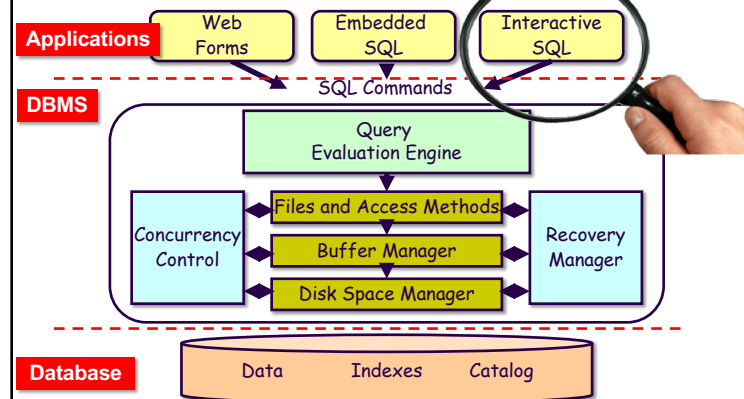


## SQL Views

- ◆ Views
- ◆ Materialized views
- ◆ Temporary tables

## Database Management System (DBMS)

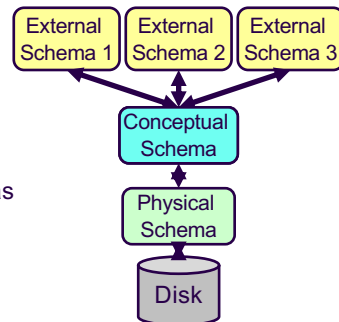


## Levels of Abstraction in a DBMS

- The data in a DBMS is described at three levels of abstraction:

1. *Conceptual Schema*
2. *Physical Schema*
3. *External Schema*

- Many external schemas
- plus one conceptual
- plus one physical



## External Schema

Course Name	Total Enrollment
DB	2
SW	0
OS	1



SID	Name	Age	GPA
546007	Peter	18	3.8
546100	Bob	19	3.65
546500	Bill	20	3.7

CID	CName
CSC 343	DB
CSC 207	SW
CSC 369	OS

SID	CID	Grade
546007	CSC 343	A
546007	CSC 369	B+
546100	CSC 343	B

## Create View

- ❑ A view is a table **derived** from base tables and other views
- ❑ Views can be queried as if they were base tables

```
CREATE VIEW CS_STUDENT
AS SELECT *
FROM STUDENT
WHERE Major = 'CS';
```

```
SELECT Class, Count(*)
FROM CS_STUDENT
GROUP BY Class;
```

## What is a view?

- ❑ It is a table:
  - as it can be queried just like a table!
- ❑ It is not a table:
  - as it does not physically exist!
- ❑ A view is a **“virtual table”** derived from **base tables**
- ❑ A view is a **“named query”**



## Advantages of Views

1. Logical independence
2. For **convenience** and clarity when writing queries
  - Views can be used just like tables
3. For **security**
  - Different data **access privileges** can be given to different users (i.e., **authorization**)

## Query Rewriting

### View

```
CREATE VIEW CS_Students AS
SELECT name, age
FROM Student
WHERE Major = 'CS';
```



### Original Query (user)

```
SELECT name
FROM CS_Students
where age > 19;
```



### Modified Query (DBMS)

```
SELECT name
FROM Student
WHERE Major = 'CS'
AND age > 19;
```

## Modify & Drop a View

### ❑ Modify a view

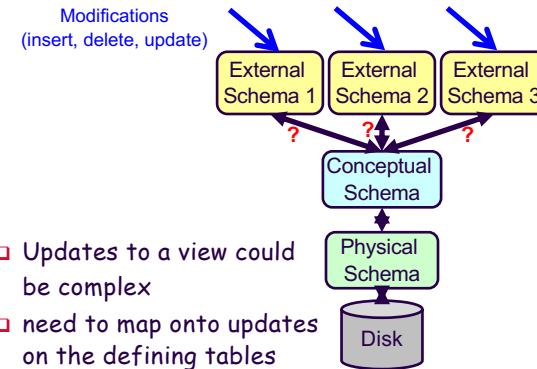
```
CREATE OR REPLACE VIEW CS_STUDENT (Class, Num)
AS SELECT Class, COUNT(*)
FROM STUDENT
WHERE Major = 'CS'
GROUP BY Class;
```

### ❑ Note:

- The new query must generate the same schema
- No REPLACE only CREATE MATERIALIZED VIEW

### ❑ Dropping a view: DROP VIEW CS\_STUDENT;

## View Updatability



- ❑ Updates to a view could be complex
- ❑ need to map onto updates on the defining tables

## Example 1 [Poll]

### Student

SID	Name	Age	GPA	Major
546007	Peter	18	3.8	CS
546100	Bob	19	3.6	EE
546500	Peter	20	3.0	CS

Updating a view:

```
UPDATE HStudents
SET gpa= 4.0
WHERE sid=546007;
```



```
CREATE VIEW HStudents AS
SELECT sid, gpa
FROM Student
WHERE gpa> 3.5;
```

## Example 2 [Poll]

### Student

SID	Name	Age	GPA	Major
546007	Peter	18	3.8	CS
546100	Bob	19	3.6	EE
546500	Peter	20	3.0	CS

Updating a view:

```
INSERT INTO
HStudents
VALUES (Jane, 4.0);
```



```
CREATE VIEW HStudents AS
SELECT name, gpa
FROM Student
WHERE gpa> 3.0;
```

What is Jane's SID?!

### Example 3 [Poll]

#### Student

SID	Name	Age	GPA	Major
546007	Peter	18	3.8	CS
546100	Bob	19	3.6	EE
546500	Peter	20	3.0	CS

Updating a view:

```
UPDATE HStudents
SET gpa= 4.0
WHERE name='Peter';
```



```
CREATE VIEW HStudents AS
SELECT name, gpa
FROM Student
WHERE gpa > 3.0;
```

How many updates  
will succeed?

### Example 4 [Poll]

#### Student

SID	Name	Age	GPA	Major
546007	Peter	18	3.8	CS
546100	Bob	19	3.6	EE
546500	Peter	20	3.0	CS

Updating a view:

```
UPDATE Majors
SET agpa= 3.6
WHERE major='CS';
```



```
CREATE VIEW Majors (major, agpa) AS
SELECT major, avg(gpa)
FROM Student
GROUP BY Major;
```

Infinite possibilities  
of values!

### View Updateability



- ❑ In general, a view is called **updateable** if:
  - all updates on the view can be **unambiguously** translated back to tuples in the base tables
- ❑ A view update is **unambiguous** if:
  - Only one** update on the base tables can accomplish the desired update effect on the view
- ❑ In general, a view is **not updateable** if:
  - an update on a view can be mapped to **more** than one possible update on the base tables

### SQL Standard for View Updateability

1. A view with a single defining table is updateable if the view attributes contain the **primary key**
2. Views defined using **aggregate** functions are not updateable
3. Views defined on **multiple tables** using joins are generally not updateable

## Updating a View

```
CREATE VIEW CS_STUDENT
AS SELECT *
FROM STUDENT
WHERE Major = 'CS';
```

```
INSERT INTO CS_STUDENT (128, 'Ping Chen', 'CS');
UPDATE CS_STUDENT
SET Name = 'Shimin Chen'
WHERE SID = 128;
```

## Migrating Tuples

- What is the outcome of the update:

```
UPDATE CS_STUDENT
SET Major = 'MATH'
WHERE SID = 128;
```



- Migrating tuples out of updateable views: an update or insert may eliminate a tuple from the view
- Prevent migration with WITH CHECK OPTION

```
CREATE VIEW CS_STUDENT
AS SELECT * FROM STUDENT WHERE Major = 'CS'
WITH CHECK OPTION;
```

## Efficient View Implementation

- A DBMS implements views in two ways:
  - ✓ Query Rewriting / Modification
  - ✓ View Materialization
- With expected trade-offs...

## Query Rewriting

- Query rewriting:
  - presents the view query in terms of a query on the underlying base tables
- Disadvantage:
  - re-compute** the view with every query
    - E.g., multiple queries `SELECT name FROM IT_Students` where age > 19, 20, 21, ...
  - inefficient for views defined via complex queries (e.g., aggregate queries)

## Virtual vs. Materialized Views

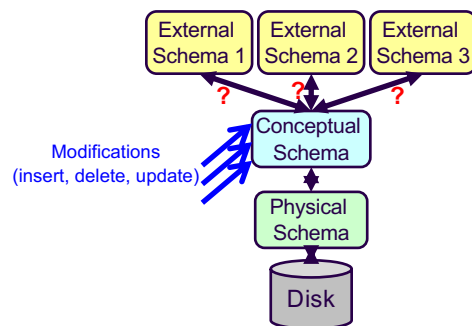
- Views:
  - Virtual tables
  - Evaluating a view (query) creates its data
- Materialized Views:
  - **Stored tables**
  - Physically store the view (query) and its data

## Materialized Views

- Advantage:
  - Avoid re-computing the view with every query
  - Assumption: more queries can use the same view
- **But**, materialized **view maintenance** is needed
  - A materialized view should be updated when any **base table** used in the view definition is updated



## Updating Materialized Views



## Example of View Materialization

### STUDENT

<i>SID</i>	<i>Name</i>	<i>Age</i>	<i>GPA</i>	<i>Major</i>
546007	Peter	18	3.8	CS
546100	Bob	19	3.65	CoE
546500	Bill	20	3.7	CS

Update on base table:

```
INSERT INTO Student
VALUES (456, ..., CoE);
COMMIT;
```

```
CREATE MATERIALIZED VIEW
Majors (major, mtotal)
AS SELECT major, count(*)
FROM Student
GROUP BY Major;
```

### Materialized View

<i>major</i>	<i>mtotal</i>
CS	2
CoE	2

## Updating Materialized Views

- ❑ Efficient strategies for automatically updating the materialized view when base tables are updated
  - Avoid re-computing the view from “scratch”
  - **Incremental update:**
    - determines what new tuples must be inserted, deleted, or modified in the view when an update is applied to the base tables

## Trade-offs in view implementation



	(Virtual) Views	Materialized Views
Queries on Views		
Updates on Base Tables		

## Trade-offs in view implementation

	(Virtual) Views	Materialized Views
Queries on Views	Re-compute view	Re-use view
Updates on Base Tables	Do nothing	View Maintenance

## Materialized View (ANSI/Oracle)

```
CREATE MATERIALIZED VIEW Majors (major, mtotal)
[BUILD METHOD][REFRESH OPTION METHOD]
AS SELECT major, count(*)
FROM Student
GROUP BY Major;
```

- ❑ No **REPLACE** *only* **CREATE MATERIALIZED VIEW**
- ❑ Build Method:
  - **IMMEDIATE:** Create view and populate it with data
  - **DEFERRED:** Create view but do not populate it
- ❑ Refresh Method:
  - **ON DEMAND:** Manually - DBMS\_MVIEW.REFRESH('<MV-name>');
  - **ON COMMIT:** Automatic after a commit
- ❑ Refresh Option:
  - **COMPLETE** (re-computation), **FAST** (incremental), **NEVER**

## Materialized View in PostgreSQL

```
CREATE MATERIALIZED VIEW [IF NOT EXISTS]
Majors (major, mtotal)
AS SELECT major, count(*)
FROM Student
GROUP BY Major
WITH [NO] DATA;
```

- ❑ No REPLACE *only* CREATE MATERIALIZED VIEW
- ❑ Build Method:
  - WITH DATA: Create view and populate it with data
  - WITH NO DATA: Create view but do not populate it
- ❑ Refresh Method:
  - ON DEMAND: Manually - REFRESH MATERIALIZED VIEW '<MV-name>';
  - ON COMMIT (with Trigger): Automatic after a commit

## Full Materialized View Construction

(ANSI/Oracle)

```
CREATE MATERIALIZED VIEW Majors (major, mtotal)
[WITH ENCRYPTION, SCHEMABINDING, VIEW_METADATA]
[BUILD METHOD][REFRESH OPTION METHOD]
AS SELECT major, count(*)
FROM Student
GROUP BY Major
WITH CHECK OPTION;
```

- ❑ ENCRYPTION: The definition of the view is stored encrypted
- ❑ SCHEMABINDING: Prevents the drop of defining tables/views
- ❑ VIEW\_METADATA: It makes visible the metadata on the view but hides the metadata of the defining tables/views.
- ❑ WITH CHECK OPTION: Prevent migration of tuples out of updateable views

## Views Vs Temporary Tables

- ❑ No standard but Temporary Tables are
  - visible to the current SQL session
  - automatically dropped at the end of session
  - cannot have foreign key constraints
- ❑ Temporary tables are local (Postgres, Oracle, MySQL)
  - Create Temporary Table Yahoo (YID int, YNM Char(3));
  - SQLServer: Create Table #Yahoo (YID int, YNM Char(3));
- ❑ Postgres & Oracle Server: global temporary tables
  - Example: Create Global Temporary Table Yahoo  
on commit preserve rows  
AS Select YID, YNM From TahoeBase;