

Recitation 13

B+trees, Linear Hashing, and Extendible Hashing

CS1555/2055, Panos K. Chrysanthis & Constantinos Costa – University of Pittsburgh

1

1

B⁺-tree Index

- A node is of the form:

$[p_0, k_1, p_1, k_2, p_2, \dots, k_i, p_i, k_{i+1}, \dots, k_n, p_n]$

- p_i 's are pointers and k_i 's are field values (keys)
- Tree Order** is the number of pointers, e.g., n
- For every field value k in a node pointed to by p_i
 $k_i < k \leq k_{i+1}$ (alternative $k_i \leq k < k_{i+1}$)
- Every node, except for the root, has between $n/2$ and n children or pointers
 - internal: $\lceil n/2 \rceil$ **tree pointers**; leaf: $\lfloor n/2 \rfloor$ **data pointers**
- Leaf nodes are chain to form a link list (**fast sequential access**)

CS1555/2055, Panos K. Chrysanthis & Constantinos Costa – University of Pittsburgh


2

2

Examples of B⁺ trees


B⁺-tree with $n=3$

- Internal nodes $\lceil 3/2 \rceil$
= 2 **tree pointers**
- Leaf nodes $\lfloor 3/2 \rfloor$
= 1 **data pointers**



B⁺-tree with $n=5$

- Internal nodes $\lceil 5/2 \rceil$
= 3 **tree pointers**
- Leaf nodes $\lfloor 5/2 \rfloor$
= 2 **data pointers**

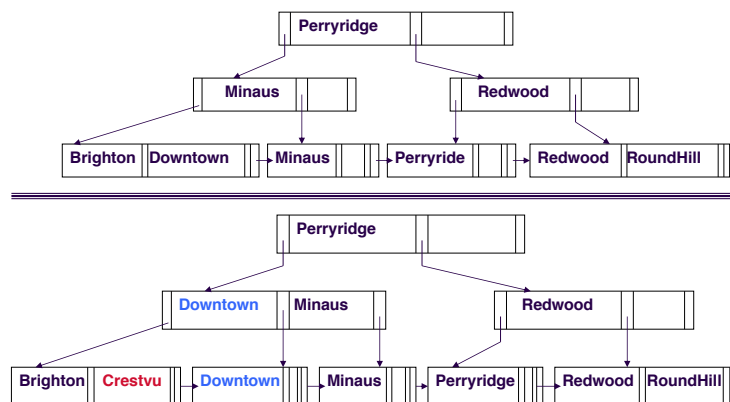


CS1555/2055, Panos K. Chrysanthis & Constantinos Costa – University of Pittsburgh

3

3

Insertion of "Crestvu"

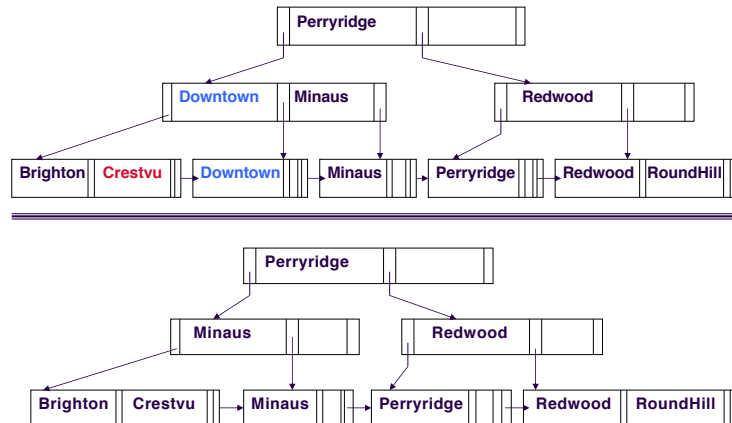


CS1555/2055, Panos K. Chrysanthis & Constantinos Costa – University of Pittsburgh

4

4

Deletion of "Downtown"




CS1555/2055, Panos K. Chrysanthis & Constantinos Costa – University of Pittsburgh

5

5

Deletion of "Perryridge"



CS1555/2055, Panos K. Chrysanthis & Constantinos Costa – University of Pittsburgh

6


6

Insert Examples of B+ trees

B⁺-tree with $n=4$, i.e.,


- internal nodes should have $\lceil 4/2 \rceil = 2$ tree pointers;
- Leaf nodes should have $\lfloor 4/2 \rfloor = 2$ data pointers

Step 1: insert 1, 2, 4



Step 2: insert 9

Although you push 9 on the stack & inserted afterwards when allocating a new block, it is considered when deciding the split: 1 2 4. So the split is at 1 2 | 4 9 and copy up 4



CS1555/2055, Panos K. Chrysanthis & Constantinos Costa – University of Pittsburgh

8

8


CS1555/2055, Panos K. Chrysanthis & Constantinos Costa – University of Pittsburgh

7

7


Insert Examples of B+ trees

Step 3: insert 3, 11



Step 4: insert 7

Although you push 7 on the stack when allocating a new block, it is considered when deciding the split: 4 9 11. So the split is at 4 7 | 9 11




CS1555/2055, Panos K. Chrysanthos & Constantinos Costa – University of Pittsburgh

9

9


Insert Examples of B+ trees

Step 5: insert 8



Step 6: insert 5

Although you push 5 on the stack when allocating a new block, it is considered when deciding the split: 4 7 8. So the split is at 4 5 | 7 8




CS1555/2055, Panos K. Chrysanthos & Constantinos Costa – University of Pittsburgh

10

10


Insert Examples of B+ trees

Step 7: insert 12



Step 8: insert 13

Push 13 on the stack when allocating a new block, it is considered when deciding the split: 9 11 12. So the split is at 9 11 | 12 13. then copy up 12.




CS1555/2055, Panos K. Chrysanthos & Constantinos Costa – University of Pittsburgh

11

11

Insert Examples of B+ trees

Step 8: insert 13




CS1555/2055, Panos K. Chrysanthos & Constantinos Costa – University of Pittsburgh

12


12

Delete Examples of B+ trees

Step 9: delete 11



Merge



CS1555/2055, Panos K. Chrysanthis & Constantinos Costa – University of Pittsburgh



13

13

Examples of B+ trees

Step 9: delete 11 (cont)

Borrow from the left subtree




CS1555/2055, Panos K. Chrysanthis & Constantinos Costa – University of Pittsburgh

14


14

Examples of B+ trees

Step 10: delete 9



Step 11: delete 12



CS1555/2055, Panos K. Chrysanthis & Constantinos Costa – University of Pittsburgh


15

15

Examples of B+ trees



Step 11: delete 12 (cont)

Merge



We can't borrow a node from the left subtree

Merge the root



CS1555/2055, Panos K. Chrysanthis & Constantinos Costa – University of Pittsburgh

16

16

17

Linear Hashing

- The file size grows linearly, bucket by bucket, using a family of hash functions, each has double the range of its predecessor
 - The file starts with s main buckets, where s is a power of 2, and k overflow buckets.
 - Buckets are numbered from 0 to $s-1$
 \Rightarrow initial hashing function $h_0(\text{key}) = \text{key} \bmod s$.
 - Collisions are handled thru **chaining**.
- We keep the following info:
 - B_{last} : A pointer to the **current** last bucket.
Initially, $B_{last} = s-1$.
 - B_{split} : A pointer to the bucket that should be split next.
Initially, $B_{split} = 0$

18

Insertion of a Record (Method 1)

- ❑ if there is a collision, push tuple to an overflow bucket.
- ❑ if there are no overflow buckets available, proceed as follows until one becomes available:
 - A new bucket is appended at the end of the hash table and $B_{last} = B_{last} + 1$
 - Records in B_{split} bucket are hashed again using $h_1(\text{key}) = \text{key} \bmod (2s)$,
 \Rightarrow these will either remain in B_{split} or stored in B_{last}
 \Rightarrow this may free an overflow bucket.
 - B_{split} becomes $B_{split} + 1$.
- ❑ if $B_{last} = 2s-1$
 set $s = 2s$, $B_{last} = s-1$, $B_{split} = 0$, $h_0(\text{key}) = h_1(\text{key})$
- ❑ proceed as above.


19

Insertion of a Record (Method 2)

- ❑ if there is a collision, push tuple to an overflow bucket.
- ❑ if there are no overflow buckets available, proceed as follows until one becomes available:
 - A new bucket is appended at the end of the hash table (but B_{last} does not change as in Method 1)
 - Records in B_{split} bucket are hashed again using $h_1(\text{key}) = \text{key} \bmod (2s)$,
 \Rightarrow these will either remain in B_{split} or stored in B_{last}
 \Rightarrow this may free an overflow bucket.
 - B_{split} becomes $B_{split} + 1$.
- ❑ if $B_{last} < B_{split}$
 set $s = 2s$, $B_{last} = s-1$, $B_{split} = 0$, $h_0(\text{key}) = h_1(\text{key})$
- ❑ proceed as above.

20

Example




| 0 | 1 | 2 | 3 |
|---|---|---|---|
| 4 | 1 | 6 | 3 |
| 0 | 5 | | 7 |
| 8 | | | |

$B_{split}=0$ $B_{last}=3$

$$h_0(k)=k \bmod 2^2$$

1, 4, 3, 0, 6, 8, 5, 7, 16 ?



| 0 | 1 | 2 | 3 | 4 |
|---|---|---|---|---|
| 0 | 1 | 6 | 3 | 4 |
| 8 | 5 | | 7 | |
| | | | | |


$B_{split}=1$ $B_{last}=3$

$$h_0(k)=k \bmod 2^2$$

$$h_1(k)=k \bmod 2^3$$

21

Example



| 0 | 1 | 2 | 3 | 4 |
|----|---|---|---|----|
| 0 | 1 | 6 | 3 | 4 |
| 8 | 5 | | 7 | 12 |
| 16 | 9 | | | |

$B_{split}=1$ $B_{last}=3$

$$h_0(k)=k \bmod 2^2$$


$$h_1(k)=k \bmod 2^3$$

1, 4, 3, 0, 6, 8, 5, 7, 16, 12, 9, 13 ?

Is $B_{split} > B_{last}$?

22

Example



| 0 | 1 | 2 | 3 | 4 | 5 |
|----|---|---|----|----|----|
| 0 | 1 | 6 | 3 | 4 | 5 |
| 8 | 9 | | 7 | 12 | 13 |
| 16 | | | 11 | | |

$B_{split}=2$ $B_{last}=3$

$$h_0(k)=k \bmod 2^2$$


$$h_1(k)=k \bmod 2^3$$

1, 4, 3, 0, 6, 8, 5, 7, 16, 12, 9, 13, 11, 15 ?

Is $B_{split} > B_{last}$?

23

Example



| 0 | 1 | 2 | 3 | 4 | 5 | 6 |
|----|---|---|----|----|----|---|
| 0 | 1 | | 3 | 4 | 5 | 6 |
| 8 | 9 | | 7 | 12 | 13 | |
| 16 | | | 11 | | | |

$B_{split}=3$ $B_{last}=3$

$$h_0(k)=k \bmod 2^2$$


$$h_1(k)=k \bmod 2^3$$

1, 4, 3, 0, 6, 8, 5, 7, 16, 12, 9, 13, 11, 15 ?

Is $B_{split} > B_{last}$?

24

Example



| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |
|----|---|---|----|----|----|----|---|
| 0 | 1 | | 3 | 4 | 5 | 6 | 7 |
| 8 | 9 | | 11 | 12 | 13 | 15 | |
| 16 | | | | | | | |

$B_{split} = 0$

$B_{last} = 7$

~~$h_0(k) = k \bmod 2^2$~~

$h_0(k) = k \bmod 2^3$

$h_1(k) = k \bmod 2^3$

$h_1(k) = k \bmod 2^4$

1, 4, 3, 0, 6, 8, 5, 7, 16, 12, 9, 13, 11, 15

25

Search for a record

➤ Search for a Record

- $I = h_0(\text{key})$
- if $I < B_{split}$ then $I = h_1(\text{key})$
- search the bucket whose hash value is I (and its overflow, if any).

26

Dynamic Hashing Methods


- ❑ Allow the file size to change as records are added or deleted.
 - Linear Hashing
 - No additional structure
 - **Extendible Hashing**
 - Binary Hashing

27

28

Extendible Hashing

- The file is structured into two levels:
directory and buckets.



Extendible Hashing

- Directory contains only pointers to buckets (no keys)
- Use some hash function to generate a **pseudokey** of b -bits (typically $b=32$)
- Use the first (or last) d bits to find the offset of the bucket pointer in the directory.
- d is called the (global) **directory depth**. It may be stored in the directory
- In each bucket a local depth d' is stored indicating the most (or least) significant bits common to all keys in that bucket ($d' \leq d$)


Extendible Hashing - File Growth

- No overflow buckets. New buckets are added one by one.
- When a bucket B with local depth d' overflows, B is split into two buckets and all keys are rehash using $d'+1$ bits
- If after a split, $d' > d$, double the size of the directory ($d=d+1$) to accommodate the growth

Example

- Use the least significant bits as bucket keys
- Each bucket can hold three records
- The keys to insert are 1,4,3,0,6,8,5,7,16,12,9,13


Example



1, 4, 3, 0, 6, 8/1000

33


Example



1, 4, 3, 0, 6, 8/1000

34


Example



1, 4, 3, 0, 6, 8, 5, 7/0111

36

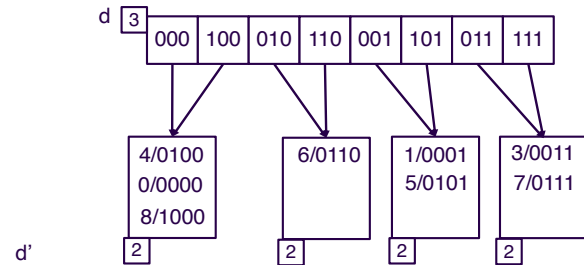
Example



1, 4, 3, 0, 6, 8, 5, 7, 16/10000

37


Example



1, 4, 3, 0, 6, 8, 5, 7, 16/10000

38


Example



1, 4, 3, 0, 6, 8, 5, 7, 16, 12, 9, 13/1101

39

Example



1, 4, 3, 0, 6, 8, 5, 7, 16, 12, 9, 13

40