
PyTorch **Tutorial**

Ahmad Diab
ahd23@pitt.edu
Mingda Zhang

Roadmap

- PyTorch - general intro
- PyTorch vs Numpy
- PyTorch vs TensorFlow
- Power of GPU (cuda)
- Automatic Differentiation
- Neural Network Models
 - Linear Regression
 - Multi-Layer Perceptrons (MLP)
 - Convolutional Neural Nets (CNN)
 - Off-the-shelf models
- TensorBoards

General Intro

- What is PyTorch?
 - Free, open-source Machine Learning library.
 - Applications in Computer Vision and Natural Language Processing.
 - Development led by Facebook AI Research.
- What makes it special?
 - Array computations with GPU compatibility.
 - Deep Learning Networks with Automatic Differentiation.
- Quick facts:
 - Python and C++ interfaces.
 - Core of many DL softwares such as Autopilot (Tesla), Pyro (Uber), Transformers (HuggingFace).
 - It is called Torch (not PyTorch) when used/imported, was born from Lua Torch package.

PyTorch vs. Numpy

- PyTorch introduced a new class called Tensor.
- Tensor is an multi dimensional array.
- Numpy also uses multi dimensional arrays supported by fast-implemented methods.
- Where is the catch?
 - Tensors can run on GPUs, utilizing powerful parallel processing.
 - Numpy arrays can only use CPUs.
- Good news! They are interchangeable.

PyTorch is not alone!



PyTorch vs. TensorFlow

- Popularity: PyTorch is the most common library in academics (based on paper implementations), where TensorFlow is still dominating industry (based on job listing). (source: <https://www.stateof.ai/>)
- Easy to learn: Subjective.
- Computation Graphs: PyTorch utilizes a dynamic perspective where TensorFlow uses a static view.
- Debugging: Harder for TensorFlow, needs external tools.

GPU and cuda

- Graphics Processing Unit (GPU) performs operations in parallel/fast way.
- CPUs use MIMD, GPUs use SIMD
 - $A = B * C$ (think about types and sizes)
- How to use it?
- CUDA = API.
- A software layer that gives access to GPU virtual instruction set.
- Created by NVIDIA, 2007.

```
1 a = torch.tensor(5)
2 b = torch.tensor(5).cuda()
3 a = a ** 2
4 b = b ** 2
5 print(a)
6 print(b)
```

```
tensor(25)
tensor(25, device='cuda:0')
```

Automatic Differentiation

- Simple Example: $X = [1, 2, 3]$
 $Y = [4, 5, 6]$
 $Z = 5X + 2Y$
 $m = \text{sum}(Z) = z_0 + z_1 + z_2$
- Calculate the derivative using the chain rule.

$$\frac{\partial m}{\partial x_i} = \frac{\partial m}{\partial z_i} \cdot \frac{\partial z_i}{\partial x_i} = 1 \times 5 = 5$$

$$\frac{\partial m}{\partial y_i} = \frac{\partial m}{\partial z_i} \cdot \frac{\partial z_i}{\partial y_i} = 1 \times 2 = 2$$

$$\frac{\partial m}{\partial X} = [5, 5, 5]$$

$$\frac{\partial m}{\partial Y} = [2, 2, 2]$$

Automatic Differentiation - more example

- What changes if:
 - $X=[1,2,3,4]$ and $Y=[5,6,7,8]$
 - $Z = 5X^2 + 2Y^3$

$$\frac{\partial m}{\partial x_i} = \frac{\partial m}{\partial z_i} \cdot \frac{\partial z_i}{\partial x_i} =$$

$$\frac{\partial m}{\partial y_i} = \frac{\partial m}{\partial z_i} \cdot \frac{\partial z_i}{\partial y_i} =$$

$$\frac{\partial m}{\partial X} =$$

$$\frac{\partial m}{\partial Y} =$$

Automatic Differentiation - PyTorch

```
4 x = torch.tensor([1., 2., 3.], requires_grad=True)
5 y = torch.tensor([4., 5., 6.], requires_grad=True)
6 z = 5 * x + 2 * y
7
8 m = z.sum()
9 m.backward()
10 print(x.grad, '\n', y.grad)
```

```
tensor([5., 5., 5.])
tensor([2., 2., 2.])
```

```
4 x = torch.tensor([1., 2., 3., 4.], requires_grad=True)
5 y = torch.tensor([5., 6., 7., 8.], requires_grad=True)
6 z = 5 * x + 2 * y
7
8 m = z.sum()
9 m.backward()
10 print(x.grad, '\n', y.grad)
```

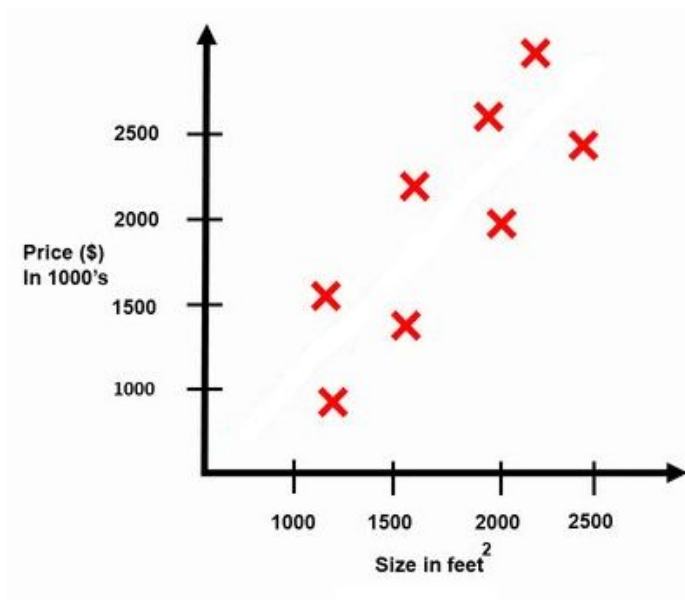
```
tensor([5., 5., 5., 5.])
tensor([2., 2., 2., 2.])
```

```
4 x = torch.tensor([1., 2., 3.], requires_grad=True)
5 y = torch.tensor([4., 5., 6.], requires_grad=True)
6 z = 5 * x ** 2 + 2 * y ** 3
7
8 m = z.sum()
9 m.backward()
10 print(x.grad, '\n', y.grad)
```


```
tensor([10., 20., 30.])
tensor([ 96., 150., 216.])
```

Linear Regression

- Entry-level technique in Machine Learning Realm.
- Example: *House Prices*




Linear Regression - in action



Loss Function

- Loss is critical, it serves as the supervision for training a neural network.
- Different types of loss functions, such as:
 - Mean Squared Error Loss (Regression)
 - Cross-Entropy Loss (Classification)
 - Triplet Loss (Representation Learning)



- All implemented off-the-shelf with unified API, which is easy to use.

Optimizers

- Lower the Loss function by updating model parameters for future iterations based on loss value of current iteration.
- *torch.optim* is a package implementing various optimization algorithms.
- It supports common methods, and provides interface for future integrations.

Example:

```
for input, target in dataset:
    optimizer.zero_grad()
    output = model(input)
    loss = loss_fn(output, target)
    loss.backward()
    optimizer.step()
```

Linear Regression - Module

```
1 import torch.nn as nn
2
3 class ManualLinearRegression(nn.Module):
4     def __init__(self):
5         super().__init__()
6
7         self.weight = nn.Parameter(torch.randn(1, requires_grad=True))
8         self.bias = nn.Parameter(torch.randn(1, requires_grad=True))
9
10    def forward(self, x):
11        return self.weight * x + self.bias
```

```
15 # Build a Model
16 model = ManualLinearRegression()
17 loss_fn = nn.MSELoss()
18 optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
```

Linear Regression - manual implementation

```
1 weight = torch.rand(1, requires_grad=True)  # Assuming  $y = x * w + b$ 
2 bias = torch.rand(1, requires_grad=True)
3
4 learning_rate = 0.1 # Hyperparameter
5
6 total_steps = tqdm_notebook(range(10000))
7
8 for step in total_steps:
9     prediction = weight * data + bias
10    loss = ((prediction - label) ** 2).mean()  # Mean squared error
11    if (step + 1) % 100 == 0:
12        total_steps.set_description("Loss: %.4f" % loss.detach().item())
13
14    loss.backward()
15    # plt.plot(x, weight * x + bias)
16
17    with torch.no_grad():  # Stops tracking variable
18        weight -= learning_rate * weight.grad
19        bias -= learning_rate * bias.grad
20
21    weight.grad.zero_()  # Avoid gradients being accumulated
22    bias.grad.zero_()
```


PyTorch built-in module: *torch.nn*

torch.nn

Parameters

- + Containers
- + Convolution layers
- + Pooling layers
- + Padding layers
- + Non-linear activations (weighted sum, nonlinear)
- + Non-linear activations (other)
- + Normalization layers
- + Recurrent layers
- + Transformer layers
- + Linear layers
- + Dropout layers
- + Sparse layers
- + Distance functions
- + Loss functions
- + Vision layers
- + DataParallel layers (multi-GPU, distributed)
- + Utilities
- Quantized Functions

– Linear layers

Identity

Linear

Bilinear

Linear

CLASS `torch.nn.Linear(in_features, out_features, bias=True)`

[\[SOURCE\]](#)

Applies a linear transformation to the incoming data: $y = xA^T + b$

Parameters

- **in_features** – size of each input sample
- **out_features** – size of each output sample
- **bias** – If set to `False`, the layer will not learn an additive bias. Default: `True`

Shape:

- Input: $(N, *, H_{in})$ where $*$ means any number of additional dimensions and $H_{in} = \text{in_features}$
- Output: $(N, *, H_{out})$ where all but the last dimension are the same shape as the input and $H_{out} = \text{out_features}$.

Variables

- **-Linear.weight** – the learnable weights of the module of shape $(\text{out_features}, \text{in_features})$. The values are initialized from $\mathcal{U}(-\sqrt{k}, \sqrt{k})$, where $k = \frac{1}{\text{in_features}}$
- **-Linear.bias** – the learnable bias of the module of shape (out_features) . If `bias` is `True`, the values are initialized from $\mathcal{U}(-\sqrt{k}, \sqrt{k})$ where $k = \frac{1}{\text{in_features}}$

Examples:

```
>>> m = nn.Linear(20, 30)
>>> input = torch.randn(128, 20)
>>> output = m(input)
>>> print(output.size())
torch.Size([128, 30])
```

MNIST (dataset)

- Dataset of handwritten digits.
- 60k training images, 10k testing images.
- Why?

Good for learning ML/DL techniques on real-world data *while* spending minimal efforts on preprocessing and formatting.

- From *torchvision*; package consists of popular datasets, models, and image transformations for computer vision.



Modularized Networks: MLP and CNN

- All the modules can be stacked together arbitrarily, which makes it convenient to build complicated network architecture using simpler ones.
 - Here the Linear module is stacked to build multi-layer perceptrons, together with ReLU (a non-linear activation module).

```
1 # Define our model (3-layer MLP)
2 class MultilayerPerceptron(nn.Module):
3     def __init__(self, input_size, hidden_size, num_classes):
4         super().__init__()
5         self.fc1 = nn.Linear(input_size, hidden_size)
6         self.relu = nn.ReLU()
7         self.fc2 = nn.Linear(hidden_size, num_classes)
8
9     def forward(self, x):
10         out = self.fc1(x.reshape(-1, 28 * 28))
11         out = self.relu(out)
12         out = self.fc2(out)
13         return out
```

Modularized Networks: MLP and CNN

```
1 class ConvNet(nn.Module):
2     def __init__(self, num_classes=10):
3         super().__init__()
4
5         # We can further simplify the code if a data will be sequentially processed
6         # by multiple modules: just use `nn.Sequential` to link them all.
7         self.layer1 = nn.Sequential(
8             nn.Conv2d(1, 16, kernel_size=5, stride=1, padding=2),
9             nn.BatchNorm2d(16),
10            nn.ReLU(),
11            nn.MaxPool2d(kernel_size=2, stride=2))
12
13        self.layer2 = nn.Sequential(
14            nn.Conv2d(16, 32, kernel_size=5, stride=1, padding=2),
15            nn.BatchNorm2d(32),
16            nn.ReLU(),
17            nn.MaxPool2d(kernel_size=2, stride=2))
18
19        self.fc = nn.Linear(7*7*32, num_classes)
20
21    def forward(self, x):
22        out = self.layer1(x)
23        out = self.layer2(out)
24        out = out.reshape(out.size(0), -1)
25        out = self.fc(out)
26
27    return out
```

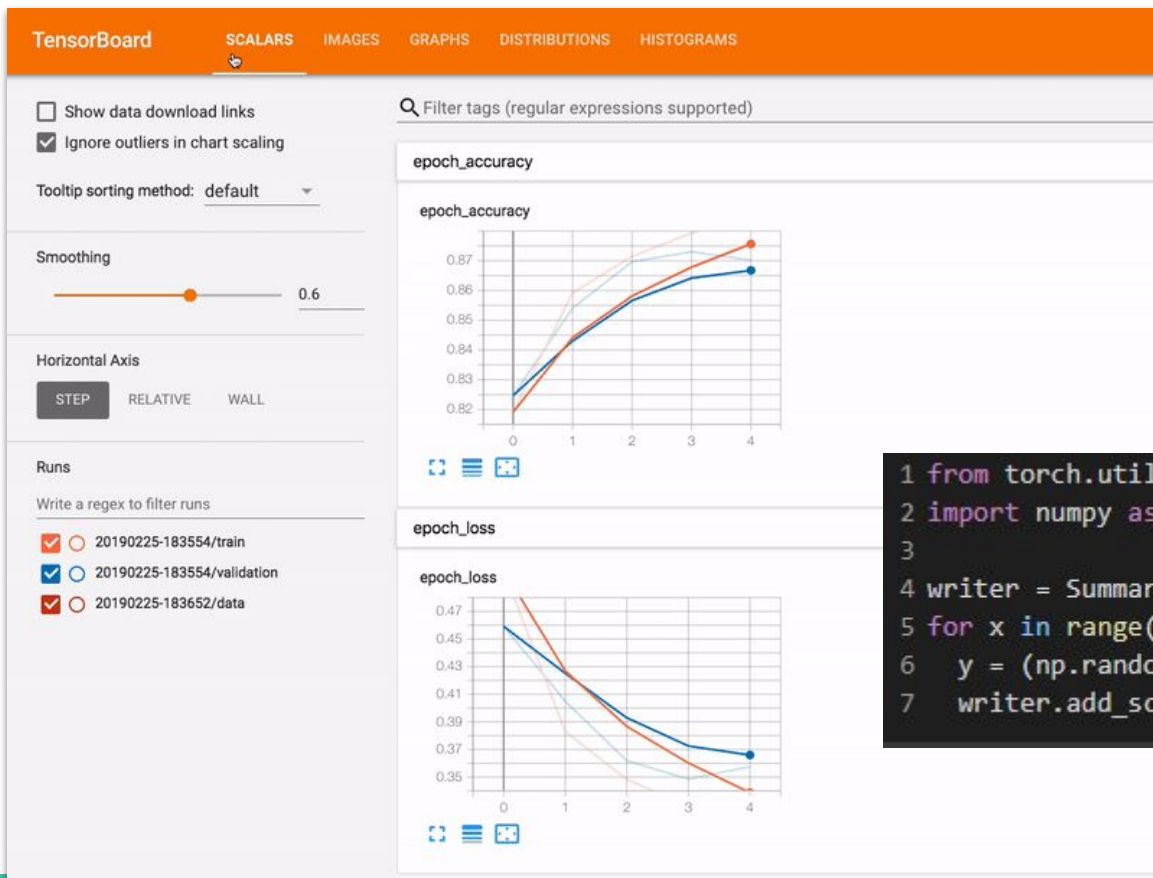
Off-the-shelf Model and Checkpoints - example

```
1 import torchvision
2 import torchvision.models as models
3
4 print(models.mobilenet_v2())

MobileNetV2(
  (features): Sequential(
    (0): ConvBNReLU(
      (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU6(inplace=True)
    )
    (1): InvertedResidual(
      (conv): Sequential(
        (0): ConvBNReLU(
          (0): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), groups=32, bias=False)
          (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
          (2): ReLU6(inplace=True)
        )
        (1): Conv2d(32, 16, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (2): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (2): InvertedResidual(
      (conv): Sequential(
        (0): ConvBNReLU(
          (0): Conv2d(16, 96, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
          (2): ReLU6(inplace=True)
        )
        (1): ConvBNReLU(
          (0): Conv2d(96, 96, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), groups=96, bias=False)
          (1): BatchNorm2d(96, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
          (2): ReLU6(inplace=True)
        )
        (2): Conv2d(96, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (3): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (3): InvertedResidual(
      (conv): Sequential(
```

MobileNet is a light-weight network designed for mobile devices, which has 17 *InvertedResidual* Blocks each with multiple Convolution and Normalization layers.

TensorBoards: Experiment Visualization.



Tracks some key metrics during model training such as loss, accuracy and histogram of gradients, which is super useful for development.

```
1 from torch.utils.tensorboard import SummaryWriter
2 import numpy as np
3
4 writer = SummaryWriter()
5 for x in range(100):
6     y = (np.random.random() + x) / 100
7     writer.add_scaler("Random Values", y, x)
```

PyTorch Resources

- Step-by-step live demo:

https://colab.research.google.com/drive/1L2US3PF6j2kzoOemkeedAb_E_hbYjVMP

- Official Documentation:

<https://pytorch.org/docs/stable/index.html>

- PyTorch Tutorial from its developers:

<https://pytorch.org/tutorials/>

- Plenty of text/visual resources.