

Transactions in SQL

Queries and Transactions

- ❑ Queries: requests to the DBMS to retrieve data from the database
- ❑ Updates: requests to the DBMS to insert, delete or modify existing data
- ❑ Transactions: logical grouping of query and update requests to perform a task
 - Logical unit of work (like a function/subroutine)

The chicken and the egg problem...

```
CREATE TABLE Chicken (ID INT PRIMARY KEY,
    eID INT NOT NULL REFERENCES Egg(ID));
CREATE TABLE Egg(ID INT PRIMARY KEY,
    cID INT NOT NULL REFERENCES Chicken(ID));
```

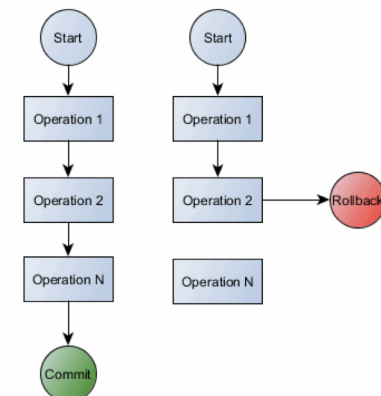


- ❑ Do we know commands that could create these tables?
- ❑ But how can we insert values into either table??
 - Need to treat two inserts into both tables as one logical unit of work...



SQL TRANSACTIONS

- ❑ Start:
 - Each SQL statement *implicitly* start a transaction, unless one is active.
 - Multi-SQL statement transaction is within **BEGIN;**
...
END;
- ❑ **COMMIT ;**
- ❑ **ROLLBACK** default action



Standard SQL TRANSACTIONS

- ❑ SET TRANSACTION [transaction characteristics];
 - It does not start a transaction
 - It can be invoked between transactions to set the next transaction to be activated
- ❑ Transaction characteristics: READ WRITE | READ ONLY
 - SQL1: DECLARE TRANSACTION [READ WRITE | READ ONLY];
 - SQL2/SQL3: SET TRANSACTION [READ WRITE | READ ONLY];
- ❑ SQL3 introduced START TRANSACTION [transaction characteristics];
 - It starts a transaction explicitly if one is not active
- ❑ COMMIT ;
- ❑ ROLLBACK; -- ROLLBACK default action

CS1555/2055, Panos K. Chrysanthis & Constantinos Costa – University of Pittsburgh

5

SQL transactions in PostgreSQL

- ❑ Basic transaction statements (SQL3 plus):
 - START TRANSACTION [READ WRITE | READ ONLY];
 - BEGIN [TRANSACTION] [READ WRITE | READ ONLY];
 - ... should be unnecessary according to the SQL standard. Each SQL statement should implicitly start a transaction
 - COMMIT: Unless START TRANSACTION is issued, PostgreSQL implicitly issues a COMMIT after each SQL statement
 - This functionality is sometimes referred to as *autocommit*
- ❑ You cannot effectively have a multi-statement transaction without issuing a START TRANSACTION

CS1555/2055, Panos K. Chrysanthis & Constantinos Costa – University of Pittsburgh

6

Transaction Atomicity (Atomic Block)

- ❑ “All or nothing” can be achieved with *begin-end block*
- ❑ Consider a transaction:

```
begin;  
  insert into Student values (23, 'John', 'CS');  
  insert into Dept values ('CS', 501);  
end;
```
- ❑ What happens if the first insert fails, e.g., due to a referential constraint violation?
 - Is the new tuple inserted into Department? No/Yes?
- ❑ The transaction fails and both inserts are aborted!!!

CS1555/2055, Panos K. Chrysanthis & Constantinos Costa – University of Pittsburgh

7

Modes of Constraints Enforcement

- ❑ **NOT DEFERRABLE** or **IMMEDIATE**
 - Evaluation is performed at input time
 - By default constraints are created as NON DEFERRABLE
 - It *cannot* be changed during execution
- ❑ **DEFERRED**
 - Constraints are not evaluated until commit time
- ❑ **DEFERRABLE**
 - It can be changed within a transaction to be DEFERRED using SET CONSTRAINTS
- ❑ Modes can be specified when a table is created.
 - INITIALLY IMMEDIATE: constraint validation to happen immediate
 - INITIALLY DEFERRED: constraint validation to defer until commit

CS1555/2055, Panos K. Chrysanthis & Constantinos Costa – University of Pittsburgh

8

Specifying Initial Eval. Mode in Tables

- ❑ `CREATE TABLE dept(
 dno VARCHAR(50),
 did INTEGER,
 CONSTRAINT dept_PK PRIMARY KEY (dno));`
- ❑ `CREATE TABLE student(
 sid INTEGER,
 name VARCHAR(50),
 dno VARCHAR(50),
 CONSTRAINT student_PK PRIMARY KEY (sid)
 INITIALLY DEFERRED DEFERRABLE,
 CONSTRAINT student_FK FOREIGN KEY (dno)
 REFERENCES dept(dno) INITIALLY IMMEDIATE DEFERRABLE);`
- ❑ Primary key could be DEFERRED/DEFERRABLE **only** if it's not referenced by another table in PostgreSQL

Changing Constraint Evaluation Mode

- ❑ It is permitted only for deferrable constraints
- ❑ Setting the constraint validation mode within a transaction
 - set mode of all deferrable constraints
`SET CONSTRAINTS ALL IMMEDIATE;`
`SET CONSTRAINTS ALL DEFERRED;`
 - set mode of specific deferrable constraints (list)
`SET CONSTRAINTS student_FK IMMEDIATE;`
`SET CONSTRAINTS student_FK DEFERRED;`

Specifying Transaction Atomicity

- ❑ Errors at commit time: only when **deferred constraints** are violated
 - Constraints can be deferred if specified as **deferrable** in the table schema, and
 - deferred in the scope of the transaction
- ❑ E.g., *assume the constraints are deferrable*

```
start transaction read write;  
set constraints all deferred;  
insert into Student values (23, 'John', 'CS');  
insert into Dept values ('CS', 501);  
Commit;
```
- ❑ No constraint violation of the first insert is detected at *commit time* → the whole transaction is committed

Specifying Transaction Atomicity (2)

- ❑ E.g. 2, *assume the constraints are deferrable and assume SID 23 exists in that Database*

```
insert into Student values (23, 'John', 'CS');
```



```
start transaction read write;  
set constraints all deferred;  
insert into Student values (23, 'John', 'CS');  
insert into Dept values ('CS', 501);  
Commit;
```
- ❑ The constraint violation of the first insert is detected at *commit time* → the whole transaction is rollback

The chicken and the egg problem...

```
CREATE TABLE Chicken(ID INT PRIMARY KEY, eID INT NOT NULL);
CREATE TABLE Egg(ID INT PRIMARY KEY, cID INT NOT NULL);

ALTER TABLE Chicken ADD CONSTRAINT Chicken_FK
    FOREIGN KEY (eID) REFERENCES Egg(ID)
    DEFERRABLE INITIALLY IMMEDIATE;

ALTER TABLE Egg ADD CONSTRAINT Egg_FK
    FOREIGN KEY (cID) REFERENCES Chicken(ID)
    DEFERRABLE INITIALLY IMMEDIATE;
```

The chicken and the egg problem...

```
START TRANSACTION READ WRITE;
SET CONSTRAINTS ALL DEFERRED;
    INSERT INTO Chicken VALUES (1, 2);
    INSERT INTO Egg VALUES (2, 1);
COMMIT;
```

SQL Savepoints

- ❑ In long transactions, isolates set of “safe” operations from set of “Risky” Operations
- ❑ Operations
 - SAVEPOINT NAME < name > ;
 - RELEASE SAVEPOINT < name > ;
 - ROLLBACK [WORK] [TO SAVEPOINT < name >];



ANSI SQL2 Isolation Levels

- ❑ SET TRANSACTION [READ ONLY | READ WRITE]
[, ISOLATION LEVEL READ UNCOMMITTED |
READ COMMIT |
REPEATABLE READ |
SERIALIZABLE]
- ❑ Isolation (alias concurrency atomicity / serializability)
 - transactions are independent,
 - the result of the execution of *concurrent transactions* is the same as if transactions were executed serially, one after the other

Example Transaction

- ❑ **CLASS** (classid, max_num_students, cur_num_students)
- ❑ Consider the following transaction

```

START TRANSACTION READ WRITE;
SELECT max_num_students, cur_num_students
FROM CLASS
WHERE classID = 1;
-- Read(classid =1)

If (cur_num_students < max_num_students)
    update CLASS
    set cur_num_students = cur_num_students + 1
    where classID = 1;
-- Write(classid =1)
else
    print 'The class is full';
COMMIT;
    
```

CS1555/2055, Panos K. Chrysanthos & Constantinos Costa – University of Pittsburgh

17

Transactions t1 and t2 - Serial Execution

- ❑ Assume *max_num_students* = 40, *cur_num_students* = 39
- ❑ Execution of transaction t1 and t2:


```

r1(max_num_students)
r1(cur_num_students)           -- cur_num_students = 39
If (cur_num_students < max_num_students)
    w1(cur_num_students++)      -- cur_num_students = 40

r2(max_num_students)
r2(cur_num_students).          -- cur_num_students = 40
If (cur_num_students < max_num_students) -- 40 < 40 - false
    "The class is full"
            
```

CS1555/2055, Panos K. Chrysanthos & Constantinos Costa – University of Pittsburgh

18

Concurrent Transactions

```

START TRANSACTION READ WRITE;
SELECT max_num_students, cur_num_students
FROM CLASS
WHERE classID = 1;

sleep...

If (cur_num_students < max_num_students)
    update CLASS
    set cur_num_students = cur_num_students + 1
    where classID = 1;
else
    print 'The class is full';
COMMIT;
    
```

CS1555/2055, Panos K. Chrysanthos & Constantinos Costa – University of Pittsburgh

19

Transactions t1 and t2 - Concurrent Execution

- ❑ Assume *max_num_students* = 40, *cur_num_students* = 39
- ❑ Execution:


```

r1(max_num_students)
r1(cur_num_students)           -- cur_num_students = 39
... sleep1

r2(max_num_students)
r2(cur_num_students).          -- cur_num_students = 39
... sleep2

If (cur_num_students < max_num_students)
    w1(cur_num_students++)      -- cur_num_students = 40
If (cur_num_students < max_num_students)
    w1(cur_num_students++)      -- cur_num_students = 41
            
```

CS1555/2055, Panos K. Chrysanthos & Constantinos Costa – University of Pittsburgh

20

Write/Exclusive Lock

- ❑ Example:

```
SELECT max_num_students, cur_num_students  
FROM CLASS  
WHERE classID = 1555  
FOR UPDATE OF cur_num_students;
```

- ❑ Alternative just specify **FOR UPDATE;**

- ❑ Error Messages:

- No lock: "Cannot serialize access for this transaction"
- With lock: "The class is full"