

## Relation Schema

STUDENT	SID	LName	Name	Class	Major
	123	Smith	John	3	CS
	395	Aiken	Mary	4	CS

← Schema

### What is the meaning?

- A **relation schema** R specifies
  - The name of the relation
  - the attribute names  $A_i$  of R
  - the domain  $D_i$  (data type + format) for each attribute  $A_i$
- data type is a set of **atomic data** values:
  - no attribute is a set-valued (**1st Normal Form**, 1-NF)
  - no attribute is composite
- format specifies the representation of a data value

## Example Table Schema

Schema of STUDENT(SID, Name, Major, GPA)

```
CREATE TABLE STUDENT
(
  SID INTEGER,
  Name CHAR(20),
  Major CHAR(4),
  GPA DEC(3,2)
);
```

## Creating a Schema



- Corresponding database is at an **empty** state!
- **Initial state** when the database is **populated** (loaded)
- Domain (type) of each field is **specified** and **enforced** by the DBMS whenever tuples are added or modified

## Example: Domain Constraints

SID	Name	Login	Age	GPA
546007	Jones	jones@cs	18	3.4
546100	Smith	smith@ee	18	3.2
546500	Smith	smith@math	19	3.8

### Example of IC Violation:

UPDATE Students

SET Age = `Eighteen`

x x x

WHERE Name = Jones;

## Useful Terms

- **Cardinality** of a relation  $r(R)$ : # of tuples in  $r(R)$  (denoted by  $|r(R)|$ )
- **Arity** or **degree** of  $r(R)$ : # of attributes in  $R$  (denoted by  $|R|$ )

$$|R| = 4$$

$$|r(R)| = 3$$

SID	Degree	Major	Year
123	BS	Math	1992
064	BA	History	1991
445	PhD	CS	1999

- ◆  $|R| > 0$  and  $|r(R)| \geq 0$
- ◆ Cardinality is property of a relation
- ◆ Arity is property of relation schema or a relation

## Relational Database Schema

- A *database schema* is a set of relation schemas and a set of **integrity constraints**



- Integrity Constraints
  - **Structural** Integrity Constraints
    - key constraints: uniqueness of keys
    - entity integrity constraint: no primary key value can be **NULL**
    - referential integrity constraint
  - **Semantic** Integrity Constraints
    - E.g., ??

## Integrity Constraints (ICs)

- **IC**: condition that must be true for any instance of the database (e.g., domain constraints)
  - A **legal** instance of a relation is one that satisfies all specified ICs
  - ICs are specified when schema is **defined**
  - ICs are enforced when tables are **modified**



## Primary Key Constraint

- A set of fields is a **key** for a relation if :
  - No two distinct tuples can have same values in all key fields



- If there is more than one key for a relation:
  - Each is called a candidate key
  - One candidate key is designated as the **primary key**
  - Other candidate key(s) are designated as **alternative** or **unique key(s)**



## Example of Keys

<i>SID</i>	<i>Name</i>	<i>Login</i>	<i>Age</i>	<i>GPA</i>
546007	Jones	jones@cs	18	3.4
546100	Smith	smith@ee	18	3.2
546500	Smith	smith@math	19	3.8

- ❑ **Candidate Keys:** *SID*, and *Login*
- ❑ **Primary Key:** *SID*
- ❑ **Unique Key:** *Login*

## Example Table Schema in SQL

Schema of STUDENT(*SID*, *Login*, Name, Major, GPA)

```
CREATE TABLE STUDENT
(
  SID INTEGER NOT NULL,
  Login CHAR(15),
  Name CHAR(20),
  Major CHAR(4),
  GPA DEC(3,2),

  CONSTRAINT STUDENT_PK
    PRIMARY KEY (SID),
  CONSTRAINT STUDENT_UN
    UNIQUE (Login) -- UNIQUE can take NULL values
);
```

## Example Table Schema in SQL (2)

Schema of STUDENT(*SID*, *Login*, Name, *SSN*, GPA)

```
CREATE TABLE STUDENT
(
  SID INTEGER NOT NULL,
  Login CHAR(15),
  Name CHAR(20),
  SSN CHAR(9),
  GPA DEC(3,2),

  CONSTRAINT STUDENT_PK PRIMARY KEY (SID),
  CONSTRAINT STUDENT_UN_SSN
    UNIQUE (SSN),
  CONSTRAINT STUDENT_UN_Login
    UNIQUE (Login)
);
```

## Example Table Schema in SQL (3)

Schema of STUDENT(*SID*, *SSN*, Name, *Login*, GPA)

```
CREATE TABLE STUDENT
(
  SID INTEGER
    CONSTRAINT STUDENT_PK_NOT_NULL NOT NULL,
  SSN CHAR(9)
    CONSTRAINT STUDENT_PK PRIMARY KEY,
  CONSTRAINT STUDENT_UN_SSN UNIQUE,
  Name CHAR(20),
  Login CHAR(15),
  GPA DEC(3,2),

  CONSTRAINT STUDENT_UN_Login
    UNIQUE (Login)
);
```

## Identifying the Key

- What is the key in relation GRADUATE=(SID, Degree, Major, Year) ?

<i>SID</i>	<i>Degree</i>	<i>Major</i>	<i>Year</i>
123	BS	CS	1992
123	MS	CS	1993
064	BA	History	1991
445	PhD	CS	1999
123	BS	Math	1992
123	MS	Math	1992



## Foreign Keys

- Foreign key (FK)** in relation  $R_2$  is a set of attributes of  $R_2$  that forms a primary key (PK) of another relation  $R_1$ .
  - Attributes in FK and PK have the **same domain**

### STUDENT

<i>SID</i>	<i>Degree</i>	<i>Major</i>	<i>Year</i>
------------	---------------	--------------	-------------

PK

### COURSE

<i>CID</i>	<i>Name</i>
------------	-------------

PK

### Enrolled

<i>SID</i>	<i>CID</i>	<i>Grade</i>
------------	------------	--------------

PK

FK

FK

## Foreign Key & Primary Key

<i>SID</i>	<i>Name</i>	<i>Age</i>	<i>SID</i>	<i>CID</i>	<i>Grade</i>	<i>CID</i>	<i>Name</i>
546007	Alex	18	546007	CS1555	A	CS1555	DB
546100	Bob	19	546007	CS1550	B+	CS1530	SW
546500	Bill	20	546500	CS1555	B	CS1550	OS

Primary Key

Students

Foreign Key

Enrolled

Courses

- Foreign key:** Set of fields in one relation that is used to “refer” to a tuple in another relation
  - Must correspond to primary key of the referred relation
  - E.g. *SID* is a foreign key referring to *Students*

## Foreign Key Constraints

- If foreign key constraints are enforced, **referential integrity** is achieved
  - E.g.: Only students can enroll in a class
    - Only students listed in the “Students” relation should be allowed to enroll for courses
- Like a “logical pointer”
  - There shouldn’t be dangling references
    - Either valid PK or NULL

## Any Attribute can be a Foreign Key

Faculty			Courses		
<u>FID</u>	Name	Area	<u>CID</u>	Name	Instructor
007	Panos	DB	CS1555	DB	007
100	Daniel	OS	CS1530	SW	NULL
500	Adriana	AI	CS1550	OS	100

Primary Key

Primary Key

Foreign Key

- Foreign key: Set of fields in one relation that is used to "refer" to a tuple in another relation
  - Must correspond to primary key of the referred relation
  - If not part of a key, it could be NULL

## Foreign Keys in SQL

<u>SID</u>	Name	Major	GPA	<u>CID</u>	Name	<u>SID</u>	<u>CID</u>	Grade
Students				Courses		Enrolled		

```

CREATE TABLE Enrolled (
    SID CHAR(20), CID CHAR(20), Grade CHAR(2),
    CONSTRAINT Enrolled_PK PRIMARY KEY (SID, CID),
    CONSTRAINT Enrolled_FK_sid
        FOREIGN KEY (SID) REFERENCES Students (SID),
    CONSTRAINT Enrolled_FK_cid
        FOREIGN KEY (CID) REFERENCES Courses
);
    
```

## Referential Integrity Constraints

<u>SID</u>	Name	Age	<u>SID</u>	<u>CID</u>	Grade	<u>CID</u>	CName
546007	Susan	18	546007	CS1555	A	CS1555	DB
546100	Bob	19	546007	CS1530	B+	CS1530	SW
546500	Bill	20	546500	CS1555	B	CS1550	OS

Students

Enrolled

Courses

- Any IC Violation? [Poll]

DELETE

FROM Enrolled

WHERE SID = 546500;

No Violations! ✓

## Referential Integrity Constraints

<u>SID</u>	Name	Age	<u>SID</u>	<u>CID</u>	Grade	<u>CID</u>	CName
546007	Susan	18	546007	CS 1555	A	CS 1555	DB
546100	Bob	19	546007	CS1530	B+	CS1530	SW
546500	Bill	20	546500	CS 1555	B	CS 1550	OS

Students

Enrolled

Courses

- Any IC Violation? [Poll]

DELETE

FROM Students

WHERE SID = 546500;

Violation! ✗


## Referential Integrity Enforcement

- ❑ What are the alternatives when a “Students” tuple is deleted?
  1. **Delete all** Enrolled tuples that refer to it
  2. **Disallow** deletion of a Students tuple that is referred to
  3. **Set** *SID* in Enrolled tuples that refer to it to some “default” *SID* (e.g., 000000)
  4. If *SID* was not part of the primary key, **Set** *SID* to a special value “**NULL**”, denoting “unknown” or “inapplicable”

## Referential Integrity in SQL

- ❑ SQL/92 and SQL/99 support all 4 options on delete and update:
  - **NO ACTION** (default)
    - delete/update is rejected
  - **CASCADE**
    - also delete all tuples that refer to deleted tuple
  - **SET NULL / SET DEFAULT**
    - sets foreign key value of referencing tuple

## RI Trigger Actions in SQL



- ❑ **CREATE TABLE Enrolled (**  
*SID* CHAR(20), *CID* CHAR(20), *Grade* CHAR(2),  
**CONSTRAINT Enrolled\_PK PRIMARY KEY (SID, CID),**  
**CONSTRAINT Enrolled\_FK\_sid**  
**FOREIGN KEY (SID) REFERENCES Students (SID),**  
**CONSTRAINT Enrolled\_FK\_cid**  
**FOREIGN KEY (CID) REFERENCES Courses**  
**ON UPDATE CASCADE ON DELETE NO ACTION**  
**);**

## Enforcing Integrity Constraints [Poll]



- ❑ What would be the outcome?
  - Insert (585811, 'Jie', 19, 3.95) into Students ✓
  - Insert (585811, NULL, NULL) into Enrollment ✗
  - Insert (546100, 'CS 1555', NULL) into Enrollment ✓
  - Insert (546100, 'Mary', 18, 3.65) into Students ✗
  - Delete ('CS 1530') from Courses ✗

<i>SID</i>	<i>Name</i>	<i>Age</i>	<i>GPA</i>
546007	Susan	18	3.8
546100	Bob	19	3.65
546500	Bill	20	3.7

**Students**

<i>CID</i>	<i>Name</i>
CS1555	DB
CS1530	SW
CS1550	OS

**Courses**

<i>SID</i>	<i>CID</i>	<i>Grade</i>
546007	CS1550	A
546007	CS1530	B+
546100	CS1550	B

**Enrollment**