

COE 1195: Advanced Digital Design

Lab 5 – 2D Image Convolution

Dr. Amr Mahmoud



Exercise: 2D Convolutional Accelerator using Vivado HLS

2D Convolution

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

Input Image

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Kernel

x1	x0	x1
x0	x1	x0
x1	x0	x1

*

=

Output Image

	4			

Access Pattern

- Software Approach
 - Two nested loops
 - Less efficient in hardware
 - Pixels are *read* more than once
- Stream Approach
 - One single loop
 - Efficient in hardware
 - Pixels are *read* only once

Software Approach

```
for (x = 0; x < WIDTH; x++)
```




Stream Approach

```
for (i = 0; i < WIDTH*HEIGHT; i++)
```



Memory Path (1/3)

- Kernel
 - Need $K \times K$ registers
 - Hold values to compute kernel to produce output pixel
- Line Buffer
 - Need $K - 1$ line buffers of length $Width - K$
 - Temporarily store image line to access previous pixels without re-accessing from stream source




Stream Approach



Memory Path (2/3)

- Kernel
 - Need $K \times K$ registers
 - Hold values to compute kernel to produce output pixel
- Line Buffer
 - Need $K - 1$ line buffers of length $Width - K$
 - Temporarily store image line to access previous pixels without re-accessing from stream source

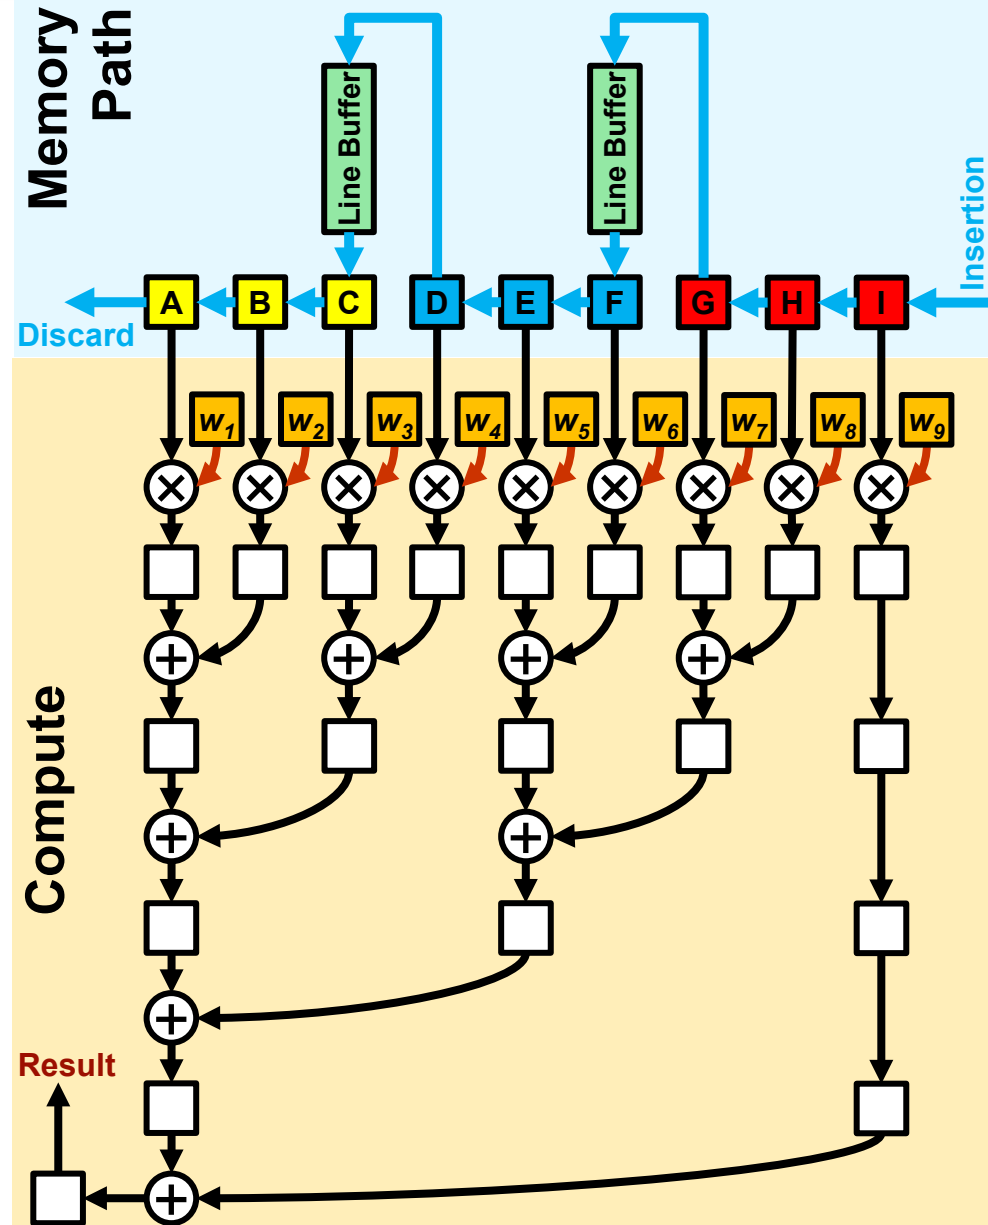



Stream Approach



Data Path

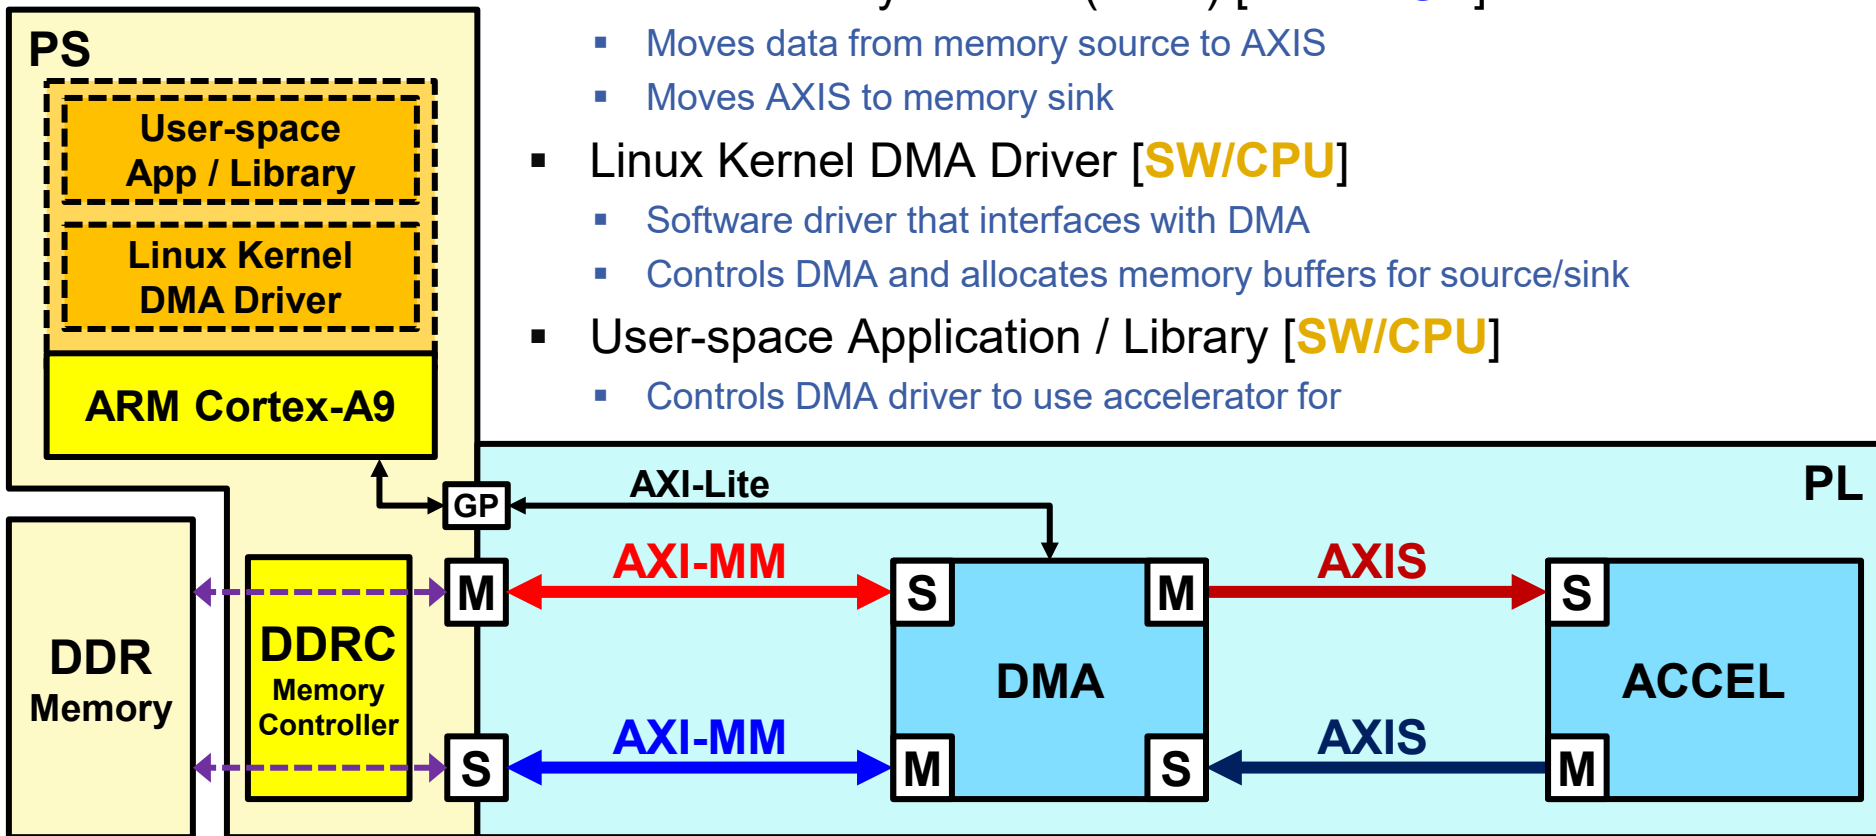
- Multiplications
 - Unrolled
 - $K \times K$ multiplications per iteration
- Additions
 - Implicit reduction
 - Pipelined binary tree



Acceleration Framework



- Accelerator [**HW/FPGA**]
 - AXI-Stream (AXIS) input and output
- Direct Memory Access (DMA) [**HW/FPGA**]
 - Moves data from memory source to AXIS
 - Moves AXIS to memory sink
- Linux Kernel DMA Driver [**SW/CPU**]
 - Software driver that interfaces with DMA
 - Controls DMA and allocates memory buffers for source/sink
- User-space Application / Library [**SW/CPU**]
 - Controls DMA driver to use accelerator for



DMA Architectures

DMA Architectures

- Direct Memory Access (DMA)

- “Data Mover” Behavior


- Asynchronously move data from MM/stream to stream/MM

- Register Space

- Control/Status (for controller)
 - MM2S/S2MM addresses
 - MM2S/S2MM lengths

- Operations

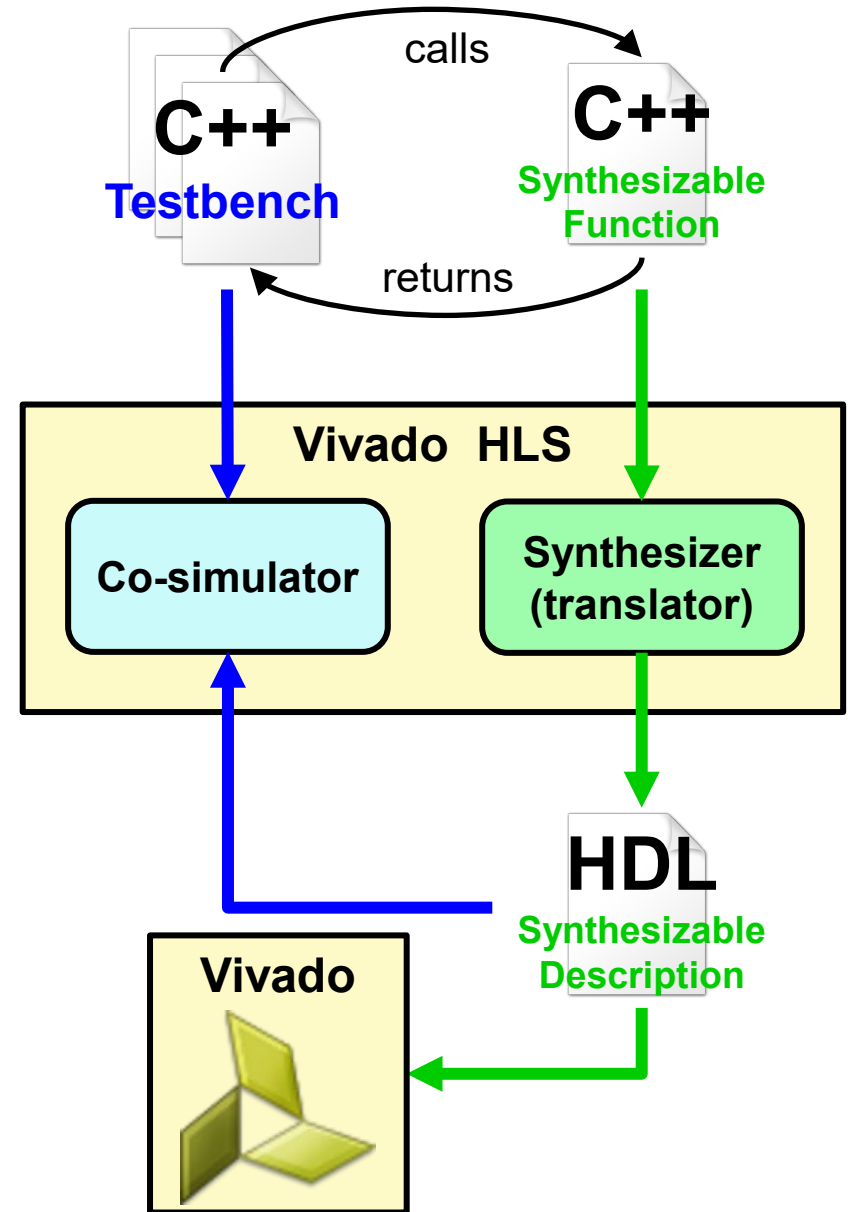
- Stream data in MM starting at *MM2S address* of *MM2S length*
 - Stream data of *S2MM length* to MM starting at *S2MM address*



Vivado HLS


Vivado HLS

- HLS translator
 - Input: **synthesizable function**
 - C++ function
 - HLS directives and rules
 - Output: HDL
- Co-simulator
 - **Testbench** (C++) that calls synthesizable function
 - Executes testbench and simulates translated synthesizable function on PC



INTERFACE (1/2)

- AXI-Stream Interfaces
 - Hint to translator that arguments are AXIS interfaces




```
void hw_conv(stream_t &sin, stream_t &sout) {  
# pragma HLS INTERFACE ap_ctrl_none port=return  
# pragma HLS INTERFACE axis port=sin  
# pragma HLS INTERFACE axis port=sout  
  ...  
}
```



INTERFACE (2/2)

- AXI-Stream Beats

- Contain **Data**, **Keep**, **Last**, **Valid**, **Ready**
- Read (>>) and write (<<)
- Data** (value), **Keep** (byte strobe), **Last** (last beat) are explicit
- Valid** and **Ready** are implicit on >> and <<



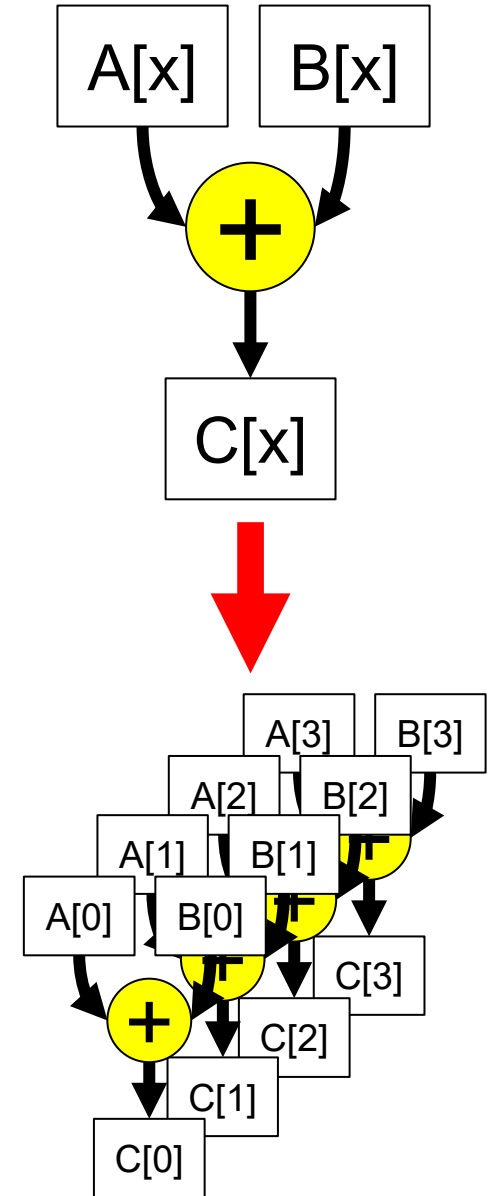
```
... // read beat from sin
beat_t beat;
sin >> beat;
value = beat.data(7, 0);
... // write beat to sout
beat_t beat;
beat.data(7, 0) = value;
sout << beat;
...
```



UNROLL

- Parallelize Loop Iterations
 - Without unrolling
 - Iterations occur one at a time (N cycles each)
 - Trade-off: less area, less bandwidth
 - With unrolling
 - Iterations occur in parallel
 - Trade-off: more area, more bandwidth
 - factor** \triangleq unrolling factor (# of iterations)

```
for (int x; x < 3; x++) {  
# pragma HLS UNROLL factor=4  
  C[x] = A[x] + B[X]  
}
```



PIPELINE

■ Pipeline Loop Iterations

■ Without pipelining

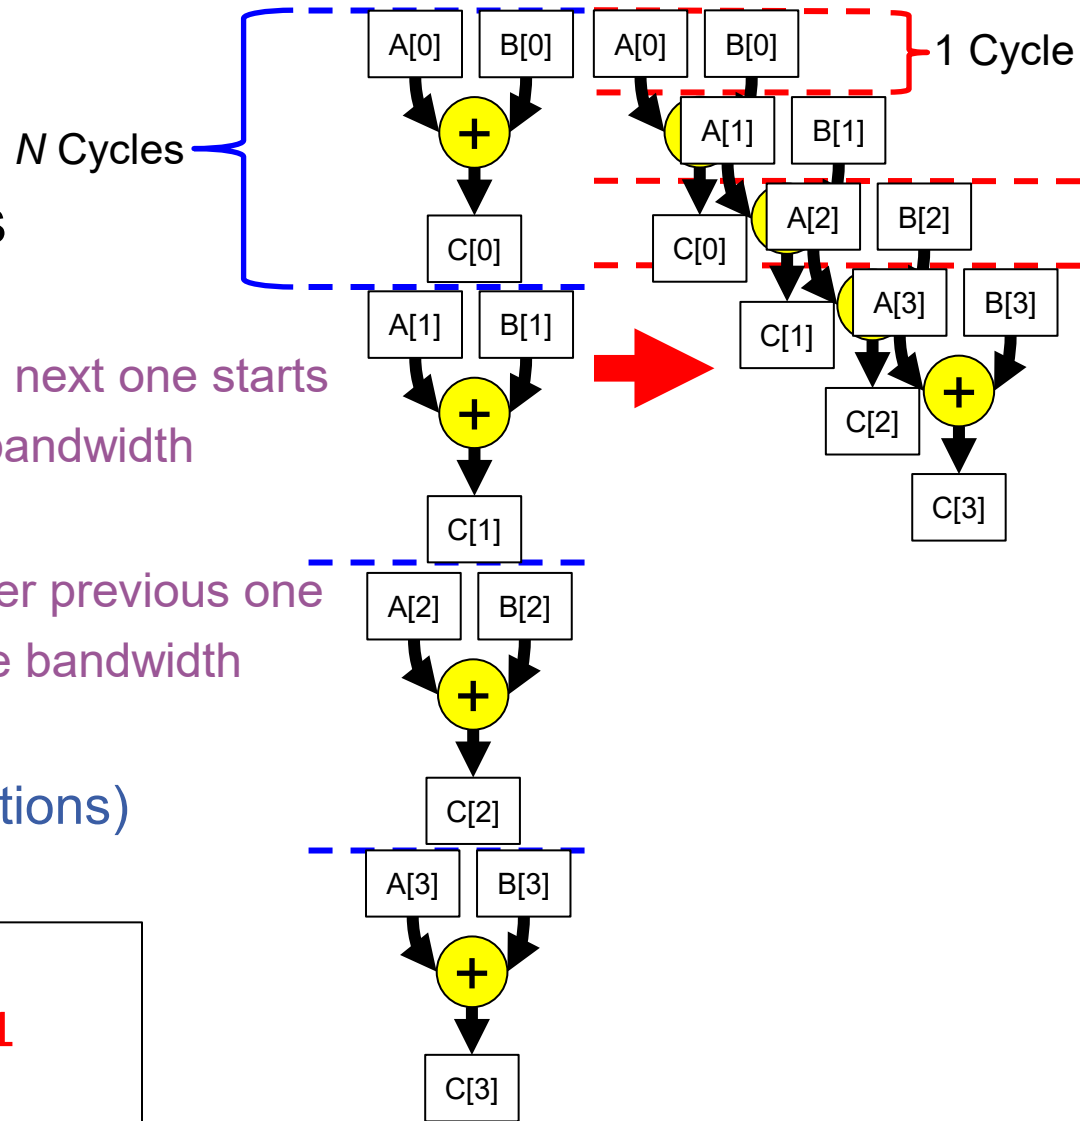
- Iteration must finish before next one starts
- Trade-off: less area, less bandwidth

■ With pipelining

- Iteration starts // cycles after previous one
- Trade-off: more area, more bandwidth


- $II \triangleq$ initiation interval
(# of cycles between iterations)

```
for (int x; x < 3; x++) {
#  pragma HLS PIPELINE II=1
  C[x] = A[x] + B[X]
}
```



ARRAY_PARTITION

- Partition Array into smaller arrays or individual elements
 - Cyclic (partition into sub-arrays with cyclic assignment)
 - Block (partition into sub-arrays)
 - Complete (partition into individual elements)
- Factor
- Dimensionality



```
char array[8];
```

```
#pragma HLS ARRAY_PARTITION variable=array cyclic factor=3 dim=0
```

```
#pragma HLS ARRAY_PARTITION variable=array block factor=3 dim=0
```

```
#pragma HLS ARRAY_PARTITION variable=array complete dim=0
```