



University of Pittsburgh

ECE 2195: Special Topics – Computers Machine Learning

Classification – Multiple Classes, LDA & QDA

Mai Abdelhakim, PhD

Assistant Professor of ECE

Swanson School of Engineering

University of Pittsburgh

maia@pitt.edu



Discriminant Analysis for Classification

- Linear Discriminant Analysis
- Quadratic Discriminant Analysis

Bayes Theorem

$$\Pr(Y = k|X = x) = \frac{\Pr(X = x|Y = k) \cdot \Pr(Y = k)}{\Pr(X = x)}$$

X contains Features (e.g., number of words in an email). Y is the class label (e.g. Y=0 means not spam, Y=1 means spam email).

- Another way to write the equation above is:

$$\Pr(Y = k|X = x) = \frac{\pi_k f_k(x)}{\sum_{l=1}^K \pi_l f_l(x)}$$

K: number of classes
k: is a class label {1,2, ... K}

$f_k(x) = \Pr(X = x|Y = k)$: is the **density of x** given that it is an observation from class k

$\pi_k = \Pr(Y = k)$: is the **prior probability** of class k

Discriminant Analysis

- **Approach:** Model the **distribution of each class** separately, then use Bayes Rule
- Recall Bayes Classifier: Assign a new observation with features x_0 to class k that has largest $\Pr(Y = k | X = x_0)$
 $\Pr()$ and $P()$ both stands for probability; the sign " $|$ " reflects conditional (information is given)
- Both Linear and Quadratic Discriminant analysis **use normal (Gaussian) distribution to model features in each class** → $P(X|y)$ is Gaussian
- More popular than logistic regression when there is **more than 2 classes**

Discriminant Analysis (LDA) with One Feature

- Assume multiple classes, and one feature ($p=1$)
- With one feature, the Gaussian density is given by:

$$f_k(x) = \Pr(X = x|Y = k) = \frac{1}{\sqrt{2\pi}\sigma_k} e^{-\frac{1}{2}\left(\frac{x-\mu_k}{\sigma_k}\right)^2}$$

μ_k is the *mean* of x in class k (class specific mean), σ_k^2 is the *variance* of x in class k

$$\Pr(Y = k|X = x) = \frac{\pi_k f_k(x)}{\sum_{l=1}^K \pi_l f_l(x)}$$

- Using training data, we estimate μ_k, σ_k, π_k for each class k

Linear Discriminant Analysis (LDA) assumes equal variance (or covariance) in all classes

Use Training Data for Estimation

- **Prior** of class k is estimated as the training observations n_k that belong to the k^{th} class divided by the total number of observations n

$$\hat{\pi}_k = \frac{n_k}{n} \quad .$$

- The **mean** can be estimated by the average of all training observations from the k^{th} class

$$\hat{\mu}_k = \frac{1}{n_k} \sum_{i:y_i=k} x_i$$

- The **variance** can be estimated as weighted average of variances of all k classes

$$\hat{\sigma}^2 = \frac{1}{n - K} \sum_{k=1}^K \sum_{i:y_i=k} (x_i - \hat{\mu}_k)^2$$

LDA - discriminant score

- $\text{Log} (P(x|y=k) P(y=k))$ when x is Gaussian with mean μ_k and variance σ_k^2

$$\delta_k(x) = x \cdot \frac{\mu_k}{\sigma^2} - \frac{\mu_k^2}{2\sigma^2} + \log(\pi_k)$$

LDA - discriminant score

- **Build classifier:** have an estimation for $\Pr(Y = k|X = x)$
- **Prediction:** assign X to class k with largest $\Pr(Y = k|X = x)$
 - Equivalent to **finding the class** that gives the largest discriminant score given by:

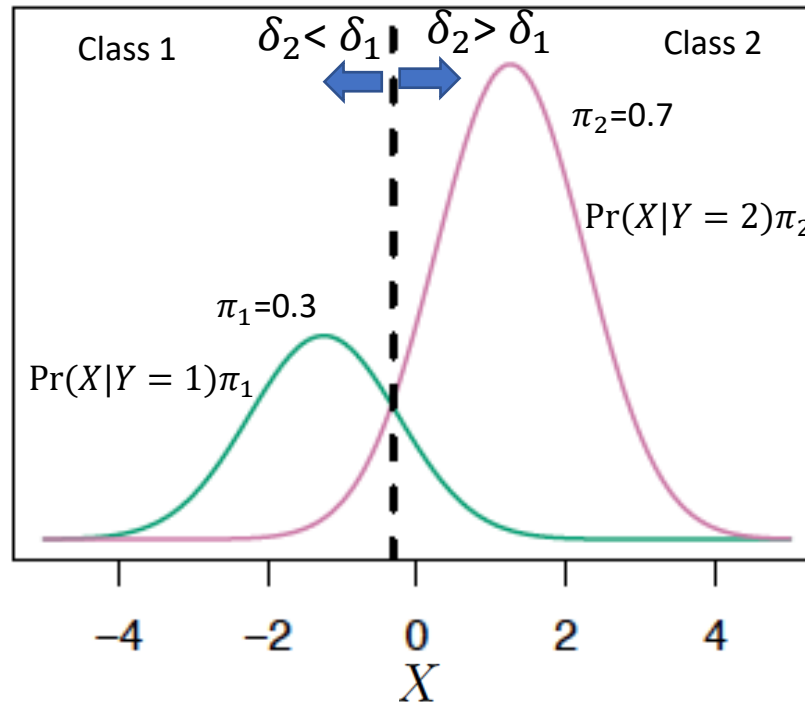
$$\delta_k(x) = x \cdot \frac{\mu_k}{\sigma^2} - \frac{\mu_k^2}{2\sigma^2} + \log(\pi_k)$$

You can get this by taking log and discarding the terms that do not depend on k . Try it!

- The discriminant score is **linear function of x** \rightarrow thus called **Linear Discriminant Analysis**
- Decision boundary is linear

One Feature and Two Classes Example

Classification: new observation is assigned to class with highest $\Pr(Y = k|X = x)$



$$\pi_k = \Pr(Y = k)$$

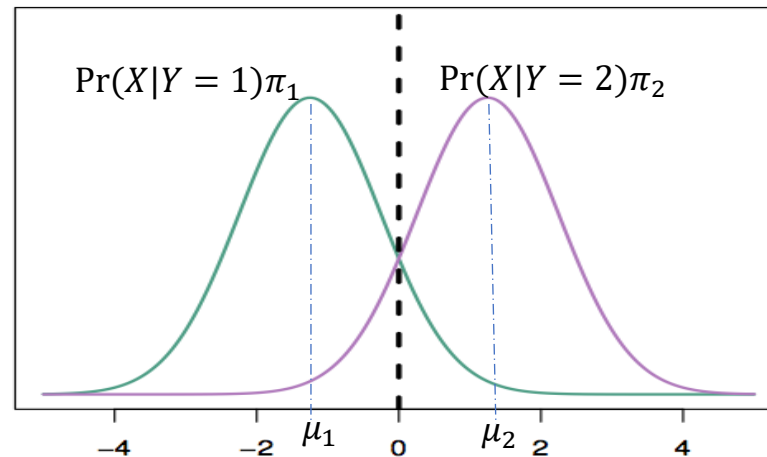
The priors can be estimated by finding the fraction of training samples that belong to each class

Special Case: Two classes (K=1,2), equal priors, one feature

- Assume equal priors $\pi_1 = \pi_2 = 0.5$, we can get a decision boundary as follows:

$$\text{Maximize } \delta_k(x) = x \cdot \frac{\mu_k}{\sigma^2} - \frac{\mu_k^2}{2\sigma^2} + \log(\pi_k) \quad \xrightarrow{\text{Try it!}} \quad x = \frac{\mu_1 + \mu_2}{2}$$

This means that, for $\mu_2 > \mu_1$:



$$\begin{array}{l} \text{Class 2} \\ X \geq \frac{\mu_1 + \mu_2}{2} \\ \text{Class 1} \end{array}$$

i.e.,
choose class 1 if $X < \frac{\mu_1 + \mu_2}{2}$
and class 2 if $X > \frac{\mu_1 + \mu_2}{2}$

Linear Discriminant Analysis When $P > 1$

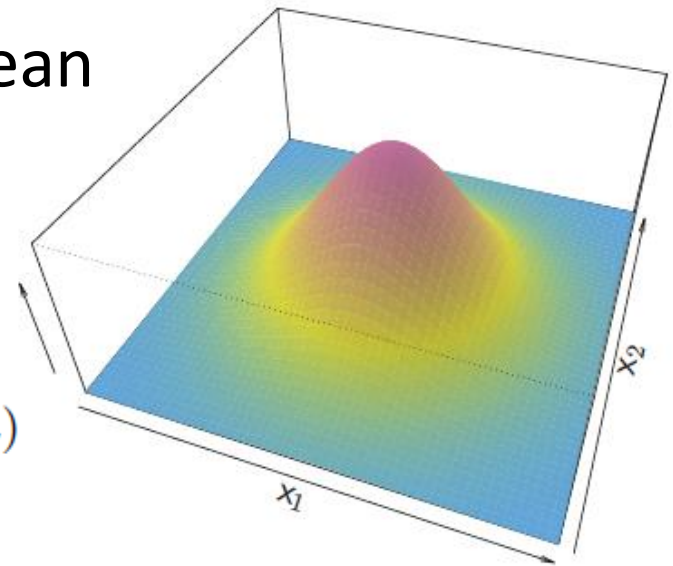
- If X contains multiple features ($p > 1$), same approach is used except that the density function $f(x)$ is modeled using the **multivariate normal density**, i.e.,

$$X \sim \mathcal{N}(\mu_k, \Sigma),$$

- μ_k is the $p \times 1$ mean vector
- Σ is the $p \times p$ covariance matrix (LDA assume same for all classes)

- The multivariate normal density of mean μ and covariance Σ is given by

$$f(x) = \frac{1}{(2\pi)^{p/2} |\Sigma|^{1/2}} e^{-\frac{1}{2}(x-\mu)^T \Sigma^{-1} (x-\mu)}$$



Linear Discriminant Analysis When $P > 1$

- In this case, the discriminant functions will take the form

$$\delta_k(x) = x^T \Sigma^{-1} \mu_k - \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log \pi_k$$

- Still LINEAR with feature vector x

Example: classify Iris flower to one of three possible species ($K=3$):

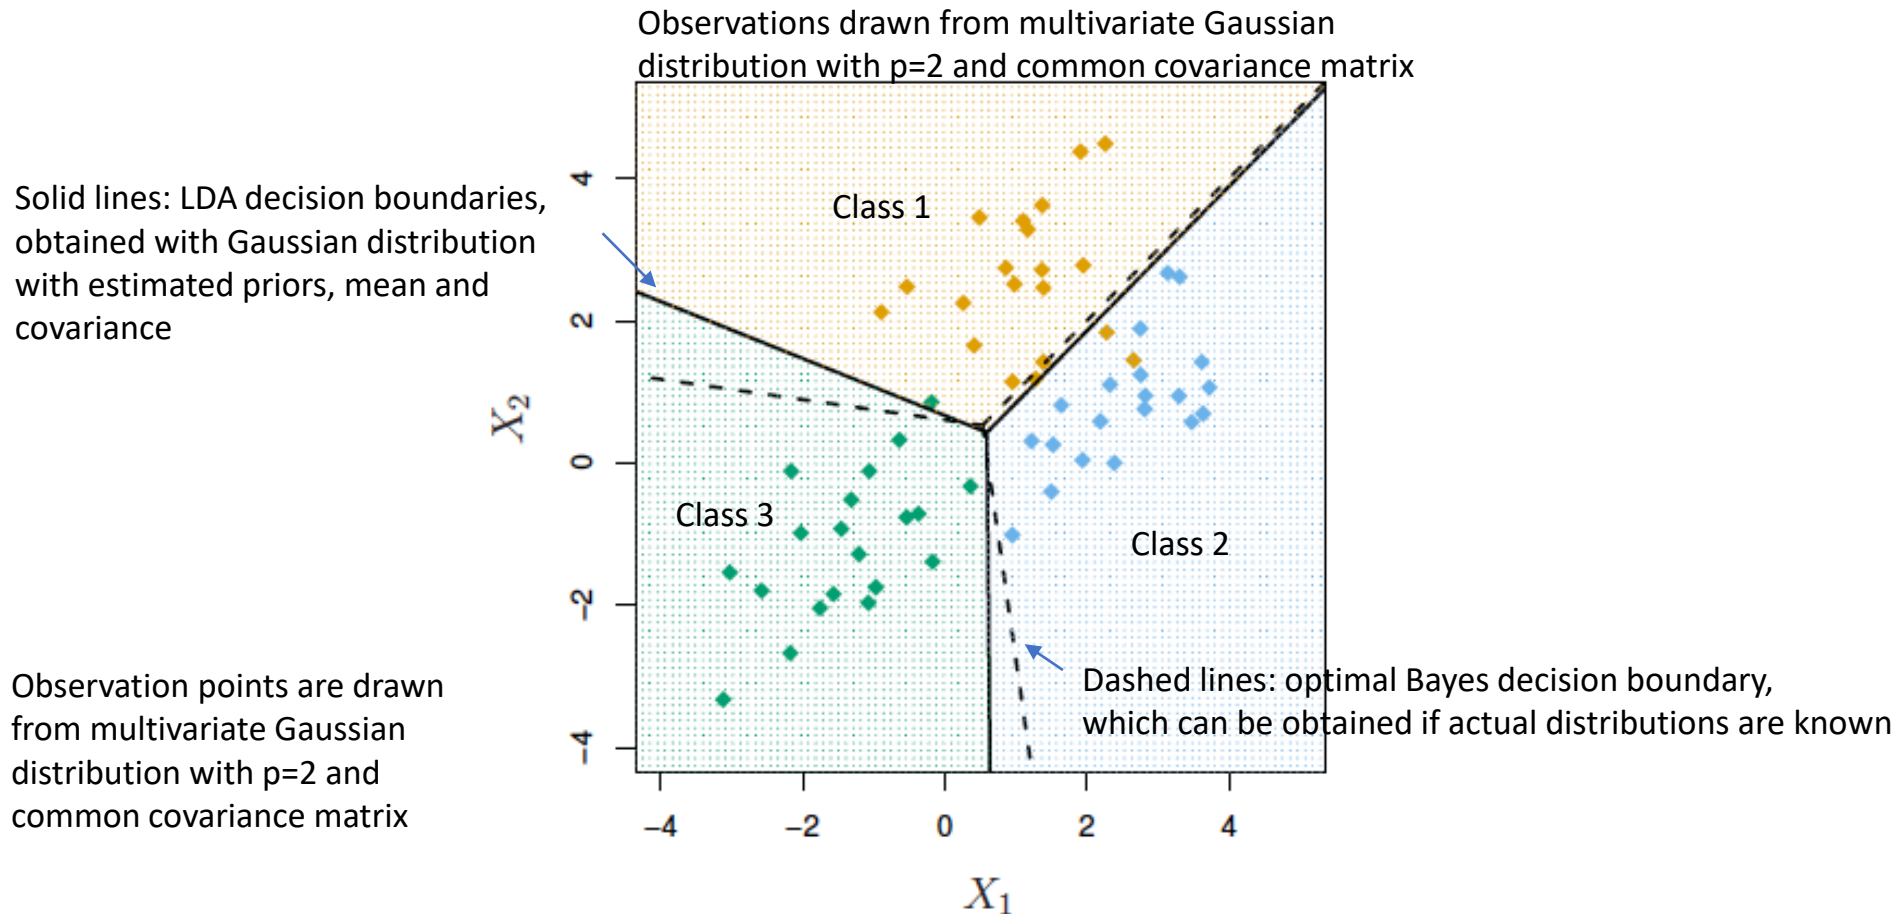
Setosa ($Y=1$), Versicolor ($Y=2$), Virginica ($Y=3$)

Using features: $X = [\text{sepal length}, \text{sepal width}, \text{petal length}, \text{petal width}]$

Build classifier: From the training data, we estimate mean and covariance matrix of this feature vector when $Y=1, 2$ and $3 \Rightarrow$ find discriminant functions / decision boundaries

Classify: Apply the discriminant function to classify new observations

Illustration: Assume 3 classes ($K=3$), and 2 features



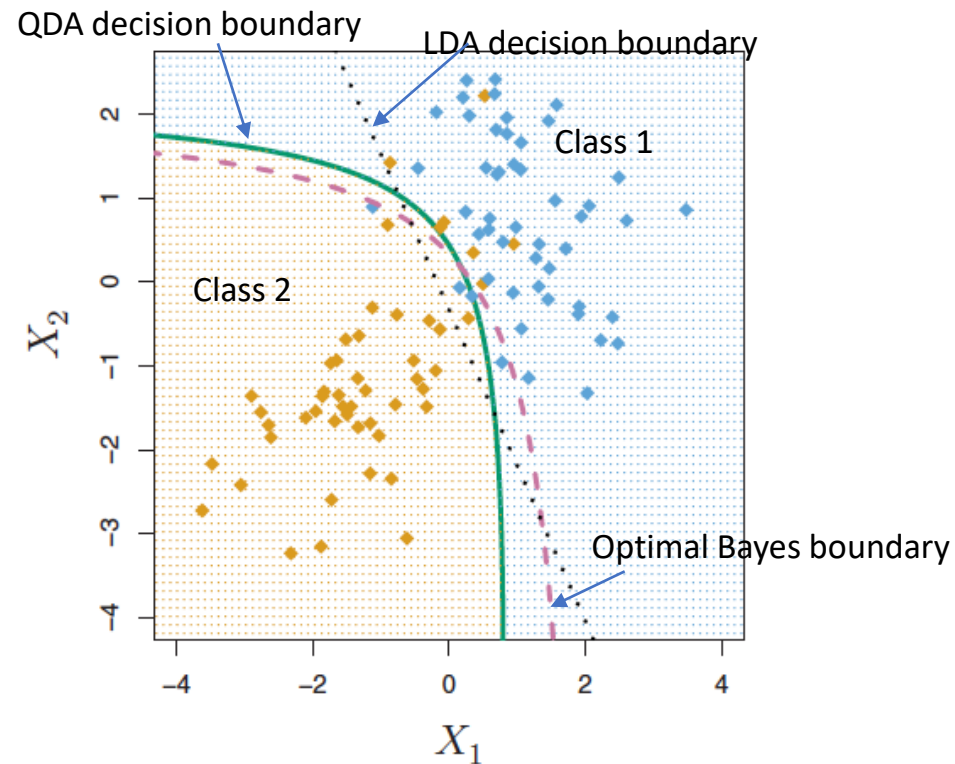
Quadratic Discriminant Analysis (QDA)

- In QDA, the density $f_k(x)$ is assumed to also be Gaussian for each class, but each **class has a different covariance matrix Σ_k**
- In this case we get discriminate functions that are **quadratic** with x (hence the name)

$$\begin{aligned}\delta_k(x) &= -\frac{1}{2}(x - \mu_k)^T \Sigma_k^{-1}(x - \mu_k) - \frac{1}{2} \log |\Sigma_k| + \log \pi_k \\ &= -\frac{1}{2} x^T \Sigma_k^{-1} x + x^T \Sigma_k^{-1} \mu_k - \frac{1}{2} \mu_k^T \Sigma_k^{-1} \mu_k - \frac{1}{2} \log |\Sigma_k| + \log \pi_k\end{aligned}$$

Quadratic Discriminant Analysis – Quadratic Decision Boundary

- Here, 2 features, and 2 classes
 - Two classes have different covariance in this example, hence QDA is better
 - QDA, the covariance matrix is estimated for each class
 - Each matrix has $p(p+1)/2$ parameters
 - Need large training data to avoid overfitting (high variance)
- Bias-variance trade-off!**
- LDA is simpler, but could lead to high bias
 - Estimate single covariance matrix



Naïve Bayes Classifier

- Assumes:
 - Gaussian densities (same as LDA and QDA)
 - Each class has its **own covariance** (same as QDA),
 - Covariance matrix of each class is diagonal matrix (features are statistically independent)
 - Now we need to estimate only p parameters for each covariance matrix

$$\begin{aligned}\delta_k(x) &= -\frac{1}{2}(x - \mu_k)^T \Sigma_k^{-1}(x - \mu_k) - \frac{1}{2} \log |\Sigma_k| + \log \pi_k \\ &= -\frac{1}{2}x^T \Sigma_k^{-1}x + x^T \Sigma_k^{-1}\mu_k - \frac{1}{2}\mu_k^T \Sigma_k^{-1}\mu_k - \frac{1}{2} \log |\Sigma_k| + \log \pi_k\end{aligned}$$

LDA and QDA in Python

- Linear Discriminant Analysis (LDA)

```
from sklearn.discriminant_analysis import  
LinearDiscriminantAnalysis
```

```
LDAmodelFitted = LinearDiscriminantAnalysis().fit(X_train,  
Y_train)
```

- Quadratic Discriminant Analysis

```
from sklearn.discriminant_analysis import  
QuadraticDiscriminantAnalysis
```

```
QDAmodelFitted =  
QuadraticDiscriminantAnalysis().fit(X_train, Y_train)
```