

ECE 0402 - Pattern Recognition

LECTURE 1

Today (1/10): Introduction and Learning Challenge—feasibility of learning.

Wednesday(1/12): Why it is important to do inference from data using statistical perspective?

Example: The Netflix prize: Predict how a user will rate a movie

10% improvement = \$1 million prize

This problem captures the essence of ML:

- a pattern exists
- cannot formulate it analytically
- have data

what we look for in a problem

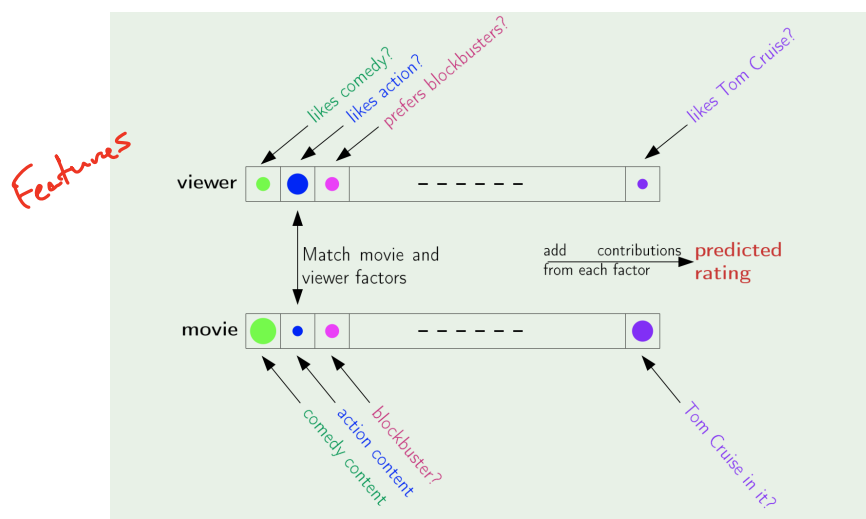


Figure 1: * A solution




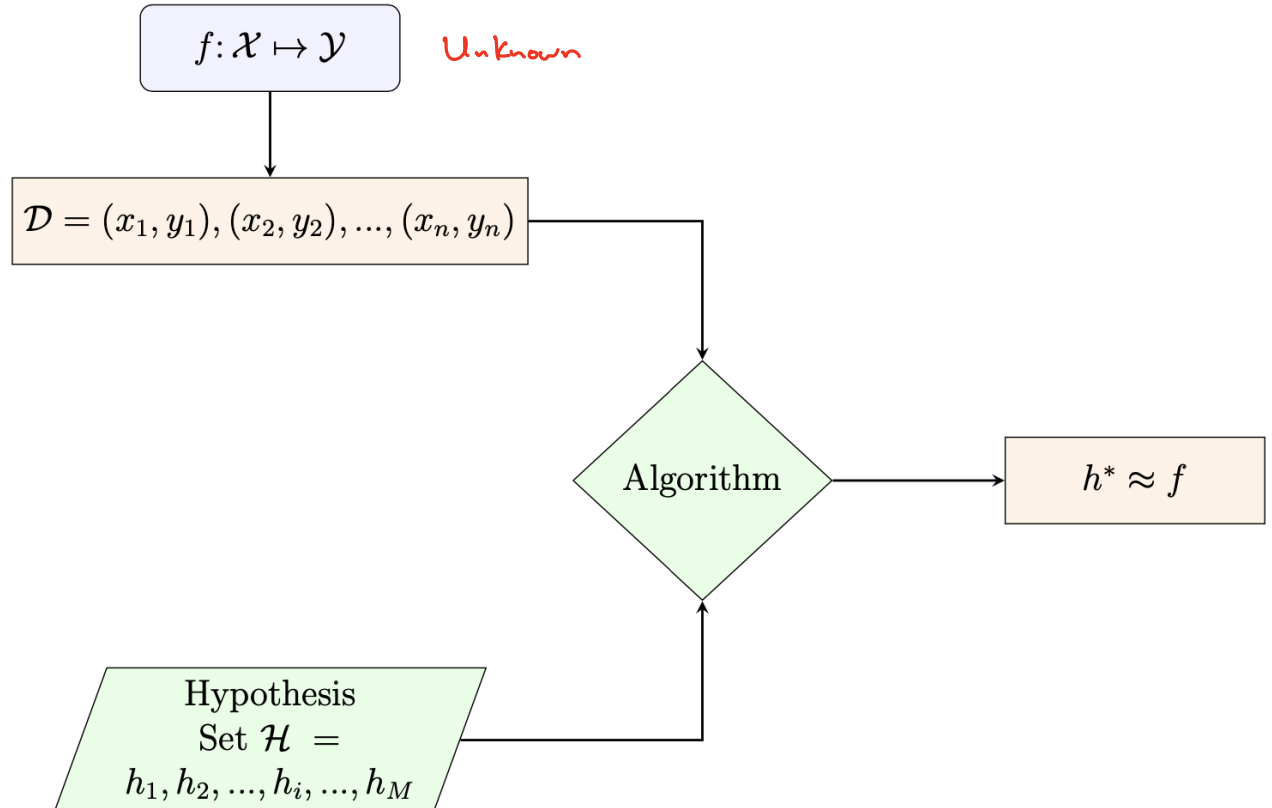
Figure 2: * Learning approach

Notation:

Supervised learning
Unsupervised learning
Reinforcement learning

Harder (with a red arrow pointing from 'Supervised learning' to 'Reinforcement learning')

Input : x
 Output : y
 classification : $y \in \mathbb{Z}$ *integer*
 regression : $y \in \mathbb{R}$ *real number*
 Unknown : $f: \mathcal{X} \mapsto \mathcal{Y}$ *target function*
 Data : $\mathcal{D} = (x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
 Hypothesis : $h^* : \mathcal{X} \mapsto \mathcal{Y}$



Linear model

n : number of samples
 d : number of features

$$x \in \mathbb{R}^d \text{ or } \mathbb{Z}^d$$

Example: Perceptron Model Let $x = (x_1, \dots, x_d)$ be 'attributes of a customer'


$$\begin{aligned} \text{approve credit if } \sum_{i=1}^d w_i x_i &> \text{threshold} \\ \text{deny } \sum_{i=1}^d w_i x_i &< \text{threshold} \end{aligned}$$

* Introduce (artificially) $x_0 = 1$, and combine

$$w_0 = b = \text{threshold}$$


$$x_0 = 1$$

$$h(x) = \text{sign} \left(\sum_{i=0}^d w_i x_i \right)$$



Learning algorithm:

- takes the training set $(x_1, y_1), \dots, (x_N, y_N)$
- picks a **misclassified point**, i.e., $\text{sign}(w^T x_n) \neq y_n$
- updates the weight vector: $w \leftarrow w + y_n x_n$



Interesting fact: it works as long as the data is linearly separable.




Types of Learning

Proposition: "using a set of observations to uncover an underlying process"


Big promise - large territory, many types...An instance - supervised learning

Learning Challenge




Example 1*:



All have shade in upper left square



→ Fits both criteria



All are symmetric

Example 2*:

x	y	h^*	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8
0 0 0	0	0	0	0	0	0	0	0	0	0
0 0 1	1	1	1	1	1	1	1	1	1	1
0 1 0	1	1	1	1	1	1	1	1	1	1
0 1 1	0	0	0	0	0	0	0	0	0	0
1 0 0	1	1	1	1	1	1	1	1	1	1
1 0 1		?	0	0	0	0	1	1	1	1
1 1 0		?	0	0	1	1	0	0	1	1
1 1 1		?	0	1	0	1	0	1	0	1

How well do you generalize to unseen examples

If we remain true to the notion of unknown target function, we cannot exclude any f_1, \dots, f_8 from being the true f . It is easy to show that any 3 bits that replace the question marks are as good as any other 3 bits.

Some thoughts:

- Learning vs. Memorizing

- The purpose of learning f is to be able to predict the value of f on points that we haven't seen before.
- The quality of the learning is going to be determined by how close our prediction is to the true value.

I References

1. Chapter 1, Elements of Statistical Learning (by Hastie, Tibshirani, and Friedman)
2. Chapter 1, Learning From Data, by Abu-Mostafa, Magdon-Ismail, Lin. * Example 1&2, and this lecture's figures are taken from this textbook.

https://web.stanford.edu/~hastie/ElemStatLearn/printings/ESLII_print12_toc.pdf