

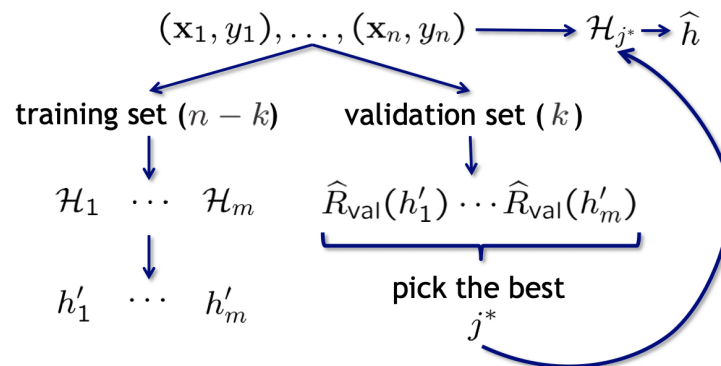
# ECE 0402 - Pattern Recognition

## Lecture 14

### Recap:

We have talked about model selection dilemma:


- we need to select appropriate values for the free parameters
- these free parameters usually control the balance between underfitting and overfitting
- but all we have training data to select parameters
- they were left “free” precisely because we don’t want to let the training data influence their selection, as this almost always leads to overfitting.
- **Using validation for model selection:** Suppose we have  $m$  models  $\mathcal{H}_1, \dots, \mathcal{H}_m$ :
  - Take the training data ( $n$  samples)
  - Split it up into two pieces:
    - \* training set:  $n - k$  samples
    - \* validation (holdout) set:  $k$  samples
    - \* train  $m$  models using the  $n - k$  samples :  $h'_1, \dots, h'_m$ .
    - \* using  $k$  samples validation set, get an estimate of risk:  $\hat{R}_{val}(h'_1), \dots, \hat{R}_{val}(h'_m)$
    - \* pick the model that gives the best result, lowest  $\hat{R}_{val}(h'_{j^*}) \Rightarrow j^*$ .
    - \* after picking the model  $h_{j^*}$ , re-train this on the entire data set ( $n$ -samples) to output  $\hat{h}$ .



We also talked about more complicated ideas like “**cross-validation**”. Try a different split each time and repeat the whole process (the one depicted in the figure) again, and finally average all of the empirical estimates  $\hat{R}_{val}$ .

- Leave one out: Train  $n$  times on  $n - 1$  points each. This is ideal, but computationally demanding
- $k - fold$  cross validation: Train  $k$  times on  $n - \frac{n}{k}$  points each

Notation clarification: in here  $k$  is the number of folds  $k' = \frac{n}{k}$  is the size of the validation set. Iterate over all 5 choices of validation set, and average. **Remarks:**



- For  $k - fold$  cross validation, the estimate depends on the particular choice of partition.
- It is common to form several estimates based on different partitions and then average them
- One thing to keep in mind: when using  $k - fold$  cross validation for classification, you should ensure that each of the sets  $\mathcal{D}_j$  contain training data from each class in the same proportion as in the full data set

In practice, cross-validation is a common choice to pick these free parameters in an automatic way. But there are other strategies. I wanna mention just one of them here briefly because it can be useful when you have smaller datasets.

### The bootstrap:

- What else can you do when your training set is really small?
- You really need as much training data as possible to get reasonable result
- Fix  $B \geq 1$
- For  $b = 1, \dots, B$  let  $\mathcal{D}_b$  be a subset of size  $n$  obtained by sampling with replacement from the full data set  $\mathcal{D}$ .

**Example:**  $n = 5$

$$\begin{aligned}\mathcal{D}_1 &= (x_4, y_4), (x_3, y_3), (x_5, y_5), (x_4, y_4), (x_1, y_1) \\ \mathcal{D}_2 &= (x_1, y_1), (x_2, y_2), (x_5, y_5), (x_5, y_5), (x_2, y_2) \\ &\vdots\end{aligned}$$

- Define  $h_b :=$  model learned base on the data  $\mathcal{D}_b$
- And define  $\mathcal{D}_b^c := \mathcal{D} \setminus \mathcal{D}_b$ 
  - In the example this could of been for  $\mathcal{D}_1, (x_2, y_2)$
  - for  $\mathcal{D}_1, (x_3, y_3)$  and  $(x_4, y_4)$
- Set  $e_b = \frac{1}{|\mathcal{D}_b^c|} = \sum_{(x_i, y_i) \in \mathcal{D}_b^c} 1_{\{h_b(x_i) \neq y_i\}}$

The **bootstrap error estimate** is given by

$$\hat{R}_B := \frac{1}{B} \sum_{b=1}^B e_b$$

Typically, in practice,  $B$  must be large, several hundred (say  $B \approx 200$ ) (or more if you can afford it) for the estimate to be accurate. It is very computationally demanding, but also you can get nice confidence intervals in addition to  $\hat{R}_B$ .

- The bootstrap error estimate  $\hat{R}_B$  tends to be pessimistic, so it is common to combine the training and bootstrap error estimates.
  - A common choice is the “**0.632 bootstrap estimate**”. This is really it’s name and these weights are actually derived from theory (not a rule of thumb, even though it looks totally like one). And I am sure you can find 5k+ papers on bootstrap error estimates...

$$0.632 \hat{R}_B + 0.368 \hat{R}_{train}$$

- The “balanced” bootstrap chooses  $\mathcal{D}_1, \dots, \mathcal{D}_B$  such that each input-output pair appears exactly  $B$  times.

**Linear Methods for Supervised Learning** In terms of linear techniques, so far, we have talked about:

- LDA
- Logistic Regression
- Naive Bayes
- PLA
- Maximum margin hyperplanes
- Soft-margin hyperplanes

- Least squares regression
- Ridge regression

Sometimes linear methods (in both regression and classification) just don't work. One way to create "nonlinear estimators" or classifiers is to first transform the data via a nonlinear feature map  $\phi : \mathbb{R} \rightarrow \mathbb{R}^p$ .


After applying  $\phi$ , we can then try applying a linear method to transformed data  $\phi(x_1), \dots, \phi(x_n)$ . We actually talked about couple of examples of this already. For example, in the case of regression, our model becomes,  $f(x) = \beta^T \phi(x) + \beta_0$  where  $\beta \in \mathbb{R}^p$ . Fitting linear models to nonlinear features...

**Example:** Suppose  $d = 1$  but  $f(x)$  is a cubic polynomial, how do we find a least squares estimate of  $f$  from training data?

$$\phi_k(x) = x^k \implies A = \begin{bmatrix} 1 & x_1 & x_1^2 & x_1^3 \\ 1 & x_2 & x_2^2 & x_2^3 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & x_n & x_n^2 & x_n^3 \end{bmatrix}$$

And we have talked about classification long time ago too...

Pictured data set is not linearly separable but can be in a higher dimension with a transform:



This dataset is linearly separable after applying such transformation with  $w = [-1, 0, 0, 1, 1]^T$

## Potential Problems with nonlinear feature maps

Suppose we transform our data via

$$x = \begin{bmatrix} x(1) \\ \vdots \\ x(d) \end{bmatrix} \xrightarrow{\phi} \Phi(x) = \begin{bmatrix} \Phi^{(1)}(x) \\ \vdots \\ \Phi^{(p)}(x) \end{bmatrix}$$

where typically  $p \gg d$ . The challenges that could arise:

- If  $p \gtrsim n$ , then this can lead to an ill-conditioned problem

- can be mitigated via regularization
- in addition, when  $p$  is very large, this can cause an increase in computational burden
  - \* (e.g., inverting  $p \times p$  matrix)

## The “kernel trick”

Shortly, do the mapping without the cost.

There is a clever way to get around this computational challenge by exploiting two facts:

- Many ML algorithms only involve the data through inner products.
- For many interesting feature maps  $\phi$ , the function

$$k(x, x') := \langle \phi(x), \phi'(x) \rangle$$

has a simple, closed form expression that can be evaluated **without explicitly calculating**  $\phi(x)$  and  $\phi'(x)$ .

**Example:** Quadratic kernel  $k(u, v) = (u^T v)^2$

Suppose  $d = 2$  and try to think backwards: “what was the feature map that this kernel is computing an inner product for?”

$$\begin{aligned}
 k(\mathbf{u}, \mathbf{v}) &= (\mathbf{u}^T \mathbf{v})^2 \\
 &= \left( [u(1) \ u(2)] \begin{bmatrix} v(1) \\ v(2) \end{bmatrix} \right)^2 \\
 &= (u(1)v(1) + u(2)v(2))^2 \\
 &= u(1)^2 v(1)^2 + 2u(1)u(2)v(1)v(2) + u(2)^2 v(2)^2 \\
 &= \langle \phi(\mathbf{u}), \phi(\mathbf{v}) \rangle
 \end{aligned}$$

$$\phi(\mathbf{u}) = \begin{bmatrix} u(1)^2 \\ \sqrt{2}u(1)u(2) \\ u(2)^2 \end{bmatrix}$$

Now suppose  $d$  is arbitrary, and

$$\begin{aligned}
k(\mathbf{u}, \mathbf{v}) &= (\mathbf{u}^T \mathbf{v})^2 \\
&= \left( \sum_{i=1}^d u(i)v(i) \right)^2 \\
&= \left( \sum_{i=1}^d u(i)v(i) \right) \left( \sum_{j=1}^d u(j)v(j) \right) \\
&= \sum_{i=1}^d \sum_{j=1}^d u(i)v(i)u(j)v(j)
\end{aligned}$$

What is the feature mapping in this case? And what is the dimension  $p$  of the corresponding feature space?

$$\begin{aligned}
&\sum_{i=1}^d \sum_{j=1}^d u(i)v(i)u(j)v(j) \stackrel{?}{=} \langle \phi(\mathbf{u}), \phi(\mathbf{v}) \rangle \\
\phi(\mathbf{u}) &= \left[ u(1)^2, \dots, u(d)^2, \dots, \sqrt{2}u(1)u(2), \dots, \sqrt{2}u(d-1)u(d) \right]^T
\end{aligned}$$

$$\therefore p = d + \frac{d(d-1)}{2}$$

So we can do this for quadratic. what happens if we have cubic kernel?

**Example:** Cubic kernel  $k(u, v) = (u^T v)^3$

In  $\mathbb{R}^2$ , we can work this out, it is not that much worse, but boring...

$$k(\mathbf{u}, \mathbf{v}) = (u(1)v(1) + u(2)v(2) + u(3)v(3))^3 \quad (1)$$

$$= u(1)^3 v(1)^3 + 3u(1)^2 u(2)v(1)^2 v(2) + 3u(1)u(2)^2 v(1)v(2)^2 + u(2)^3 v(2)^3 \quad (2)$$

$$= \sum_{i=0}^3 \binom{3}{i} u(1)^{3-i} u(2)^i \cdot v(1)^{3-i} v(2)^i \quad (3)$$

$$\phi(\mathbf{u}) = \left[ u(1)^3, \sqrt{3}u(1)^2 u(2), \sqrt{3}u(1)u(2)^2, u(2)^3 \right]^T$$

and we could do this in  $d$  dimensions if we felt like it...But it is probably okay, if we don't...

So let me state what happens in general,

## Polynomial Kernels:

$$\begin{aligned}k(\mathbf{u}, \mathbf{v}) &= (\mathbf{u}^T \mathbf{v})^m \\&= \sum_{\substack{\text{partitions} \\ j_1, \dots, j_d}} \binom{m}{j_1, \dots, j_d} u(1)^{j_1} v(1)^{j_1} \dots u(d)^{j_d} v(d)^{j_d} \\ \phi(\mathbf{u}) &= \left[ \dots, \sqrt{\binom{m}{j_1, \dots, j_d}} u(1)^{j_1} \dots u(d)^{j_d}, \dots \right]^T\end{aligned}$$

In English, all possible monomials of degree  $m$ .

This is one example of a kernel. There are lots of different kernels we could use. The general kind of kernel we will use in this class is called **inner product kernel**.

**Definition:** A inner product kernel is a mapping

$$k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$$

for which there exists an inner product space  $\mathcal{F}$  and a mapping  $\phi : \mathbb{R}^d \rightarrow \mathcal{F}$  such that

$$k(u, v) = \langle \phi(\mathbf{u}), \phi(\mathbf{v}) \rangle_{\mathcal{F}}$$

for all  $u, v \in \mathbb{R}^d$ .

Given a function  $k(u, v)$ , how can we tell when it is an inner product kernel?

- Mercer's theorem
- Positive semidefinite property

## Positive semidefinite kernels

We say that  $k(u, v)$  is a positive semidefinite kernel

- if  $k$  is symmetric.
- for all  $n$  and all  $x_1, \dots, x_n \in \mathbb{R}^d$ , the **Gram matrix**  $\mathbf{K}$  is defined by

$$K(i, j) = k(x_i, x_j)$$

is positive semidefinite, i.e.,  $x^T K x \geq 0$  for all  $x$ .

**Theorem:**  $k$  is an inner product kernel iff  $k$  is a positive semidefinite kernel.

### Examples:

- Homogeneous polynomial kernel

$$k(u, v) = (u^T v)^m \quad m = 1, 2, 3, \dots$$

- Inhomogeneous polynomial kernel

$$k(u, v) = (u^T v + c)^m \quad m = 1, 2, 3, \dots$$
$$c > 0$$

$\phi$  maps to the set of all monomials of degree  $\leq m$ .

- Gaussian/Radial basis function (RBF) kernel:

$$k(u, v) = (2\pi\sigma^2)^{-d/2} \exp\left(-\frac{\|u - v\|^2}{2\sigma^2}\right)$$

You can just do (it doesn't matter if it is a PDF or not):

$$k(u, v) = \exp\left(-\frac{\|u - v\|^2}{2\sigma^2}\right)$$

One can show that  $k$  is a positive semidefinite kernel. but what is  $\mathcal{F}$ ?

–  $\mathcal{F}$  is **infinite dimensional**!