

ECE 1390/2390

## Image Processing and Computer Vision – Fall 2021

---

### Lecture 4: Linear Filtering – Edge detection

**Ahmed Dallal**

Assistant Professor of ECE

University of Pittsburgh

### Reading

- FP 5.1, 5.2




## Filters for features

- Previously, thinking of filtering as a way to remove or reduce **noise**
- Now, consider how filters will allow us to abstract higher-level **“features”**.
  - Map raw pixels to an intermediate representation that will be used for subsequent processing
  - Goal: reduce amount of data, discard redundancy, preserve what's useful



K. Grauman


## Reduced images



## Reduced images




*Edges seem to be important...*



---

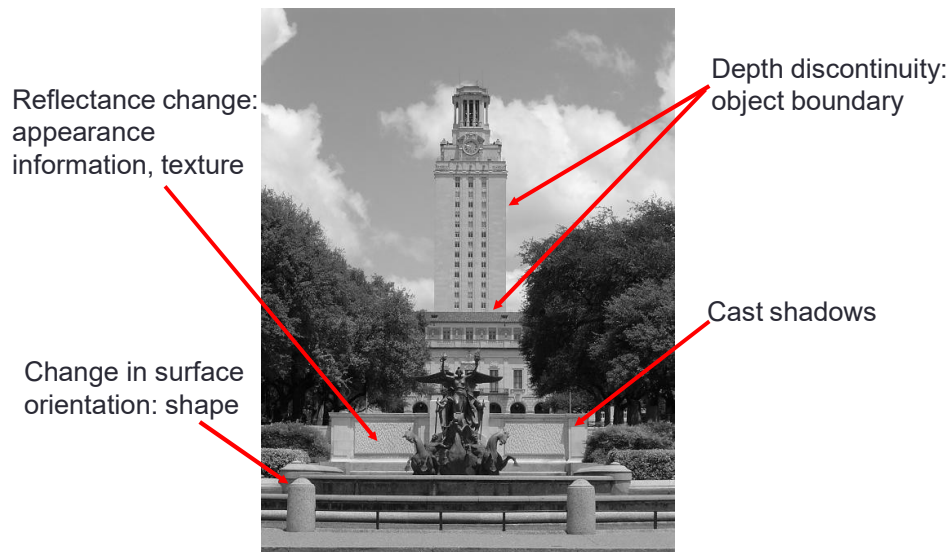
*Edge detection: Gradients*

## Origin of Edges





- Edges are caused by a variety of factors
- Information theory view: edges encode change, change is what is hard to predict, therefore edges efficiently encode an image

## In a real image





## Edge detection




- Convert a 2D image into a set of curves
  - Extracts salient features of the scene
  - More compact than pixels

## Edge detection



- How can you tell that a pixel is on an edge?

Recall images as functions...



*Edges look like steep cliffs → very large change in gray level*

## Edge Detection

Basic idea: look for a neighborhood with strong signs of change.


Problems:

- neighborhood size
- how to detect change

81	82	26	24
82	33	25	25
81	82	26	24

## Derivatives and edges

An edge is a place of rapid change in the image intensity function.



## Differential Operators


- Differential operators –when applied to the image returns some derivatives.
- Model these “operators” as masks/kernels that compute the image **gradient function**.
- Threshold the this gradient function to select the edge pixels.
- Which brings us to the question:

## What's a gradient?

### Image gradient

The gradient of an image:

$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$



The gradient points in the direction of most rapid increase in intensity

The gradient direction is given by:

$$\theta = \tan^{-1} \left( \frac{\partial f / \partial y}{\partial f / \partial x} \right)$$

- how does this relate to the direction of the edge?

The *edge strength* is given by the gradient magnitude  $\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$



## Discrete gradient


- For 2D function,  $f(x,y)$ , the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

- For discrete data, we can approximate using finite differences:

$$\begin{aligned} \frac{\partial f(x, y)}{\partial x} &\approx \frac{f(x+1, y) - f(x, y)}{1} \\ &\approx f(x+1, y) - f(x, y) \quad \text{"right derivative"} \end{aligned}$$

## Finite differences



Source: D.A. Forsyth

## Differentiation and convolution

- For 2D function,  $f(x,y)$ , the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

- For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

- To implement above as correlation, what would be the associated filter?

## Partial derivatives of an image



Which shows changes with respect to x?  
(showing correlation filters)

## The discrete gradient

- We want an “operator” (mask/kernel) that we can apply to the image that implements:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

How would you implement this as a cross-correlation?  
(not flipped)

0	0
-1	+1
0	0

$H$

Not symmetric  
around image  
point; which is  
“middle” pixel?

0	0	0
-1/2	0	+1/2
0	0	0

$H$

Average of  
“left” and  
“right”  
derivative.  
See?

## Example: Sobel operator

$\frac{1}{8}$	-1	0	1
	-2	0	2
	-1	0	1

$S_x$

$\frac{1}{8}$	1	2	1
	0	0	0
	-1	-2	-1

$S_y$

On a pixel of the image  $I$

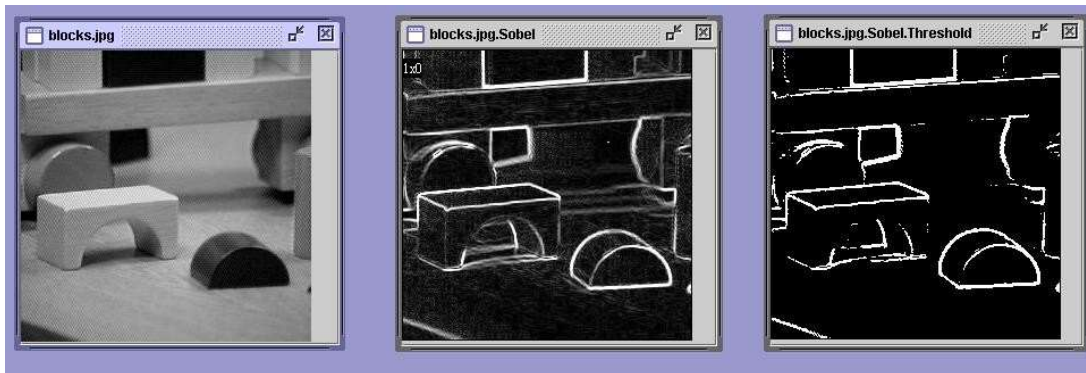
- Let  $g_x$  be the response to mask  $S_x$  (sometimes  $\ast 1/8$ )
- Let  $g_y$  be the response to mask  $S_y$

What is the gradient?

(Sobel) Gradient is  $\nabla I = [g_x \ g_y]^T$

$g = (g_x^2 + g_y^2)^{1/2}$  is the gradient magnitude.  
 $\theta = \text{atan2}(g_y, g_x)$  is the gradient direction.

## Sobel Operator on Blocks Image



original image

gradient  
magnitudethresholded  
gradient  
magnitude

## Some Well-Known Gradients Masks

	S <sub>x</sub>	S <sub>y</sub>																		
• Sobel:	<table><tr><td>-1</td><td>0</td><td>1</td></tr><tr><td>-2</td><td>0</td><td>2</td></tr><tr><td>-1</td><td>0</td><td>1</td></tr></table>	-1	0	1	-2	0	2	-1	0	1	<table><tr><td>1</td><td>2</td><td>1</td></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>-1</td><td>-2</td><td>-1</td></tr></table>	1	2	1	0	0	0	-1	-2	-1
-1	0	1																		
-2	0	2																		
-1	0	1																		
1	2	1																		
0	0	0																		
-1	-2	-1																		
• Prewitt:	<table><tr><td>-1</td><td>0</td><td>1</td></tr><tr><td>-1</td><td>0</td><td>1</td></tr><tr><td>-1</td><td>0</td><td>1</td></tr></table>	-1	0	1	-1	0	1	-1	0	1	<table><tr><td>1</td><td>1</td><td>1</td></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>-1</td><td>-1</td><td>-1</td></tr></table>	1	1	1	0	0	0	-1	-1	-1
-1	0	1																		
-1	0	1																		
-1	0	1																		
1	1	1																		
0	0	0																		
-1	-1	-1																		
• Roberts	<table><tr><td>0</td><td>1</td></tr><tr><td>-1</td><td>0</td></tr></table>	0	1	-1	0	<table><tr><td>1</td><td>0</td></tr><tr><td>0</td><td>-1</td></tr></table>	1	0	0	-1										
0	1																			
-1	0																			
1	0																			
0	-1																			


## MATLAB does edges

```
filt = fspecial('sobel')
```

```
filt =
```

```
    1    2    1
    0    0    0
   -1   -2   -1
```

```
outim = imfilter(double(im),filt);
imagesc(outim);
colormap gray;
```




## Quiz

It is better to compute gradients using:


- a) **Convolution** since that's the right way to model filtering so you don't get flipped results.
- b) **Correlation** because it's easier to know which way the derivatives are being computed.
- c) Doesn't matter.
- d) Neither since I can just write a for-loop to compute the derivatives.

But...

- Consider a single row or column of the image
  - Plotting intensity as a function of  $x$




- Apply derivative operator....



Uh, where's the edge?

## Finite differences responding to noise




Increasing noise


(this is zero mean additive Gaussian noise)

Source: D. Forsyth


## Solution: smooth first




## Solution: smooth first



Solution: smoothfirst




Solution: smoothfirst






Solution: smoothfirst



## Derivative theorem of convolution


- This saves us one operation:  $\frac{\partial}{\partial x}(h * f) = (\frac{\partial}{\partial x}h) * f$




How can we find (local) maxima of a function?

## 2<sup>nd</sup> derivative of Gaussian


- Consider  $\frac{\partial^2}{\partial x^2}(h * f)$




$f$



$\frac{\partial^2}{\partial x^2} h$



$\frac{\partial^2}{\partial x^2}(h * f)$



Where is the edge? Zero-crossings of bottom graph

---


*Edge detection: 2D operators*

## Derivative of Gaussian filter – 2D

$$(I \otimes g) \otimes h_x = I \otimes (g \otimes h_x)$$

## Derivative of Gaussian filter – 2D


$$(I \otimes g) \otimes h_x = I \otimes (g \otimes h_x)$$




$$\begin{bmatrix} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix} \otimes \begin{bmatrix} -1 & 1 \end{bmatrix} =$$

## Derivative of Gaussian filter – 2D

$$(I \otimes g) \otimes h = I \otimes (g \otimes h)$$



$$\begin{bmatrix} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix} \otimes \begin{bmatrix} -1 & 1 \end{bmatrix} =$$




Is this preferable?

## Quiz



Why is it preferable to apply  $h$  to the smoothing function  $g$  and apply the result to the Image.

- a) It's not – they are mathematically equivalent.
- b) Since  $h$  is typically smaller we take fewer derivatives so it's faster.
- c) The smoothed derivative operator is computed once and you have it to use repeatedly.
- d) B & C


## Derivative of Gaussian filter



x-direction



y-direction






*Correlation or convolution?*

*And for y it's always  
a problem!*


Source: S. Lazebnik

## Smoothing with a Gaussian




```
for sigma=1:3:10
    h = fspecial('gaussian', fsize, sigma);
    out = imfilter(im, h);
    imshow(out);
    pause;
end
```

## Smoothing with a Gaussian



Recall: parameter  $\sigma$  is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.

## Effect of $\sigma$ on derivatives



The apparent structures differ depending on Gaussian's scale parameter.

**Larger values:** larger scale edges detected

**Smaller values:** finer features detected

## Gradients -> edges

### Primary edge detection steps:

- 1. Smoothing: suppress noise
- 2. Edge “enhancement”: filter for contrast
- 3. Edge localization (“Thin”)
  - Determine which local maxima from filter output are actually edges vs. noise
    - Threshold, Thin → get a single contour
  - We may need linking to connect edge pixels.



## *Canny* edgedetector

1. Filter image with derivative of Gaussian
2. Find magnitude and orientation of gradient
3. Non-maximum suppression:
  - Thin multi-pixel wide “ridges” down to single pixel width

## *Canny* edgedetector

### 4. Linking and thresholding (hysteresis):

- Define two thresholds: low and high
- Use the high threshold to start edge curves and the low threshold to continue them

MATLAB: `edge(image, 'canny');`

`>>doc edge` (or *help edge* if doc is not supported)

Source: D. Lowe, L. Fei-Fei

## The Canny edgedetector



original image (Lena)



## The Canny edgedetector



magnitude of the gradient

## The Canny edgedetector



thresholding

## The Canny edge detector



Problem: pixels  
along this edge  
didn't survive  
the  
thresholding



## The Canny edge detector



## The Canny edgedetector



## Canny: Non-maximal suppression



- Check if pixel is local maximum along gradient direction
  - can require checking interpolated pixels p and r

## The Canny edge detector



Problem:  
pixels along this  
edge didn't  
survive the  
thresholding

thinning  
(non-maximum suppression)

## Canny threshold hysteresis

1. Apply a high threshold to detect strong edge pixels.
2. Link those strong edge pixels to form strong edges.
3. Apply a low threshold to find weak but plausible edge pixels.
4. Extend the strong edges to follow weak edge pixels.

## Result of Canny



## Effect of $\sigma$ (Gaussian kernel spread/size)



original



Canny with  $\sigma=1$



Canny with  $\sigma=2$

- Large  $\sigma$  detects large scale edges
- Small  $\sigma$  detects fine features

*The choice of  $\sigma$  depends on desired behavior*

## So, what scale to choose?

It depends what we're looking for.



Too fine of a scale...can't see the forest for the trees.

Too coarse of a scale...can't tell the maple grain from the cherry.

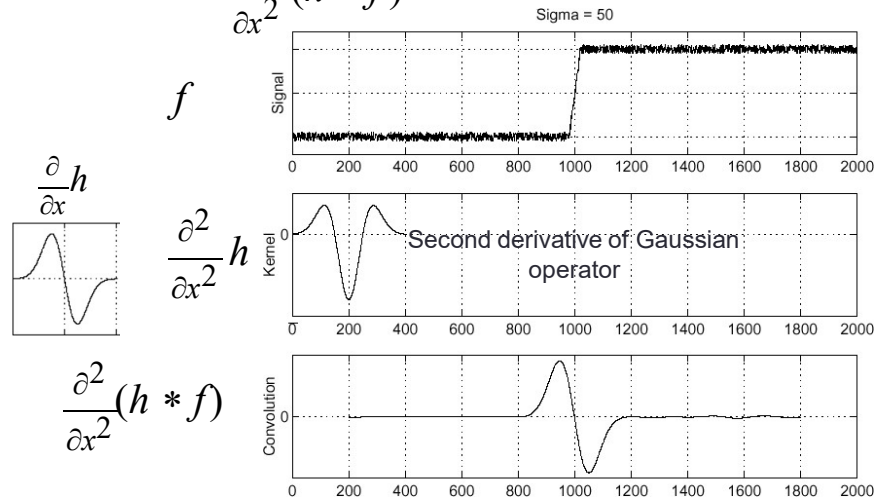
## Quiz

The Canny edge operator is probably quite sensitive to noise.

- a) True – derivatives accentuate noise
- b) False – the gradient is computed using a derivative of Gaussian operator which removes noise.
- c) Mostly false – it depends upon the  $\sigma$  chose.

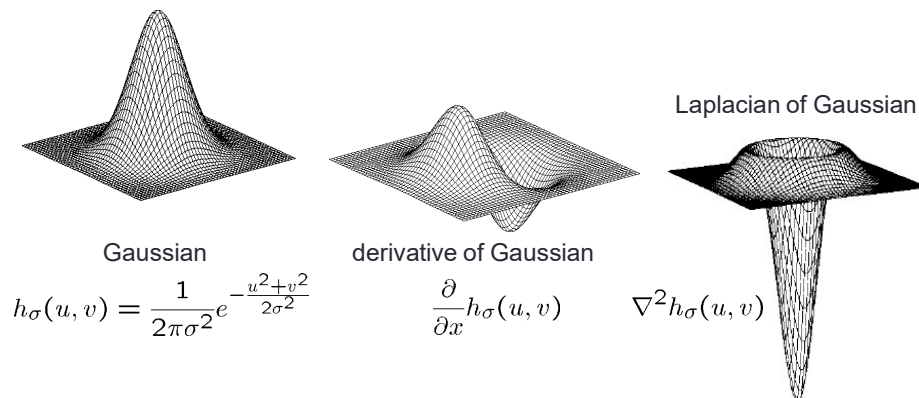
## Recall: 2<sup>nd</sup> derivative of Gaussian (1-D)

- Consider  $\frac{\partial^2}{\partial x^2}(h * f)$



Zero-crossings of bottom graph are edges

## Single 2D edge detection filter



$\nabla^2$  is the **Laplacian** operator. Edges at zero crossings:

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

## Edgedemo

```
% Edge Demo
pkg load image; % Octave only

%% Read Lena image
lena = imread('lena.png');
figure, imshow(lena), title('Original image, color');

%% Convert to monochrome (grayscale) using rgb2gray
lenaMono = rgb2gray(lena);
figure, imshow(lenaMono), title('Original image, monochrome');

%% Make a blurred/smoothed version
h = fspecial('gaussian', [11 11], 4);
figure, surf(h);
lenaSmooth = imfilter(lenaMono, h);
figure, imshow(lenaSmooth), title('Smoothed image');
```

## Edgedemo (contd.)

```
% Method 1: Shift left and right, and show diff image
lenaL = lenaSmooth;
lenaL(:, [1:(end - 1)]) = lenaL(:, [2:end]);
lenaR = lenaSmooth;
lenaR(:, [2:(end)]) = lenaR(:, [1:(end - 1)]);
lenaDiff = double(lenaR) - double(lenaL);
figure, imshow(lenaDiff, []), title('Difference between right and left shifted images');

%% Method 2: Canny edge detector
cannyEdges = edge(lenaMono, 'canny'); % on original mono image
figure, imshow(cannyEdges), title('Original edges');
cannyEdges = edge(lenaSmooth, 'canny'); % on smoothed image
figure, imshow(cannyEdges), title('Edges of smoothed image');
% notice how a lot of the detail features are now gone

%% Method 3: Laplacian of Gaussian
logEdges = edge(lenaMono, 'log');
figure, imshow(logEdges), title('Laplacian of Gaussian');
```



## Summary

- Hopefully you've learned filtering by convolution and correlation, taking derivatives by operators, computing gradients and using these for edge detection.
- We've also discussed filters as templates – something we'll use again later.
- Next we'll take a detour and do some “real” computer vision where we find structures in images. It will make use of the edges we discussed today.