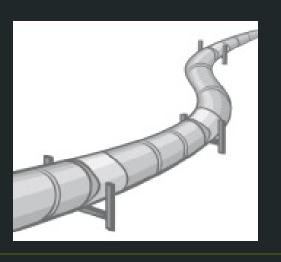
Processor Pipelining

CS/COE 1541 (Fall 2020) Wonsun Ahn



Pipelining Basics





Improving Washer / Dryer / Closet Utilization

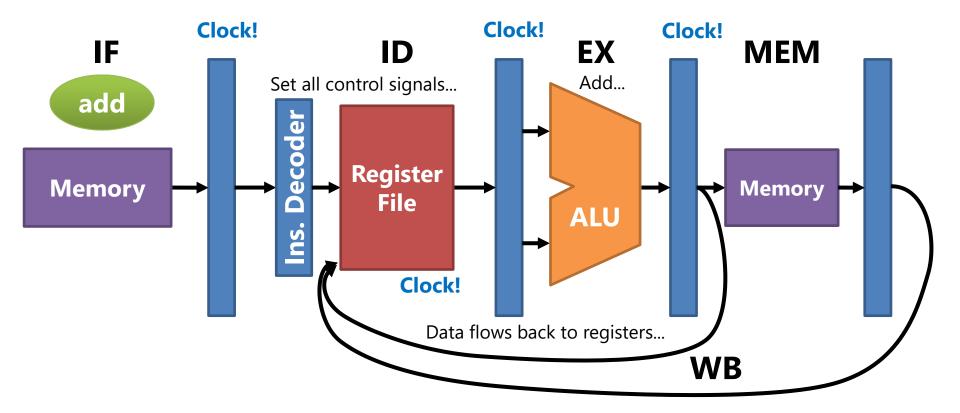
- If you work on loads of laundry one by one, you only get ~33% utilization
 If you form an "assembly line", you achieve ~100% utilization!





Multi-cycle instruction execution

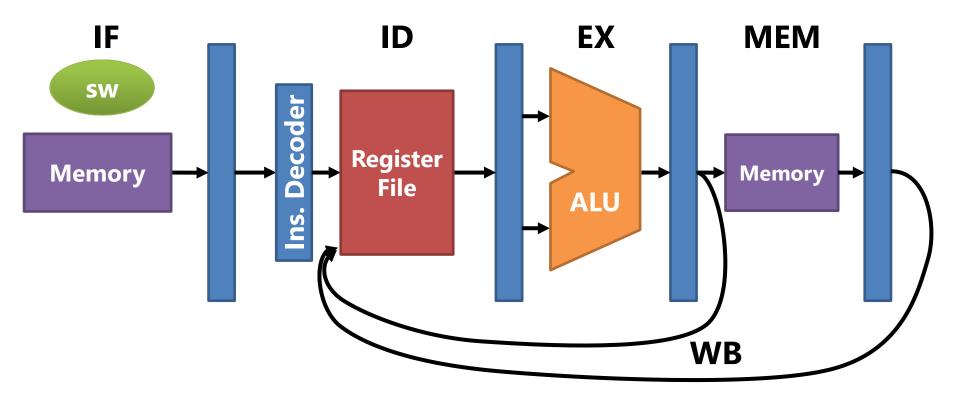
• Let's watch how an instruction flows through the datapath.





Pipelined instruction execution

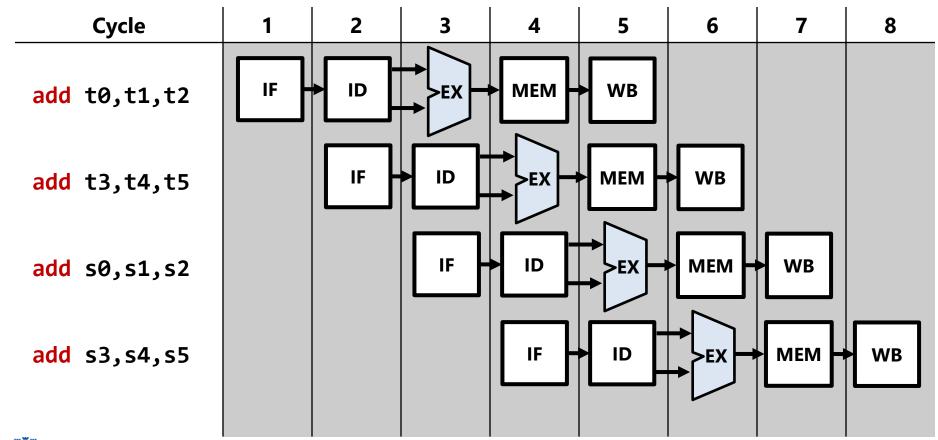
Pipelining allows one instruction to be fetched each cycle!





Pipelining Timeline

• This type of parallelism is called *pipelined parallelism*.





A Pipelined Implementation is even Faster!

- Again each instruction takes different number of cycles to complete
 - Iw takes 5 cycles: IF/ID/EX/MEM/WB
 - o add takes 4 cycles: IF/ID/EX/WB
- If each stage takes 1 ns each:
 - Iw takes 5 ns and add takes 4 ns
- Q) The average instruction execution time (given 100 instructions)?
- A) (99 ns + 5 ns) / 100 = 1.04 ns
 - Assuming last instruction is a lw (a 5-cycle instruction)
 - A ~5X speed up from single cycle!



Pipelined vs. Multi-cycle vs. Single-cycle

What happened to the three components of performance?

$$\frac{\text{instructions}}{\text{program}} \quad X \quad \frac{\text{cycles}}{\text{instruction}} \quad X \quad \frac{\text{seconds}}{\text{cycle}}$$

Architecture	Instructions	СРІ	Cycle Time (1/F)
Single-cycle	Same	1	5 ns
Multi-cycle	Same	4~5	1 ns
Pipelined	Same	1	1 ns

- Compared to single-cycle, pipelining improves clock cycle time
 - Or in other words CPU clock frequency
 - The deeper the pipeline, the higher the frequency will be

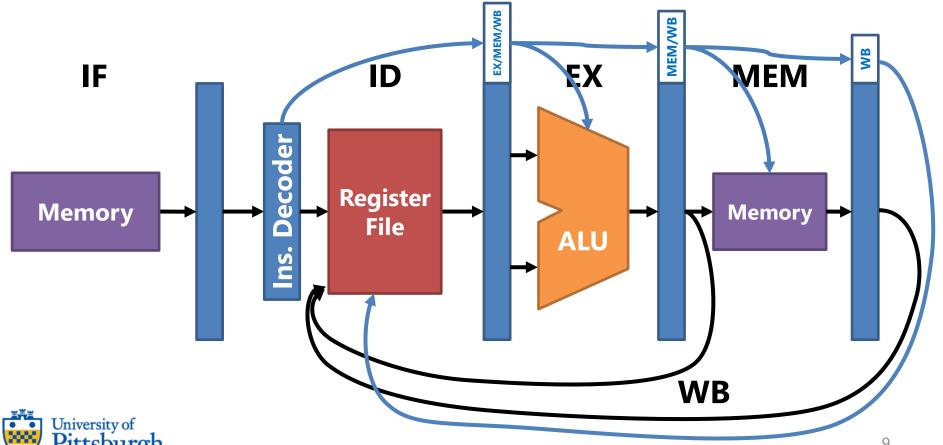
* Caveat: latch delay and unbalanced stages can increase cycle time



How about the control signals?

A new instruction is decoded at every cycle!

• Control signals must be passed along with the data at each stage



Pipeline Hazards



Pipeline Hazards

- For pipelined CPUs, we said CPI is practically 1
 - But that depends entirely on having the pipeline filled
 - o In real life, there are *hazards* that prevent 100% utilization

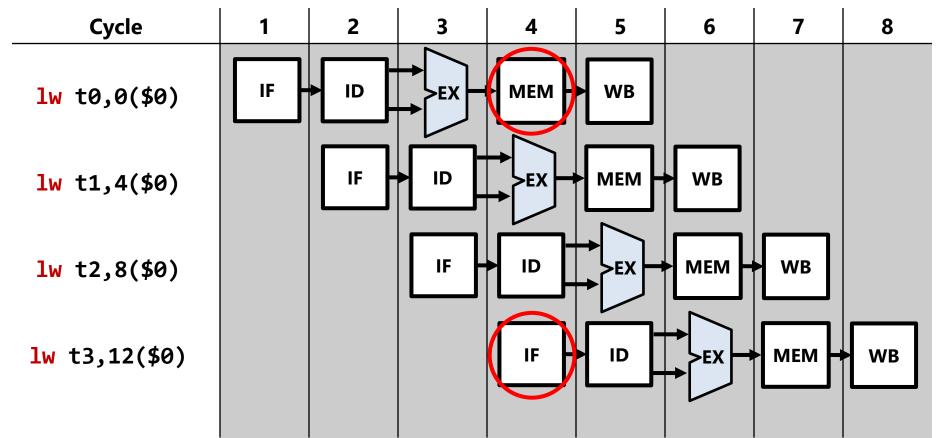
• Pipeline Hazard

- When the next instruction cannot execute in the following cycle
- o Hazards introduce **bubbles** (delays) into the pipeline timeline
- Architects have some tricks up their sleeves to avoid hazards
- But first let's briefly talk about the three types of hazards: Structural hazard, Data hazard, Control Hazard



Structural Hazards

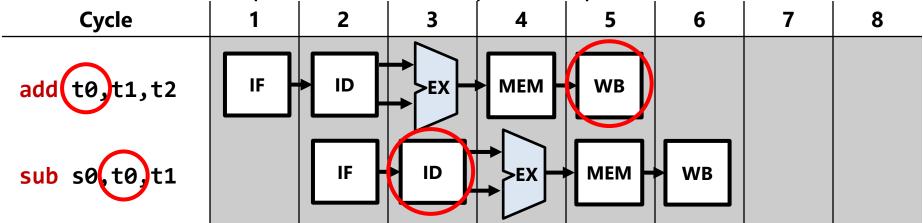
Two instructions need to use the same hardware at the same time.



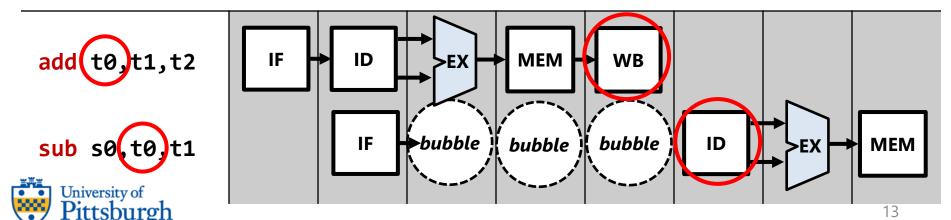


Data Hazards

An instruction depends on the output of a previous one.

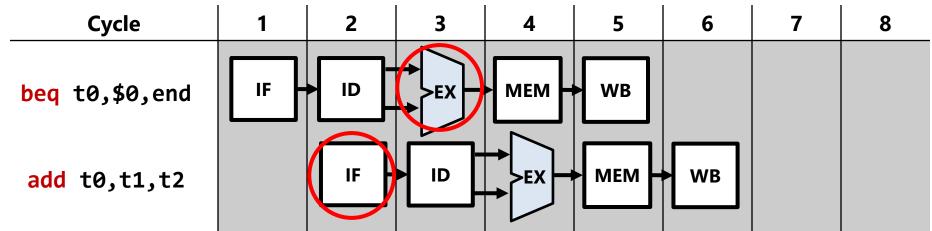


• sub must wait until add's WB phase is over before doing its ID phase

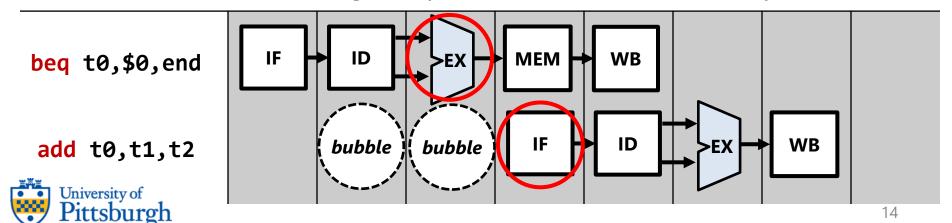


Control Hazards

You don't know the outcome of a conditional branch.



• add must wait until beq's EX phase is over before its IF phase



Dealing with Hazards

- Pipeline must be controlled so that hazards don't cause malfunction
- Who is in charge of that? You have a choice.
 - 1. Compiler can avoid hazards by inserting nops
 - Insert a nop where compiler thinks a hazard would happen
 - 2. CPU can internally avoid hazards using a *hazard detection unit*
 - If structural/data hazard, pipeline stalled until resolved
 - If control hazard, pipeline *flushed* of wrong path instructions



Compiler avoiding a data hazard

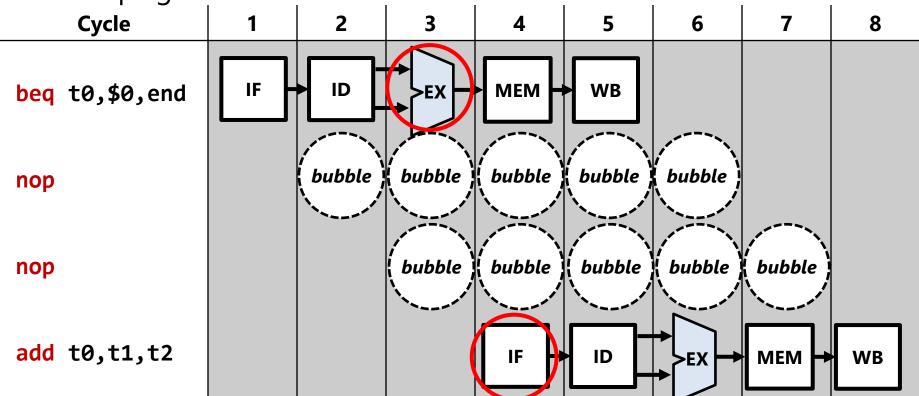
University of

 The nops flow through the pipeline not doing any work Cycle 8 add(t0,)t1,t2 IF ID MEM WB >EX bubble bubble bubble bubble bubble nop bubble bubble bubble bubble nop nop bubble bubble bubble) bubble bubble IF ID >EX **MEM** sub s0, t0, t1

16

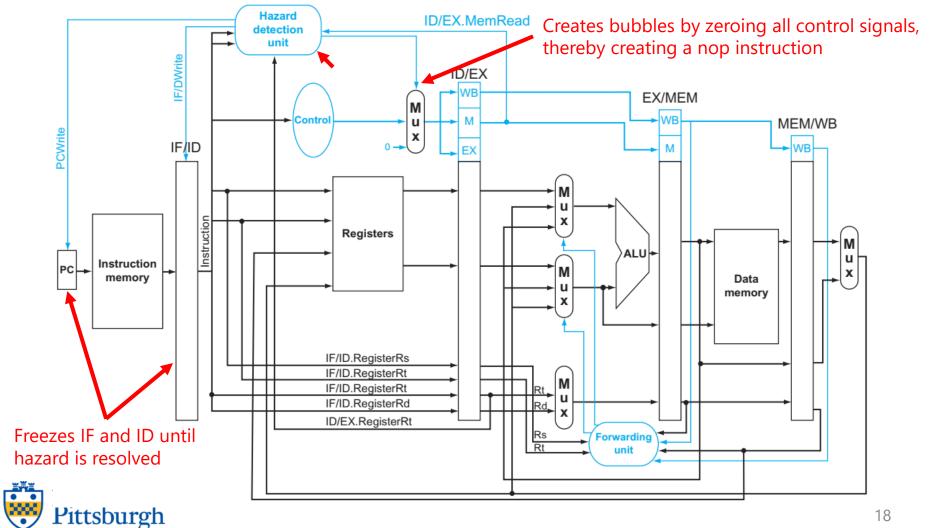
Compiler avoiding a control hazard

The nops give time for condition to resolve before instruction fetch



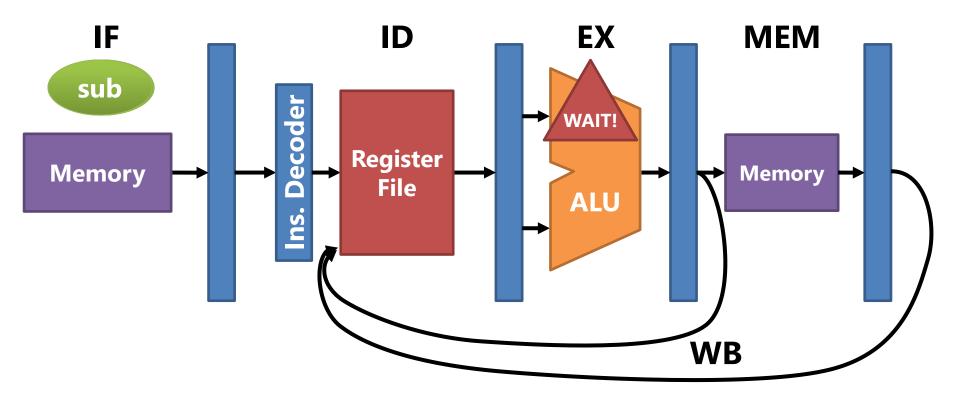


Hazard Detection Unit



Hazard Detection Unit avoiding a data hazard

• Suppose we have an **add** that depends on an **lw**.





Structural / Data Hazards cause stalls

- If HDU detects a structural or data hazard, it does the following:
 - o It **stops fetching instructions** (doesn't update the PC).
 - o It stops clocking the pipeline registers for the stalled stages.
 - The stages after the stalled instructions are filled with nops.
 - Change control signals to 0 using the mux!
 - o In this way, all following instructions will be stalled
- When structural or data hazard is resolved
 - HDU resumes instruction fetching and clocking of stalled stages
- But what about control hazards?
 - Instructions in wrong path are already in pipeline!
 - Need to *flush* these instructions



Control Hazard Example

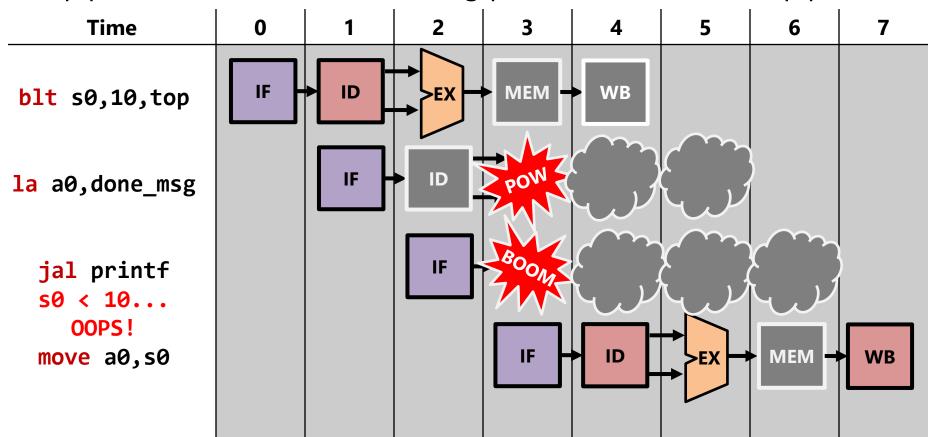
• Supposed we had this for loop followed by printf("done"):

```
for(s0 = 0 .. 10)
                                li
                                      s0, 0
    print(s0);
                            top:
                                move a0, s0
printf("done");
                                jal
                                      print
                                addi s0, s0, 1
                                blt s0, 10, top
  By the time s0, 10
  are compared at blt
  EX stage, the CPU
                                      a0, done msg
                                la
  would have already
                                      printf
                                jal
  fetched la and jal!
```



What's a flush?

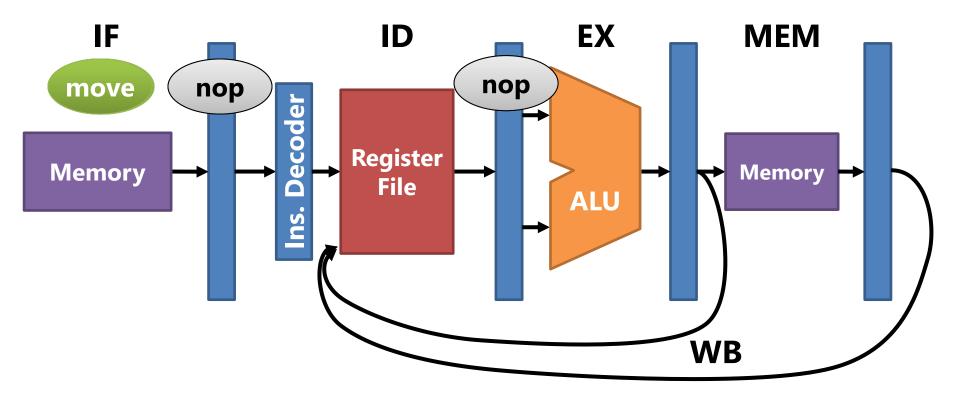
• A pipeline flush removes all wrong path instructions from pipeline





Hazard Detection Unit avoiding a control hazard

• Let's watch the previous example.





Control Hazards cause flushes

- If a control hazard is detected due to a branch instruction:
 - Any "newer" instructions (those already in the pipeline) are transformed into nops.
 - Any "older" instructions (those that came BEFORE the branch) are left alone to finish executing as normal.



Performance penalty of pipeline stalls

• Remember the three components of performance:

$$\frac{\text{instructions}}{\text{program}} \quad X \quad \frac{\text{cycles}}{\text{instruction}} \quad X \quad \frac{\text{seconds}}{\text{cycle}}$$

Architecture	Instructions	СРІ	Cycle Time (1/F)
Single-cycle	Same	1	5 ns
Ideal 5-stage pipeline	Same	1	1 ns
Pipeline w/ stalls	Same	2	1 ns

- Pipelining increases *clock frequency* proportionate to depth
- But stalls increase *CPI* (cycles per instruction)
 - \circ If stalls prevent new instructions from being fetched half the time, the CPU will have a CPI of 2 \rightarrow Only 2.5X speed up (instead of 5X)
- We'd like to avoid this penalty if possible!



Compiler nops vs. CPU Hazard Detection Unit

- Limitations of compiler nops
 - Compiler must make assumptions about processor design
 - That means processor design must become part of ISA
 - What if that design is no longer ideal in future generations?
 - Length of MEM stage is very hard to predict by the compiler
 - Until now we assumed MEM takes a uniform one cycle
 - But remember what we said about the Memory Wall?
 - MEM isn't uniform really and sometimes hundreds of cycles
- But compiler nops is very energy-efficient
 - Hazard Detection Unit can be power hungry
 - A lot of long wires controlling remote parts of the CPU
 - Adds to the **Power Wall** problem
 - Compiler scheduling via nops removes need for HDU

