

rec04

September 24, 2020

1 CS 1656 – Introduction to Data Science

1.1 Instructor: Alexandros Labrinidis, Evangelos Karageorgos

1.2 Teaching Assistant: Xiaoting Li

1.3 Additional Credits: Tahereh Arabghalizi, Agha Zuha, Anatoli Shein

1.4 SQLite in Python

In this recitation we will learn how to create SQLite Databases, create tables, populate tables, and execute SQL queries.

Start off by importing `sqlite3`, which comes installed with Anaconda's package list.

```
In [1]: import sqlite3 as lite
```

1.4.1 Introduction to SQLite

SQLite is an in-process library that implements a self-contained, serverless, zero-configuration, transactional SQL database engine. Unlike most other SQL databases, SQLite does not have a separate server process. SQLite reads and writes directly to ordinary disk files.

1.4.2 Creating and Connecting to SQLite Database

To connect to a database, use the `connect()` method which returns a connection object. If a database with that name does not exist, `connect()` method creates a database.

```
In [2]: con = lite.connect('cs1656wed.sqlite')
```

1.4.3 Create/Drop Tables & Insert Data

From the connection, we get the cursor object. The cursor is used to traverse the records from the result set. By using the `with` keyword, the Python interpreter automatically releases the resources by closing the connection, provides error handling and **commits** the changes. Otherwise, each update to the database has to be committed manually. You can think of `commit` as saving the changes.

We call the `execute()` method of the cursor to execute the SQL statements. Let's start by creating a `Rankings` table in the database.

```
In [3]: with con:
        cur = con.cursor()
        cur.execute('DROP TABLE IF EXISTS Courses')
        cur.execute("CREATE TABLE Courses(cid INT, number INT, professor TEXT, major TEXT, y

        cur.execute('DROP TABLE IF EXISTS Majors')
        cur.execute("CREATE TABLE Majors(sid INT, major TEXT)")

        cur.execute('DROP TABLE IF EXISTS Grades')
        cur.execute("CREATE TABLE Grades(sid INT, cid INT, credits INT, grade INT)")

        cur.execute('DROP TABLE IF EXISTS Students')
        cur.execute("CREATE TABLE Students(sid INT, firstName TEXT, lastName TEXT, yearStart
```

Now data can be inserted in the table using two ways. You could either insert each row one by one as shown below,

```
In [4]: import pandas
        from sqlalchemy import create_engine

        engine = create_engine("sqlite:///cs1656wed.sqlite")
        df1 = pandas.read_csv('students.csv')
        df1.to_sql('students', engine, if_exists='append', index=False)

        df2 = pandas.read_csv('grades.csv')
        df2.to_sql('grades', engine, if_exists='append', index=False)

        df3 = pandas.read_csv('courses.csv')
        df3.to_sql('courses', engine, if_exists='append', index=False)

        df4 = pandas.read_csv('majors.csv')
        df4.to_sql('majors', engine, if_exists='append', index=False)
```

Or a easier way to insert all rows together is by using executemany() method. But before we try the second method of inserting data, let's first drop the existing table and create it again.

1.4.4 Select, Where, Orderby

To select all data from the table,

```
In [5]: cur.execute("SELECT * FROM students")
```

```
Out[5]: <sqlite3.Cursor at 0x137c6868420>
```

To retrieve data after executing a SELECT statement, you can either treat the cursor as an iterator and call the cursor's fetchone() method to retrieve a single matching row, or call fetchall() to get a list of the matching rows.

```
In [6]: for row in cur.execute("select * from students"):
        print(row)
        cur.execute("select * from students")
        df5 = pandas.DataFrame(cur.fetchall(), columns=[column[0] for column in cur.description])
        df5
```

```
(555, 'Solange', 'Knowles', 2012)
(666, 'Peter', 'Weiner', 2012)
(1111, 'Ya', 'Boi', 2020)
(1234, 'Michael', 'Scott', 2009)
(1337, 'Beyonce', 'Knowles', 1985)
(1345, 'Julius', 'Caesar', -60)
(1865, 'Abraham', 'Lincoln', 2012)
(3321, 'Mark', 'Brandanowitz', 1992)
(4224, 'Michelle', 'Young', 1984)
(4444, 'Grace', 'Hopper', 1944)
(5376, 'Poverty', 'Jones', 1969)
(5432, 'Mark', 'Wahlberg', 2000)
(6969, 'Thug', 'Nugget', 1862)
(7928, 'John', 'Cash', 1950)
(9191, 'Margaret', 'Mead', 1919)
(9878, 'First', 'Last', 2014)
(9999, 'Elon', 'Musk', 1932)
(14325, 'John', 'Doe', 1999)
(69420, 'Ray', 'Zimmerman', 2017)
(90210, 'Kappa', 'Pride', 2018)
(314158, 'Mr.', 'Pie', 1000)
(999831, 'John', 'Cena', 2003)
(89990, 'BoJack', 'Horseman', 2012)
```

```
Out[6]:
```

	sid	firstName	lastName	yearStarted
0	555	Solange	Knowles	2012
1	666	Peter	Weiner	2012
2	1111	Ya	Boi	2020
3	1234	Michael	Scott	2009
4	1337	Beyonce	Knowles	1985
5	1345	Julius	Caesar	-60
6	1865	Abraham	Lincoln	2012
7	3321	Mark	Brandanowitz	1992
8	4224	Michelle	Young	1984
9	4444	Grace	Hopper	1944
10	5376	Poverty	Jones	1969
11	5432	Mark	Wahlberg	2000
12	6969	Thug	Nugget	1862
13	7928	John	Cash	1950
14	9191	Margaret	Mead	1919
15	9878	First	Last	2014

16	9999	Elon	Musk	1932
17	14325	John	Doe	1999
18	69420	Ray	Zimmerman	2017
19	90210	Kappa	Pride	2018
20	314158	Mr.	Pie	1000
21	999831	John	Cena	2003
22	89990	BoJack	Horseman	2012

Now, let's find out how many courses were passed per semester (plus year)

```
In [7]: q3a = """
SELECT year, semester, count(*)
FROM courses natural join grades
WHERE grade > 0
GROUP BY year, semester
"""

cur.execute(q3a)
cur.fetchall()
```

```
Out[7]: [(-59, 'Fall', 2),
(-58, 'Fall', 1),
(1776, 'Summer', 4),
(1920, 'Fall', 2),
(1951, 'Spring', 1),
(1966, 'Summer', 1),
(1969, 'Spring', 1),
(1986, 'Summer', 1),
(1993, 'Spring', 2),
(1994, 'Fall', 1),
(1999, 'Spring', 2),
(2002, 'Fall', 2),
(2009, 'Spring', 1),
(2013, 'Fall', 1),
(2016, 'Fall', 3),
(2016, 'Spring', 1),
(2017, 'Fall', 1),
(2017, 'Spring', 3)]
```

Let's create a view called 'alldata' that compiles student grades, and show the view using a dataframe.

```
In [8]: cur.execute("DROP VIEW IF EXISTS allgrades")
q4c = """
create view allgrades as
SELECT s.firstName, s.lastName, m.major as ms,
       c.number, c.major as mc, g.grade
FROM students as s, majors as m, grades as g, courses as c
WHERE s.sid = m.sid AND g.sid = s.sid AND g.cid = c.cid
"""
```

```
cur.execute(q4c)
```

```
pandas.DataFrame(cur.execute("select * from allgrades").fetchall(), columns=[column[0] f
```

```
Out[8]:
```

	firstName	lastName	ms	number	\
0	Peter	Weiner	Women's Studies	8	
1	Peter	Weiner	Women's Studies	13	
2	Peter	Weiner	Women's Studies	1567	
3	Peter	Weiner	Women's Studies	1111	
4	Ya	Boi	Underwater Basket Weaving	420	
5	Ya	Boi	Underwater Basket Weaving	1113	
6	Ya	Boi	Underwater Basket Weaving	2011	
7	Michael	Scott	Paper Supplies	1	
8	Julius	Caesar	Classics	1568	
9	Julius	Caesar	Classics	1567	
10	Julius	Caesar	MILT	1568	
11	Julius	Caesar	MILT	1567	
12	Abraham	Lincoln	Theatre	4	
13	Mark	Brandanowitz	Urban Planning	2	
14	Mark	Brandanowitz	Urban Planning	2	
15	Michelle	Young	Film Study	2000	
16	Grace	Hopper	American Sign Language	2	
17	Grace	Hopper	American Sign Language	420	
18	Grace	Hopper	American Sign Language	1656	
19	Grace	Hopper	American Sign Language	1313	
20	Grace	Hopper	American Sign Language	73652	
21	Grace	Hopper	Anthropology	2	
22	Grace	Hopper	Anthropology	420	
23	Grace	Hopper	Anthropology	1656	
24	Grace	Hopper	Anthropology	1313	
25	Grace	Hopper	Anthropology	73652	
26	Grace	Hopper	CS	2	
27	Grace	Hopper	CS	420	
28	Grace	Hopper	CS	1656	
29	Grace	Hopper	CS	1313	
30	Grace	Hopper	CS	73652	
31	Mark	Wahlberg	Communications	8	
32	Mark	Wahlberg	Communications	101	
33	Thug	Nugget	Basketball	1069	
34	John	Cash	Music	101	
35	John	Cash	Music	1101	
36	Margaret	Mead	Anthropology	1313	
37	Margaret	Mead	Anthropology	73652	
38	First	Last	Stuff	245	
39	Kappa	Pride	KappaPridding	13	
40	Mr.	Pie	Home Economics	1999	
41	Mr.	Pie	Seinfeld	1999	
42	John	Cena	U Can't See Me	15	
43	BoJack	Horseman	Justice	8	

6

1.4.5 Tasks

ATTENTION: Use this notebook only to test and debug your queries, NOT as the submission.

T1) Show how many courses were passed (grade>0) per student per semester (plus year). Show student id, year, semester and the count. Sort the results by student id, year and semester.

```
In [9]: cur.execute("""  
  
        """)  
        cur.fetchall()
```

```
Out[9]: []
```

T2) Same as T1, but show student first and last name instead of student id. Also only show results for students passing at least two courses for every semester. Sort the results by first name, last name, year and semester.

```
In [10]: cur.execute("""  
  
        """)  
        cur.fetchall()
```

```
Out[10]: []
```

T3) Show the students that have failed at a course in their majors (firstName, lastName, major, courseNumber), utilizing the 'allgrades' view. Sort the results by first name, last name, major and courseNumber.

```
In [11]: cur.execute("""  
  
        """)  
        cur.fetchall()
```

```
Out[11]: []
```

T4) Same as T3, but without utilizing the view.

```
In [12]: cur.execute("""  
  
        """)  
        cur.fetchall()
```

```
Out[12]: []
```

T5) Show the professors in decreasing order of 'success' (professor, success). Success will be defined as the number of students passing any of the courses with grade >= 2. Sort by success in descending order and professor in ascending order.

```
In [13]: cur.execute("""  
  
        """)  
        cur.fetchall()
```

Out[13]: []

T6) Show a report of the courses (course_number, student_names, avg_grade). Column 'student_names' will contain the first and last names (seperated by a space) of all students taking the course, each name being seperated by ',' (eg. 'John Doe, Mary Jane'). Only students that passed a specific course (grade>=2) will be considered. Also, the report should only contain courses with avg_grade > 3. Sort the results by avg_grade, student_names and course_number.

```
In [14]: cur.execute("""  
        """)  
        cur.fetchall()
```

Out[14]: []

```
In [15]: cur.close()  
        con.close()
```